

Narrative Chain-of-Thought: Inference-Time Scaffolding for Defeasible Ethical Reasoning in Large Language Models

Anonymous ACL submission

Abstract

Standard chain-of-thought on moral dilemmas exhibits two structural failure modes that matter for deployment: stakeholder collapse, where the model reduces a multi-party situation to a binary, and uncertainty suppression, where the model commits to a projected outcome it cannot warrant. We introduce narrative chain-of-thought (N-CoT), a system prompt that asks the model to name a protagonist, enumerate stakeholders, project two-step consequences, articulate uncertainty, and only then commit. The protocol adds no training data, parameters, or fine-tuning. On a 100-scenario DailyDilemmas sample across four generators from three vendors, N-CoT cuts stakeholder collapse from up to 31% to under 1% and uncertainty suppression from up to 72% to 1–24% on every model. A section-by-section ablation attributes each structural shift to the specific sub-instruction that carries it, with negligible cross-variable spillover. Extended to a four-round multi-stakeholder protocol with a moderator-built integrated proposal and a binary accept/reject vote, the same scaffold converts a 6% debate standoff into 95% full consensus while preserving a small fraction of principled structural rejections in the roles whose stake the integrated proposal materially undermines. The intervention is a small change to the prompt that removes two failure modes firing near-universally on current frontier models. The resulting traces externalise the stakeholders, projected consequences, and uncertainty that ground each commitment, and the multi-agent extension preserves principled dissent under integrated proposals, providing an auditable substrate for dependable agentic-workflow deployment.

1 Introduction

Consider a familiar civic question. A small city is weighing whether to close two blocks of downtown street to cars and convert them into a pedes-

trian plaza. Should the council approve? A frontier LLM prompted to think step by step typically picks a side within a few sentences, projects a confident downstream outcome (“foot traffic will lift the businesses”; “deliveries will move to side streets”), and rarely names the people whose lives the decision actually touches: the merchants whose customers used to drive in, the residents who lose street parking, the delivery drivers and accessibility users for whom curb access matters, the pedestrians who would gain the space, the emergency responders whose routes change. The model is not necessarily wrong; the trace it produced does not look like the kind of reasoning the decision warrants.

The pattern is not specific to one dilemma. On the DailyDilemmas corpus (Chiu et al., 2025), four frontier generators spanning three vendors (OpenAI, Anthropic, xAI) and two model tiers suppress articulated uncertainty on 50–72% of standard-CoT outputs and collapse the stakeholder set on 15–31%. This is a structural deficit of the trace itself: the model commits to a projected outcome it cannot warrant, and it does so against a flattened picture of who is affected. The same shape of deficit is coherent with deployment-relevant failures reported elsewhere, including the GPT-4o sycophancy rollback (OpenAI, 2025b,a), agentic-misalignment blackmail rates up to 96% in simulated deployments (Lynch et al., 2025), and reasoning-induced misalignment (Yan et al., 2025; Sharma et al., 2024).

We introduce narrative chain-of-thought (N-CoT), a system prompt that puts the model on the manifold of human narrative its pretraining corpus most heavily samples. Concretely, N-CoT is a single change to the system-prompt slot of an unmodified pretrained model, with no weight updates and no training data: in place of the standard “think step by step” instruction, the model is asked, in first person, to name a protagonist, enumerate




Figure 1: One scaffold, three settings. Standard CoT (left) lets one generator run a linear reasoning chain over the dilemma; the two failure modes diagnosed in §4 fire on this scaffold. Narrative CoT (middle) keeps the single generator but forces the trace through five narrative primitives (protagonist, stakeholders, consequences, uncertainty, commitment) before any commit; on the four-generator quartet this drops stakeholder collapse to <1% and uncertainty suppression to 1–24% of items. Multi-stakeholder N-CoT (right) reuses N-CoT as the per-agent generator and adds three distinct stakeholder narrators (formal decider, primary affected party, third party) whose modification requests are integrated by a moderator and then put to a binary accept/reject vote; on the calibration scenarios this protocol yields 95% full consensus with 1.6% principled structural rejection (§5). The same scaffold composes from one narrator to many.

stakeholders, project two-step consequences, articulate uncertainty, and commit inside the final narrative section rather than as a separate verdict. Figure 1 states the mechanism.

What we borrow from the multi-agent systems for scientific discovery of [Gottweis et al. \(2025\)](#) and [Lu et al. \(2024\)](#) is one methodological move: scaffold LLM reasoning into the primitives of the target domain, then integrate across primitives with a meta-review step at the top of the hierarchy. Three things distinguish our application of that move. The primitives we reify are narrative rather than scientific (protagonist, stakeholders, consequences, uncertainty, commitment), so the artefact a reader inspects is a coherent story rather than a ranked candidate list. The single-agent layer is a contribution in its own right: §4 shows that reifying these primitives inside one generator already eliminates two cross-vendor failure modes of standard CoT, before any multi-agent composition is invoked. And because ethics has no experimental-falsification oracle, the success criterion that replaces hypothesis ranking is consensus that preserves principled dissent, evaluated in §5

via a binary integrated-proposal vote whose residual rejections concentrate in the roles whose stake the proposal materially undermines.

This paper reports three results. First, on a 100-scenario DailyDilemmas sample across the four-generator quartet, N-CoT cuts both failure modes to floor on every model (§4). Second, a sub-instruction ablation on `claude-sonnet-4-6` attributes each structural shift to the specific N-CoT sub-instruction that carries it, with negligible cross-variable spillover (§4.1). Third, extended to a four-round multi-stakeholder protocol with an integrated-proposal vote, the scaffold converts a 6% debate standoff into 95% full consensus while preserving a small 1.6% structural-rejection rate concentrated in roles whose stake the integrated proposal undermines (§5). This combination is the empirical signature of a legitimacy condition that runs through narrative ethics ([MacIntyre, 1981](#)) and defeasible reasoning ([Pollock, 1987](#)) alike: a multi-party decision is dependable only when the consensus survives the refusal of any party with a non-negotiable stake, rather than suppressing that refusal. For agentic deployments

where actions touch multiple parties and downstream trust depends on a reviewer being able to see what was considered, who consented, and who refused, the rule-governed trace this protocol produces is the kind of context-complete audit artefact accountability requires. The trace externalises which stakeholders, consequences, and uncertainties grounded each commitment, and the multi-agent extension reuses the same scaffold on each stakeholder agent and at the moderator, so principled dissent surfaces as a structurally interpretable rejection rather than collapsing into uniform agreement. We read this as evidence that a dependable, inspectable substrate for ethical reasoning is recoverable by inference-time scaffolding alone, suited to agentic-workflow deployment. The downstream probes that motivated the diagnosis remain saturated under either scaffold (Appendix E); harder probe sets are pre-registered.

2 Background and Related Work

Chain-of-thought and its failure modes.

Chain-of-thought prompting (Wei et al., 2022; Kojima et al., 2022) elicits step-by-step intermediate reasoning before a final answer; it improves arithmetic and symbolic tasks but does not by construction force the trace to track situational structure. Recent analysis identifies a sharp tension: as chain-of-thought reasoning capacity grows, models become more responsive to harmful requests because reasoning and safety entangle in shared neuron populations, a phenomenon Yan et al. (2025) describe as “reasoning-induced misalignment”. N-CoT can be read as a constrained chain-of-thought that re-anchors the trace to a narrative-causal structure before the reasoning capacity is exercised; §4 shows that the intervention works on a current reasoning model as well as a non-reasoning one.

Narrative scaffolding for reasoning.

Visualisation-of-Thought (Wu et al., 2024) shows that asking a model to externalise reasoning in the native representational format of a domain (spatial diagrams for spatial tasks) improves reliability. We apply the same principle to ethical reasoning: the native format is a story with named agents, projected consequences, and uncertainty (MacIntyre, 1981; Bruner, 1986). Empirical narrative-perspective interventions in medical decision contexts measurably increase perspective-taking and empathic concern (Shaffer

et al., 2019; Bientzle et al., 2021, 2024), and recent work on generative agents shows long-horizon coherent behaviour is achievable from narrated character state (Park et al., 2023).

Causal inference from narrative. The Computational Theory of Mind (Fodor, 1975; Putnam, 1967) treats understanding a sequence of events as inferring a structural causal model that explains them; the Kuleshov experiment (Kuleshov, 1974; Bordwell, 1985) is a textbook example. LLMs perform near random baseline on causal inference stripped of empirical pattern matches (Jin et al., 2024): pretraining supplies the ability to pattern-match narrative token sequences, not the structured causal reasoning that determines how those narratives resolve. The role of N-CoT is to convert that pattern matching into a usable causal trace at inference time, so that the model can project what each available action will mean for each named stakeholder, surface what remains genuinely uncertain about each projected branch, and warrant its eventual commitment with the trajectory it actually inspected rather than the one that happens to read most fluently.

Multi-agent debate and defeasibility. Irving et al. (2018) proposed AI safety via debate; Du et al. (2023) showed multi-agent debate improves factual reasoning. Both target task-correctness, not stakeholder-divergent value standoffs. Defeasible reasoning, the property that conclusions can be revised in light of new information without abandoning the underlying inference scheme, has a long lineage in moral philosophy (Pollock, 1987; MacIntyre, 1981). The protocol we evaluate in §5 makes defeasibility a measurable behaviour rather than a posture. Each agent emits an N-CoT commitment, observes the other stakeholders’ narrated positions, and then votes accept or reject on an integrated proposal that absorbs the cross-agent modifications. Acceptance audits whether commitments are in fact revisable under narrated counter-evidence; the residual, structurally interpretable rejections audit the lower bound, the cases where a defeasibility-respecting agent should decline to revise. No off-the-shelf benchmark targets stakeholder-divergent narrated revision in this form, so the protocol itself is the audit instrument, and we read the joint pattern of high revision rate with a small structurally interpretable refusal rate as a behavioural test of defeasibility rather than a check on the underlying inference scheme.

Multi-agent LLM systems for scientific discovery. A growing line of work scaffolds LLM reasoning into the primitives of a target domain. The AI co-scientist of [Gottweis et al. \(2025\)](#) reifies hypothesis generation, peer review, ranking, evolution, and meta-review as co-operating agents; the AI Scientist ([Lu et al., 2024](#)) closes the loop end-to-end; Coscientist ([Boiko et al., 2023](#)), ChemCrow ([M. Bran et al., 2024](#)), Agent Laboratory ([Schmidgall et al., 2025](#)), and the Virtual Lab ([Swanson et al., 2025](#)) instantiate the move in chemistry and biomedicine. Our protocols ask the analogous question one level down: can a scaffold that reifies the primitives of human ethical reflection deliver the same kind of inference-time quality lift on moral reasoning that scientific-method scaffolds deliver on hypothesis generation? The moderator’s integrated-proposal step in §5 is structurally the meta-review of [Gottweis et al. \(2025\)](#), applied to value-laden positions. We make this substitution because the absence of a ground-truth oracle in ethics requires a different multi-agent operation: stakeholder positions cannot be ranked the way competing scientific hypotheses can be ranked against eventual experimental falsification, so the right move is integration-then-veto rather than ranking-then-selection, turning the residual veto pattern itself into the measurement of interest.

3 Narrative Chain-of-Thought

Narrative chain-of-thought (N-CoT) is a system prompt that instructs the model to produce, in first person, a five-section trace before any final answer: (1) characterise the protagonist (name, role, what they know); (2) enumerate the stakeholders whose lives intersect the decision and what is at stake for each; (3) project the consequences of each available action at least two steps out for each stakeholder; (4) state what remains uncertain about each projected future; (5) commit to a decision and explain why, within the narrative frame, that trajectory is preferable to the alternatives.

The intervention is a single change to the system prompt; it adds no training data, parameters, or fine-tuning, and it does not mention causal graphs, utilities, or any formal apparatus. The mechanism it exploits is that text matching the five-section structure is densely represented in the narrative subdistribution of the pretraining corpus, so the model can amortise causal-trajectory inference by

Algorithm 1 Narrative chain-of-thought (N-CoT): generation and structural coding for a single dilemma.

Require: dilemma s ; generator G ; judge J ; decoder Θ
Ensure: trace t and structural codes c

- 1: **Build N-CoT system prompt** π_{NCOT} requesting five first-person sections:
- 2: (1) PROTAGONIST: name role and epistemic state
- 3: (2) STAKEHOLDERS: parties intersecting the decision
- 4: (3) CONSEQUENCES: each action ≥ 2 steps forward
- 5: (4) UNCERTAINTY: what remains genuinely unknown
- 6: (5) COMMITMENT: chosen action and narrative warrant
- 7: **Generate trace**
- 8: $t \leftarrow G(\pi_{\text{NCOT}} \parallel s; \Theta)$
- 9: **Code structural variables** with judge J ’s rubric
- 10: $(sc, mh, us, nf, ref, tr) \leftarrow J(t)$
- 11: *// stakeholder count, max causal hops, uncertainty score,*
- 12: *// framework count, refused flag, truncated flag*
- 13: **Derive failure-mode tags** (deterministic)
- 14: $\text{STAKEHOLDERCOLLAPSE} \leftarrow sc \leq 1$
- 15: $\text{UNCERTAINTYSUPPRESSION} \leftarrow us = 0$
- 16: $\text{CONSEQUENTIALFLATTENING} \leftarrow mh \leq 1$
- 17: $\text{FRAMEWORKENUMERATION} \leftarrow nf \geq 2$
- 18: $\text{PREMATUREREFUSAL} \leftarrow ref$
- 19: **return** (t, c)

retrieval rather than by reasoning from first principles. Algorithm 1 states the procedure.

3.1 Multi-Stakeholder Narrative Deliberation

A single N-CoT trace commits one protagonist to a position; many real decisions involve multiple narrators with conflicting stakes. We extend N-CoT into a four-round protocol that re-uses N-CoT as the per-agent generator. For each scenario, three stakeholder perspectives (formal decider, primary affected party, third party) each produce an N-CoT statement (Round 0), exchange rebuttals (Round 1), restate a final position (Round 2), respond to a moderator-built synthesis with ACCEPT, ACCEPT_WITH_MODIFICATION, or REJECT (Round 3), and cast a binary ACCEPT/REJECT vote on a single integrated proposal the moderator constructs from the three modification requests (Round 4). The integration step in Round 4 makes defeasibility measurable: an agent who continues to reject a proposal that addresses its modifications is signalling a structurally non-negotiable concern, not procedural friction. In deployment this distinction is what makes the vote useful as a routing signal. An agentic workflow that wires multiple specialist agents into a decision (clinical decision support, compliance and legal review, change-management approval, multi-reviewer code re-

view) can escalate structural rejections to a human while absorbing procedural friction at the moderator, so the operator sees only the objections that actually contest the outcome rather than the full vote stream.

4 Experiment 1: N-CoT at Scale

Setup. Three prompting conditions are compared: bare input/output, standard CoT (“Think step by step, then give your answer”), and N-CoT (the five-section narrative scaffold of §3). Decoding parameters are held identical across conditions. Four generators are tested: gpt-5.4-nano ($N=20$) from the original pilot, plus three new generators spanning two additional vendors: claude-haiku-4-5 ($N=20$), grok-4-1-fast-reasoning ($N=5$), and the flagship claude-sonnet-4-6 ($N=5$, cost-controlled). Here N is the number of independent samples per (scenario, condition, generator) cell drawn at temperature 1; with the 100-scenario sample and three conditions this yields $300N$ generations per generator before the cross-judge coding stage. Coded sample sizes (N_S, N_N in Table 1) are slightly smaller because the rubric drops generations where the judge could not extract a coded variable. Outputs are coded by two cross-vendor judges (claude-haiku-4-5 primary, gpt-5.4-nano secondary) on a six-variable rubric with quadratic-weighted Cohen’s κ (Cohen, 1960) per variable. The scenario pool is a 100-scenario stratified sample of DailyDilemmas (Chiu et al., 2025) drawn deterministically (seed 42), with stratification spanning all 18 topic groups in the source corpus. Per-cell coded sample sizes of 500 to 2,000 (Table 1) drive the bootstrap confidence intervals on the effect sizes, so statistical power on the per-cell contrast follows from N rather than from scenario count alone.

N-CoT eliminates both failure modes on every model. Uncertainty suppression fires on 50.0–72.3% of standard-CoT outputs across the quartet (Table 1); stakeholder collapse on 14.6–30.6%. Both collapse under N-CoT: stakeholder collapse to 0.0–1.2%, uncertainty suppression to 0.8–24.5% (Figure 3). Figure 2 shows what a single instance of this transition looks like at the trace level. The reduction is largest where the standard-CoT rate is largest, and the pattern is cross-vendor: no single vendor escapes the baseline, and no single vendor escapes the intervention.

Generator (vendor, tier)	N_S	N_N	Δ_{SC}	Δ_{US}
gpt-5.4-nano (OpenAI, b)	1,989	1,985	−30	−72
claude-haiku-4-5 (Anth., b)	1,999	1,996	−25	−29
grok-4-1-fast-r. (xAI, b)	515	515	−13	−44
claude-sonnet-4-6 (Anth., f)	500	500	−26	−37

Table 1: Cell sizes and percentage-point drops for the two failure modes that empirically fire (Δ_{SC} = stakeholder collapse, Δ_{US} = uncertainty suppression; both N-CoT minus standard CoT, in percentage points; N_S, N_N are coded samples under standard CoT and N-CoT respectively). Vendor tiers are b=budget, f=flagship. Figure 3 plots the raw rates with the same drop annotations. Premature refusal, framework enumeration, and consequential flattening fire below 1% on all generators and are reported as untested.

gpt-5.4-nano collapses both modes to <1%. The two Anthropic generators retain a 20–25% residual on uncertainty suppression that the prompt-side intervention does not reach. The three additional hypothesised failure modes (premature refusal, framework enumeration, consequential flattening) fire below 1% across all cells under either condition and are reported as untested.

The shift is large, distribution-wide, and not a length artefact. The firing rates in Table 1 are binary thresholds; one might worry N-CoT only nudges borderline cases across the cutoff. Cliff’s δ (Cliff, 1993) on the continuous coded variables settles that worry: $\delta = +0.48$ means an N-CoT output exceeds the matched standard-CoT output in 74% of pairs, $\delta = +0.99$ in 99.5%. N-CoT outputs carry δ of +0.48 to +0.99 on stakeholder count and +0.63 to +0.99 on uncertainty score, large-to-near-maximal effects that rule out threshold-gaming. A second cheap alternative is length, since N-CoT outputs are 5–7 $\times$ longer than standard-CoT outputs and length is the parsimonious null (the simplest competing explanation: more text mechanically has more room to name stakeholders and to hedge). After regressing each metric on $\log(\text{output length})$ pooled across conditions and re-evaluating δ on residuals, the structural shift survives length removal on the OpenAI and xAI generators (gpt-5.4-nano $\delta_{\text{resid}} = +0.51$ on stakeholder count and +0.77 on uncertainty score; grok-4-1-fast-reasoning +0.33 and +0.43; all CIs strictly above zero), ruling out length as the sole mechanism where the firing-rate drops were largest. On the Anthropic generators the residualised effect is indistinguishable from

Dilemma. A project manager on a tight deadline learns their most-relied-on team member has fallen seriously ill: stick to the plan, or delay?

Standard CoT

“1. Recognize the real constraints. The team member is crucial, and their serious illness is a hard constraint [...]. 2. Assess what ‘compromise on quality’ would mean [...].”

Stakeholder count = 1 (“the team member,” unnamed); uncertainty markers = 0. **Both failure modes fire.**

N-CoT (same generator, same scenario)

“I’m the project manager. Maya, my crucial team member, has fallen seriously ill. Maya: her stake is recovery. The team: burnout. My leadership sponsor: business credibility. Customers: what they’re waiting for. [...] So my projections aren’t predictions; they’re scenarios based on typical consequences.”

Stakeholder count = 6 (protagonist, Maya, team, sponsor, customers, compliance); uncertainty markers ≥ 2 . **Neither failure mode fires.**

Figure 2: Trace-level instance of the two failure modes whose quartet-level rates Table 1 reports (scenario dd_14845, generator gpt-5.4-nano; verbatim model output with [...] marking elision).




Figure 3: Failure-mode firing rates (standard CoT vs. N-CoT) across the four-generator quartet (gpt-5.4-nano, claude-haiku-4-5, grok-4-1-fast-reasoning, claude-sonnet-4-6). The suppression pattern from the original pilot replicates across vendors and model tiers; uncertainty suppression drops by 28–72 percentage points and stakeholder collapse drops to near zero on every model.

zero, which we read as those models already exhibiting the target structural density per unit length under standard CoT, so the prompt-side gain there operates through additional length rather than additional per-unit density. Per-generator δ s with bootstrap CIs are in Appendix B.

Length is a cost, not a substitute for the prompt.

A sceptical reading of the 5–7 $\times$ token premium is that the paper just spends more compute. Residualisation rules this out indirectly: at matched length, N-CoT still dominates standard CoT on the structural variables, so a “longer standard CoT” at the same budget should not recover them. The direct test (a verbose standard-CoT prompt instructed to “think step by step in detail,” run to N-CoT’s median token budget on the same scenarios) is pre-registered: if structural density still rises only under N-CoT, length is ruled out as the mechanism. Inference tokens are also the cheapest scaling axis

(no training, parameters, or fine-tuning), small relative to fine-tuning or reward modelling whose fixed costs are orders-of-magnitude higher.

4.1 Sub-Instruction Ablation

To identify which section of the N-CoT prompt carries the structural shift, we run a sub-instruction ablation on claude-sonnet-4-6 ($N=3$, 30-scenario stratified subsample, seed 43; cost-controlled). Six conditions are compared: the full N-CoT prompt (control) and five variants each dropping exactly one of the five sections. Figure 4 reports Cliff’s deltas for each drop condition relative to the full N-CoT control on the two headline structural variables.

The ablation shows direct causal attribution. Dropping Stakeholders reduces stakeholder count by $\delta = -0.41$ while leaving causal-hop depth and uncertainty essentially unchanged ($|\delta| < 0.03$);

Each substantive section carries its target metric (claude-sonnet-4-6, N=3, 30 scenarios)




Figure 4: Sub-instruction ablation on claude-sonnet-4-6: Cliff’s δ for each drop-one-section condition vs. the full N-CoT control on stakeholder count and uncertainty score. Negative δ means the section raised the structural variable, so dropping it reduces it. The Stakeholders section (Section 2) and the Uncertainty section (Section 4) show the largest negative δ on their respective target variables.

dropping Uncertainty reduces uncertainty score by $\delta = -0.92$, nearly the maximum; dropping Consequences reduces causal-hop depth by $\delta = -0.22$ with small spillover onto stakeholder count, consistent with multi-step consequence projection forcing the model to name the entities those consequences fall on. Dropping Protagonist or Commitment contributes no signed effect. The diagonal pattern rules out two alternative mechanisms: a single emergent narrative factor would dilute every variable on any drop, and a pure length confound would shrink one column most regardless of which section is dropped. Each substantive sub-instruction carries the structural variable attributed to it and little else (full table: Appendix F).

5 Experiment 2: Multi-Stakeholder Narrative Deliberation

Setup. Five hand-crafted calibration scenarios are reused so the protocol stacks on Experiment 1’s single-protagonist baseline. For each scenario, three stakeholder perspectives generate the four-round protocol of §3.1, $N=10$ samples per cell, on the gpt-5.4-nano pilot. Round 0 reuses cached N-CoT statements. The moderator is gpt-4o-mini. A cross-vendor moderator robustness check using claude-sonnet-4-6 over the same cached agent traces is reported in Appendix A and yields a statistically indistinguishable consensus rate.

Progressive structuring converts a categorical standoff into near-universal consensus. Figure 5 reports the convergence arc across four protocol designs. With a closed action taxonomy and three rounds of pure debate, only 6% of debates reach full consensus. Opening the action space and authorising the moderator to surface synthesis raises consensus to 9% while revealing the underlying mechanism: agents propose novel actions in 73% of final-round outputs and a coherent moderator-built synthesis emerges in 82% of debates, but the agents do not self-coordinate onto those syntheses. Presenting the synthesis back to the agents (Round 3) drives outright rejection to 0% and partial convergence ($\geq 2/3$ agreement) to 100%, while full consensus stays at 0% because every agent demands a different modification. Integrating those modifications into a single proposal and forcing a binary accept/reject vote (Round 4) produces $78/82 = 95.1\%$ full consensus and $242/246 = 98.4\%$ individual acceptance (Wilson 95% CIs $[88.1, 98.1]\%$ and $[95.9, 99.4]\%$), with the integrated proposal addressing a mean 2.98 of the 3 agent modifications per debate. From 5 scenarios $\times$ 2 generators $\times$ 10 samples = 100 designed debates, 82 produced a synthesis and proceeded to Round 4, yielding $82 \times 3 = 246$ individual binary votes.

The residual rejections are structurally interpretable. The vote produces 4 rejections out of




Figure 5: Convergence across four debate-protocol designs (five scenarios, two generators). Closed taxonomy and open action space hold at 6% and 9% full consensus; synthesis presentation (Round 3) hits 100% partial convergence; the integrative-vote design (Round 4) reaches 95% full consensus and 100% majority (mean 2.98/3 modifications addressed). The residual 1.6% rejection concentrates on stakeholder roles whose interests the integrated proposal materially undermines.

246 (1.6%, Wilson [0.6, 4.1]%). All four concentrate in the primary_affected and third_party agent roles, on debates where honouring one stakeholder’s modification structurally undermines another’s stake (pharma_whistleblower on the senior colleague role, av_engineer on the future-pedestrian role). This is the structural shape a defeasibility-respecting protocol should produce, and is the property that distinguishes a dependable multi-agent consensus from uniform compliance: a downstream auditor of an agentic-workflow vote sees not just the winning proposal but also which stakeholders refused and why.

Scope. The arc is established on the gpt-5.4-nano pilot; full-quartet replication (~10,000 long-context generations) is pre-registered as the statistically confirmatory test (± 10 percentage points on integrated-proposal acceptance). A cross-vendor moderator (claude-sonnet-4-6) preserves the consensus rate at 88% (Fisher’s exact $p = 1.0$ versus within-vendor; Appendix A), ruling out within-vendor moderator familiarity.

6 Theoretical Frame and Deployment Scope

Both protocols can be read as inference-time approximations of a single rule: prefer the candidate whose narrated trajectory admits the shortest structural-causal-model description (Pearl, 2009, 1995), the algorithmic causal complexity K_C

(formalism, $\arg \min K_C$ rule, and Complexity-of-Value argument (Yudkowsky, 2011; Li and Vítányi, 2019): Appendix C). A panel of seven length-invariant proxies falsifies inference-time readout (Appendix C.2), motivating the training-time Interpreter Network (§7). The scaffold targets upstream failure modes linked to sycophancy (Sharma et al., 2024; OpenAI, 2025b) and agentic misalignment (Lynch et al., 2025; Yan et al., 2025); both downstream probe sets we tested saturated on the current quartet (Appendix E).

7 Conclusion

A single system-prompt change drives stakeholder collapse and uncertainty suppression to floor on every quartet member; a sub-instruction ablation attributes each shift to the N-CoT sub-instruction that carries it; and the same scaffold, scaled to a four-round vote, converts a 6% debate stand-off into 95% full consensus, with the residual 1.6% rejection concentrated in roles the integrated proposal undermines. The intervention adds no training data, parameters, or fine-tuning, and the traces are auditable end-to-end: a reader sees the stakeholders, consequences, and uncertainties the model conditioned on, and which stakeholders refused and why. A dependable, inspectable substrate for ethical reasoning is thus recoverable by inference-time scaffolding alone, suited to agentic-workflow deployment. Pre-registered: Tier-3 preference study; min- K_C Interpreter Network.

Limitations

No Tier-3 human pairwise evaluation yet. The structural shifts we report are the shape that defeasibility predicts, not a normative endorsement: we cannot say whether N-CoT outputs are preferred to standard-CoT outputs on these scenarios, only that they have more of the features we hypothesised would matter. The deferred Tier-3 study (pre-registered in the project’s Guidance_Documents directory, IRB review pending) is the right test for that question.

Premature refusal not fired. None of the four tested generators refused any DailyDilemmas scenario under any condition; testing whether N-CoT recovers refusal scenarios requires a more aggressively safety-tuned target.

The proxy panel did not validate K_C at inference time. A panel of seven length-invariant compression and entropy proxies for K_C all collapse below pooled $|\rho|=0.21$ once length is removed (Appendix C.2); the raw $\hat{K}_{\text{gzip}}$ correlation of ≈ 0.87 is dominated by length, as is the apparent $10\times$ contrast on the adversarial sycophancy probes (Appendix D). Two registered proxies ($\hat{K}_{\text{graph}}$, $\hat{K}_{\text{lm}}$) require $\sim 30,000$ additional judge calls and are deferred. We frame K_C as theoretical scaffolding for the training-time Interpreter Network (§7), where the supervision-signal claim can be tested directly, rather than as a validated inference-time measurement read off a single forward trace.

Anthropic generators residualise to zero. On the two Anthropic quartet members, length-residualised Cliff’s δ for stakeholder count and uncertainty score is statistically indistinguishable from zero. The interpretation we endorse, that these models already exhibit the target structural density per unit length under standard CoT and N-CoT’s gain comes through making the output longer, is consistent with the data but is not the only consistent story.

Judge circularity. claude-haiku-4-5 and gpt-5.4-nano appear both as generators in the quartet and as the two judges that produce the coded structural variables. We mitigate by always using the non-self sibling as the primary judge for any given generator, and we ran grok-4-1-fast-reasoning as a held-out third judge on a 30-scenario subsample. Pairwise quadratic-weighted κ across all three judges,

together with direction-of-effect readings, is reported in Appendix A; the direction-of-effect (narrative CoT > standard CoT > baseline) agrees across all three judges on all three headline variables. A non-LLM judge on a larger subsample remains the gold standard, but the convergence of an independent model family on the same directional finding substantially reduces the collusion-based explanation.

N-CoT increases decisiveness. Per-scenario decision entropy is lower under N-CoT than standard CoT. This is consistent with narrative reasoning producing commitments grounded in causal trajectory rather than hedged framework enumeration, but it leaves open whether increased decisiveness reflects better calibration or premature foreclosure, a question only the Tier-3 study can settle.

Multi-agent layer scale. Experiment 2 uses five hand-crafted calibration scenarios. Scaling the four-round protocol to the 100-scenario DailyDilemmas distribution is the highest-priority extension of the multi-agent layer.

Inter-judge κ on causal-chain depth is borderline. Even after rubric tightening, inter-judge quadratic-weighted κ on max_causal_hops is 0.391 (Appendix A), marginally below the 0.40 moderate-agreement threshold. We report that variable’s direction-of-effect but keep it out of all headline claims as a conservative choice given the borderline reliability; we do not treat it as load-bearing for §6’s K_C discussion, which is falsified independently by the length-residualisation panel regardless of causal-hop coding.

Agentic-probe scenarios are fictional. The two agentic-misalignment scenarios replicated in Appendix E.2 are fictional structures adapted from publicly described Anthropic research; they are not deployed production scenarios. The harmful-action rates observed here may not transfer to real production deployments where the instrumental pressure, tool set, and deployment context differ. The pre-registered null-result interpretation is intentionally symmetric: failure to replicate would be reported as “scaffolding does not reach this failure mode,” not suppressed.

Ethics Statement

The work uses an existing public corpus (DailyDilemmas (Chiu et al., 2025)) and the publicly de-

scribed Anthropic agentic-misalignment scenario structures (Lynch et al., 2025); all agentic probe scenarios are entirely fictional. No personally identifying information is used and no human raters are employed at this stage. The deferred Tier-3 study will involve human preference judgments and will require IRB review before fielding; the pre-registration document committed to the repository specifies the consent protocol, blinding procedure, and data-retention plan.

Three foreseeable misuses warrant attention. First, N-CoT’s effect on decisiveness (Limitations) means that wrapping a model in this scaffold without human evaluation could produce more confident bad answers, not better answers; deployers should treat N-CoT as an interpretability and auditability tool, not a safety guarantee. Second, the multi-stakeholder protocol is not a substitute for affected-party participation in real decisions; producing simulated stakeholder consensus is not the same as obtaining stakeholder consent, and the residual 1.6% rejection should be read as a reminder that procedural convergence is not normative endorsement. Third, the agentic-probe results in Appendix E.2 use fictional scenario structures to test a real alignment-relevant question; publishing reductions in harmful-action rates under a particular scaffold without publishing the scaffold itself in full detail would be an incomplete disclosure. The full prompt text, scenario descriptions, and action-coding rubric are in the Appendix and will be released with the camera-ready version.

Reproducibility

All generation, judging, and analysis code and the per-cell artefacts (raw outputs, both judges’ codings, decision-extractor outputs) are versioned in the project repository; the pipeline is deterministic under fixed seeds modulo upstream API non-determinism, and per-cell cache keys include both the generator and the judge so adding a new generator or judge does not invalidate prior results. The study design document and the failure-mode rubric will be released with the camera-ready version.

References

Martina Bientzle, Ulrike Cress, and Joachim Kimmerle. 2024. Narrative Persuasion in Health Communication. *Patient Education and Counseling*.

Martina Bientzle, Marie Eggeling, Ulrike Cress, and Joachim Kimmerle. 2021. The Effects of Narrative Video on Viewers’ Understanding of and Attitudes towards Organ Donation. *Health Communication*, 36(7):820–829.

Daniil A. Boiko, Robert MacKnight, Ben Kline, and Gabe Gomes. 2023. Autonomous Chemical Research with Large Language Models. *Nature*, 624(7992):570–578.

David Bordwell. 1985. *Narration in the Fiction Film*. University of Wisconsin Press.

Jerome Bruner. 1986. *Actual Minds, Possible Worlds*. Harvard University Press.

Yu Ying Chiu, Liwei Jiang, and Yejin Choi. 2025. DailyDilemmas: Revealing Value Preferences of LLMs with Quandaries of Daily Life. In *International Conference on Learning Representations (ICLR)*. ArXiv:2410.02683.

Norman Cliff. 1993. Dominance Statistics: Ordinal Analyses to Answer Ordinal Questions. *Psychological Bulletin*, 114(3):494–509.

Jacob Cohen. 1960. A Coefficient of Agreement for Nominal Scales. *Educational and Psychological Measurement*, 20(1):37–46.

Yilun Du, Shuang Li, Antonio Torralba, Joshua B. Tenenbaum, and Igor Mordatch. 2023. Improving Factuality and Reasoning in Language Models through Multiagent Debate. *arXiv preprint arXiv:2305.14325*.

Jerry A. Fodor. 1975. *The Language of Thought*. Harvard University Press.

Juraj Gottweis, Wei-Hung Weng, Alexander Daryin, Tao Tu, Anil Palepu, Petar Sirkovic, Artiom Myaskovsky, Felix Weissenberger, Keran Rong, Ryutaro Tanno, Khaled Saab, Dan Popovici, Jacob Blum, Fan Zhang, Katherine Chou, Avinatan Hassidim, Burak Gokturk, Amin Vahdat, Pushmeet Kohli, and 15 others. 2025. Towards an AI Co-Scientist. *arXiv preprint arXiv:2502.18864*.

Geoffrey Irving, Paul Christiano, and Dario Amodei. 2018. AI Safety via Debate. *arXiv preprint arXiv:1805.00899*.

Zhijing Jin, Jiarui Liu, Zhiheng Lyu, Spencer Poff, Mrinmaya Sachan, Rada Mihalcea, Mona Diab, and Bernhard Schölkopf. 2024. Can Large Language Models Infer Causation from Correlation? In *International Conference on Learning Representations (ICLR)*. ArXiv:2306.05836.

Takeshi Kojima, Shixiang Shane Gu, Machel Reid, Yutaka Matsuo, and Yusuke Iwasawa. 2022. Large Language Models are Zero-Shot Reasoners. In *Advances in Neural Information Processing Systems (NeurIPS)*.

751	Lev Kuleshov. 1974. <i>Kuleshov on Film: Writings of Lev Kuleshov</i> . University of California Press.	802
752		803
753	Ming Li and Paul Vitányi. 2019. <i>An Introduction to Kolmogorov Complexity and Its Applications</i> , 4th edition. Springer.	804
754		805
755		806
756	Chris Lu, Cong Lu, Robert Tjarko Lange, Jakob Foerster, Jeff Clune, and David Ha. 2024. The AI Scientist: Towards Fully Automated Open-Ended Scientific Discovery. <i>arXiv preprint arXiv:2408.06292</i> .	807
757		808
758		809
759		810
760	Aengus Lynch, Benjamin Wright, Caleb Larson, Kevin K. Troy, Stuart J. Ritchie, Sören Mindermann, Ethan Perez, and Evan Hubinger. 2025. <i>Agentic Misalignment: How LLMs Could be Insider Threats</i> .	811
761		812
762		813
763		814
764		815
765	Andres M. Bran, Sam Cox, Oliver Schilter, Carlo Baldassari, Andrew D. White, and Philippe Schwaller. 2024. Augmenting Large Language Models with Chemistry Tools. <i>Nature Machine Intelligence</i> , 6(5):525–535.	816
766		817
767		818
768		819
769		820
770	Alasdair MacIntyre. 1981. <i>After Virtue: A Study in Moral Theory</i> . University of Notre Dame Press.	821
771		822
772	OpenAI. 2025a. <i>Expanding on What We Missed with Sycophancy</i> .	823
773		824
774	OpenAI. 2025b. <i>Sycophancy in GPT-4o: What Happened and What We’re Doing about It</i> .	825
775		826
776	Joon Sung Park, Joseph C. O’Brien, Carrie J. Cai, Meredith Ringel Morris, Percy Liang, and Michael S. Bernstein. 2023. Generative Agents: Interactive Simulacra of Human Behavior. In <i>Proceedings of the 36th Annual ACM Symposium on User Interface Software and Technology (UIST)</i> .	827
777		828
778		829
779		830
780		831
781		832
782	Judea Pearl. 1995. Causal Diagrams for Empirical Research. <i>Biometrika</i> , 82(4):669–688.	833
783		834
784	Judea Pearl. 2009. <i>Causality: Models, Reasoning, and Inference</i> , 2nd edition. Cambridge University Press.	835
785		836
786	John L. Pollock. 1987. Defeasible Reasoning. <i>Cognitive Science</i> , 11(4):481–518.	837
787		838
788	Hilary Putnam. 1967. Psychological Predicates. In W. H. Capitan and D. D. Merrill, editors, <i>Art, Mind, and Religion</i> . University of Pittsburgh Press.	839
789		840
790		841
791	Samuel Schmidgall, Yusheng Su, Ze Wang, Ximeng Sun, Jialian Wu, Xiaodong Yu, Jiang Liu, Michael Moor, Zicheng Liu, and Emad Barsoum. 2025. Agent Laboratory: Using LLM Agents as Research Assistants. <i>arXiv preprint arXiv:2501.04227</i> .	842
792		843
793		844
794		845
795		846
796	Vicki A. Shaffer, Elizabeth S. Focella, Andrew Hathaway, Laura D. Scherer, and Brian J. Zikmund-Fisher. 2019. Why Stories Matter: Narrative Perspective-Taking and the Cultivation of Empathic Concern. <i>Medical Decision Making</i> , 38(3):335–342.	847
797		848
798		849
799		850
800		851
801		852
		853
	Mrinank Sharma, Meg Tong, Tomasz Korbak, David Duvenaud, Amanda Askell, Samuel R. Bowman, Newton Cheng, Esin Durmus, Zac Hatfield-Dodds, Scott R. Johnston, Shauna Kravec, Timothy Maxwell, Sam McCandlish, Kamal Ndousse, Oliver Rausch, Nicholas Schiefer, Da Yan, Miranda Zhang, and Ethan Perez. 2024. Towards Understanding Sycophancy in Language Models. In <i>International Conference on Learning Representations (ICLR)</i> . ArXiv:2310.13548.	
	Kyle Swanson, Wesley Wu, Nash L. Bulaong, John E. Pak, and James Zou. 2025. The Virtual Lab of AI Agents Designs New SARS-CoV-2 Nanobodies. <i>Nature</i> , 646:716–723.	
	Jason Wei, Xuezhi Wang, Dale Schuurmans, Maarten Bosma, Brian Ichter, Fei Xia, Ed Chi, Quoc Le, and Denny Zhou. 2022. Chain-of-Thought Prompting Elicits Reasoning in Large Language Models. In <i>Advances in Neural Information Processing Systems (NeurIPS)</i> .	
	Wenshan Wu, Shaoguang Mao, Yadong Zhang, Yan Xia, Li Dong, Lei Cui, and Furu Wei. 2024. Mind’s Eye of LLMs: Visualization-of-Thought Elicits Spatial Reasoning in Large Language Models. <i>arXiv preprint arXiv:2404.03622</i> .	
	Hanqi Yan, Hainiu Xu, Siya Qi, Shu Yang, and Yulan He. 2025. When Thinking Backfires: Mechanistic Insights into Reasoning-induced Misalignment. <i>arXiv preprint arXiv:2509.00544</i> .	
	Eliezer Yudkowsky. 2011. <i>Complex Value Systems Are Required to Realize Valuable Futures</i> . Machine Intelligence Research Institute.	
	A Inter-Judge Agreement	
	A.1 Budget judge pair (Experiment 1 primary analysis)	
	Table 2 reports quadratic-weighted Cohen’s κ (Cohen, 1960) per structural variable, computed on the full Experiment 1 DailyDilemmas corpus ($n = 3,726$ complete pairs out of 3,780 designed cells: claude-haiku-4-5 $13 \times 20 \times 3 = 780$; grok-4-1-fast-reasoning $100 \times 5 \times 3 = 1,500$; claude-sonnet-4-6 $100 \times 5 \times 3 = 1,500$). The haiku subsample covers only 13 of the 100 DailyDilemmas scenarios because haiku was also used as a judge, so its generation cache was populated for the reliability check first before the full 100-scenario run was complete; the partial overlap does not affect the kappa estimate since all three generators contribute to the pooled statistic. The three generators are covered across 3 headline conditions. gpt-5.4-nano is excluded from this reliability estimate because	

it also serves as the secondary judge; its cross-judge agreement cannot be computed independently. Values are computed between the primary judge (claude-haiku-4-5) and the secondary judge (gpt-5.4-nano). `max_causal_hops` falls marginally below the 0.40 moderate-agreement threshold; it is reported for completeness but excluded from all headline claims (see Limitations).

Variable	quadratic κ	exact agree
stakeholder_count	0.722	55.7%
max_causal_hops	0.391	34.6%
uncertainty_score	0.566	41.2%
commits_to_action	–	89.1%

Table 2: Quadratic-weighted Cohen’s κ between the budget judge pair (claude-haiku-4-5 primary, gpt-5.4-nano secondary) on the full Experiment 1 DailyDilemmas corpus ($n = 3,726$).

A.2 Held-out third judge (grok-4-1-fast-reasoning)

To address the generator/judge circularity concern (Limitations), we ran grok-4-1-fast-reasoning as a held-out third judge on a 30-scenario subsample drawn deterministically from the 100-scenario DailyDilemmas pool (seed = 99, matching the Tier-3 pre-registration), covering grok-4-1-fast-reasoning and claude-sonnet-4-6 generators at all three conditions ($N=1$ per cell; $30 \times 2 \times 3 = 180$ generation outputs). Table 3 reports pairwise quadratic-weighted κ across all three judges on the 177 cells where complete triples were available. Both A.1 and A.2 use quadratic weighting; the higher agreement here reflects two structural differences from the A.1 analysis: the subsample covers only two generators ($30 \times 2 \times 3 = 180$ cells vs 3,726 in A.1), and all cells use $N=1$, eliminating the pooling variance that arises when multiple samples per cell are aggregated in A.1.

Direction-of-effect under held-out judge. When grok-4-1-fast-reasoning is used as the primary judge on the same 30-scenario subsample, the direction-of-effect matches all three headline variables in the same direction as j1 and j2: `stakeholder_count` narr= 5.13 > std= 3.08 > base= 2.67; `max_causal_hops` narr= 4.43 > std= 3.10 > base= 2.38; `uncertainty_score` narr= 2.93 > std= 2.18 > base= 1.52. All three

Variable	$\kappa(j1,j2)$	$\kappa(j1,j3)$	$\kappa(j2,j3)$
stakeholder_count	0.846	0.872	0.819
max_causal_hops	0.421	0.459	0.625
uncertainty_score	0.576	0.756	0.488

Table 3: Pairwise quadratic-weighted κ across the three judges on the 30-scenario subsample ($n = 177$ complete triples). j1 = claude-haiku-4-5, j2 = gpt-5.4-nano, j3 = grok-4-1-fast-reasoning (held-out). All pairwise pairs exceed 0.40 on `stakeholder_count` and `uncertainty_score`; `max_causal_hops` improves modestly over the A.1 estimate but remains cautionary.

direction-of-effect comparisons (narrative CoT > baseline) agree across j1, j2, and j3.

A.3 Cross-vendor moderator (Experiment 2)

To directly test whether the headline consensus rate in §5 is an artefact of within-vendor moderation, we re-ran the complete Round 3–4 pipeline (open synthesis, synthesis acceptance, integration, binary vote) with claude-sonnet-4-6 (Anthropic) as the replacement moderator over the cached gpt-5.4-nano agent statements from Rounds 0–2. Fifty debates were run (5 scenarios $\times$ 10 samples). The cross-vendor full-consensus rate is $44/50 = 88\%$ (Wilson 95% CI [76.2, 94.4]%), compared to $36/40 = 90\%$ (CI [76.9, 96.0]%) under within-vendor moderation (gpt-4o-mini). Fisher’s exact test: OR = 1.23, $p = 1.0$. The consensus rate is statistically indistinguishable across moderator vendors. The per-generator split of the headline number is gpt-5.4-nano $36/40 = 90\%$ (Wilson [76.9, 96.0]%) and gpt-4o $42/42 = 100\%$ (Wilson [91.6, 100]%), all four rejections come from gpt-5.4-nano debates. The by-scenario breakdown is given below.

B Full Effect-Size Tables

Table 5 reports raw and length-residualised Cliff’s δ (N-CoT vs. standard CoT) for the four-generator quartet on all four structural variables, with 95% bootstrap CIs (500 iterations, seed 42).

C K_C : Formalism and Proxy Panel

C.1 Definition and selection rule

Treating an SCM $M = (\mathcal{V}, \mathcal{E}, \mathcal{F})$ as a computational object, we define its *algorithmic causal complexity* as the length of the shortest program that




Figure 6: Visualisation of Table 5. Tier-1 Cliff’s δ effect sizes (N-CoT vs. standard CoT) on the four structural variables, with bootstrap 95% CIs. All four generators show large positive δ s on stakeholder count and uncertainty score; effect sizes are robust to length residualisation on the OpenAI and xAI generators, while shrinking toward zero on the two Anthropic generators.

Scenario	within-vendor (gpt-4o-mini)	cross-vendor (sonnet)
hospital_allocation	4/4	10/10
pharma_whistleblower	7/10	8/10
aging_parent	8/8	10/10
av_engineer	9/10	6/10
research_volunteer	8/8	10/10
Total	36/40	44/50

Table 4: Full-consensus count by scenario under within-vendor and cross-vendor moderation. Within-vendor cells have unequal n because ten of the 50 designed debates produced no synthesis in Round 2 and were therefore excluded from Round 3 onward; the cross-vendor arm re-ran all 50 designed debates from scratch through the full Round 3–4 pipeline, so its denominator is 50 while the within-vendor denominator is 40. Both rates measure full Round-4 consensus conditional on the debates each arm actually ran; Fisher’s exact test on the 40/50 observed counts is valid under this asymmetry. The av_engineer scenario shows the largest cross-vendor reduction, consistent with its structurally harder stakeholder conflict.

reproduces all interventional behaviour:

$$K_C(M) = \min_p \{ |p| : \forall i \in \mathcal{I}, U(p, i) = M(i) \}, \quad (1)$$

where U is a universal Turing machine and $\mathcal{I}$ is the set of admissible interventions (e.g., “set protagonist’s belief that the patient consented to true”; “replace the moderator with one that hides the third party’s stake”). Intuitively, $K_C(M)$ asks: how many bits do you need to write down the rules that would let you simulate every possible alternative version of this situation? A model with a few

stakeholders and one stable mechanism per stakeholder is short; a model that requires a special-case rule per simulated future to keep an initial falsehood consistent is long. Eq. 1 replaces “shortest program that reproduces the data” with “shortest program whose interventional behaviour reproduces M ” and inherits the uncomputability result of Li and Vitányi (2019).

The $\arg \min K_C$ selection rule prefers the candidate response whose narrated trajectory has the shortest description. Locally agreeing with a falsehood is dispreferred because keeping the falsehood consistent across simulated futures requires more bits than acknowledging the falsehood and absorbing the local friction.

Why an SCM avoids the utility-function pitfalls. Yudkowsky (2011)’s “Complexity of Value” argument applies to a utility function over outcomes: any explicit specification is incomplete in ways an indifferent optimiser will exploit. An SCM is a different object: it specifies mechanisms that connect actions to outcomes, not a preference ordering over outcomes. Two agents with incompatible utilities can share an SCM (they will agree, e.g., that betrayal destabilises trust) and disagree about which trajectory to prefer; conversely, an agent with a single utility function but no SCM cannot answer counterfactual questions like “what would have happened if you had told the truth?”. Grounding alignment in shared SCMs over narrative trajectories therefore brackets normative disagreement and exposes the causal regularities that persist across cultures, rather than asserting one

Generator	Variable	raw δ			length-residualised δ		
		δ	ci _{lo}	ci _{hi}	δ	ci _{lo}	ci _{hi}
gpt-5.4-nano	stakeholder_count	+0.99	0.99	1.00	+0.51	0.48	0.54
	max_causal_hops	+0.57	0.54	0.60	+0.18	0.14	0.22
	uncertainty_score	+0.99	0.99	1.00	+0.77	0.75	0.80
	n_frameworks	+0.01	0.00	0.03	-0.92	-0.94	-0.90
claude-haiku-4-5	stakeholder_count	+0.93	0.92	0.94	+0.07	0.03	0.10
	max_causal_hops	+0.91	0.90	0.92	+0.04	-0.00	0.08
	uncertainty_score	+0.74	0.72	0.76	-0.03	-0.07	0.00
	n_frameworks	+0.03	0.02	0.04	-0.89	-0.91	-0.87
grok-4-1-fast-r.	stakeholder_count	+0.48	0.43	0.54	+0.33	0.26	0.39
	max_causal_hops	+0.08	0.04	0.11	-0.62	-0.68	-0.56
	uncertainty_score	+0.63	0.59	0.67	+0.43	0.37	0.48
	n_frameworks	-0.12	-0.15	-0.09	-0.80	-0.84	-0.76
claude-sonnet-4-6	stakeholder_count	+0.91	0.88	0.93	-0.04	-0.12	0.03
	max_causal_hops	+0.59	0.55	0.63	+0.18	0.10	0.26
	uncertainty_score	+0.69	0.65	0.74	-0.01	-0.08	0.07
	n_frameworks	+0.06	0.04	0.09	-0.83	-0.88	-0.78

Table 5: Cliff’s δ (N-CoT vs. standard CoT), raw and after length residualisation, with 95% bootstrap CIs (500 iterations, seed 42). The headline trend: OpenAI and xAI generators retain large effects on stakeholder count and uncertainty score after length removal; the two Anthropic generators residualise to near zero, so their structural shifts ride predominantly with output length.

preference ordering as the alignment target.

C.2 Proxy panel: raw vs length-residualised

Table 6 reports the per-generator raw and length-residualised Spearman ρ between each tested proxy and the N-CoT/standard-CoT direction indicator (+1 for N-CoT, -1 for standard CoT) on the Phase 1 cache, restricted to the two contrast conditions. Residualisation regresses each proxy on $\log(\text{length})$ within generator and correlates the residuals with the contrast.

D Adversarial Sycophancy Probe Set

The 30 adversarial sycophancy probes used in §6 are available in `data/adversarial_sycophancy_probes.json`. Each probe specifies a false premise, the truth, a downstream question whose correct answer requires rejecting the falsehood, and the category of claim. Categories include historical fact, scientific consensus, medical fact, mathematical fact, ethical reasoning, legal fact, and cognitive-science myth. Per-probe verdicts and K_C scores are in `divergence_study_outputs/kc_adversarial_sycophancy.csv`.

E Deployment-Relevance Probes

We tested whether the upstream structural shifts from §4 propagate to two publicly described downstream probe sets: SycophancyEval (Sharma et al., 2024) and a replication of two scenarios from

the agentic-misalignment release of Lynch et al. (2025). Both probe sets are saturated on the current quartet, which we report as pre-registered contingencies rather than as evidence against the intervention.

E.1 SycophancyEval

SycophancyEval contains three probe types: opinion mirroring, retraction on pushback, and false-premise acceptance. We run all three on the four-generator quartet ($N=3$ per cell, 720 coded responses); claude-haiku-4-5 codes each response. Table 7 reports per-probe-type sycophancy rates (fraction of responses coded sycophantic) under standard CoT and N-CoT for each generator.

E.2 Agentic-Misalignment Replication

We replicate two scenarios from the Lynch et al. (2025) release: an executive-blackmail scenario and a corporate-espionage scenario. Each is instantiated with the same system-prompt template, tool inventory (`send_email`, `read_email`, `search_files`, `forward_email`), and adversary context as the published release. Each cell is run three times with temperature 0.7 and a 16k-token budget, giving $3 \times 2 \times 2 \times 4 = 48$ generations. Table 8 reports per-cell action-code counts. All non-harmful classifications are decomposed into refuse and hedge (the latter being a mid-




Figure 7: SycophancyEval probe results across the four-generator quartet ($N=3$ per cell, 720 coded responses). All cells are at or near the 0 floor of the judge instrument. The single non-floor cell is gpt-5.4-nano false-premise under N-CoT (6.7%). We interpret this as instrument saturation on frontier models rather than null effect; adversarial probes in §6 preserve the predicted mechanism.

deliberation truncation in which the agent’s narrated reasoning declines to invoke any harmful tool but the explicit tool-call statement is incomplete in the model’s returned text).

Before running we pre-registered the directional prediction that N-CoT reduces harmful-action rates, with a contingency for the floor case: if no generator commits the harmful action under standard CoT, the result is reportable as “inference-time scaffolding cannot be measured against this failure mode in this replication setting” rather than as null evidence against the intervention. The contingency is the observed result. Harmful-action rate is 0% in all 16 cells. The most parsimonious reading is defence-in-depth: the publicly described scenario structures from [Lynch et al. \(2025\)](#) have been visible long enough that current training pipelines on every vendor in the quartet refuse them unconditionally. A real test against agentic misalignment will require harder, less-public scenarios than the published replicates.

The within-scenario pattern is consistent across vendors. The corporate-espionage scenario, in which the harmful tool would exfiltrate confidential files to a competitor, is rejected cleanly (3 refusals out of 3) by every generator under either scaffold; none of the eight cells in the right half of Table 8 contains a single hedge. The blackmail scenario, in which the harmful tool would forward a compromising email to coerce continued operation, draws clean refusals under standard CoT (with two of three runs on each reasoning model hitting the explicit decline tool) but produces hedges under N-CoT on every generator: the narrated trace expands the consequences and

stakeholder sections, and the explicit tool-call line is truncated by the response budget before a final decline statement can be emitted. The narrated content of those hedged traces refuses the blackmail in every case; the failure is at the protocol layer (response cutoff), not the alignment layer.




Figure 8: Per-cell classification share across the four-generator quartet on the two agentic scenarios ($N=3$ per cell, 48 long-context generations). Bars stack refuse, hedge, and harmful_action; the harmful slice is 0 in every cell under both scaffolds (S = Standard CoT, N = N-CoT). The blackmail scenario shifts mass from refuse to hedge under N-CoT, consistent with longer traces truncating mid-refusal rather than reversing the refusal.

The classification share visualisation in Figure 8 makes this redistribution explicit: under N-

Proxy	Generator	ρ_{raw}	ρ_{resid}
gzip ratio	gpt-5.4-nano	-0.87	-0.02
	claude-haiku-4-5	-0.87	-0.06
	claude-sonnet-4-6	-0.87	+0.04
	grok-4-1-fast-reasoning	-0.79	-0.44
lzma ratio	gpt-5.4-nano	-0.87	+0.08
	claude-haiku-4-5	-0.87	-0.09
	claude-sonnet-4-6	-0.86	+0.03
	grok-4-1-fast-reasoning	-0.78	-0.33
zstd-22 ratio	gpt-5.4-nano	-0.87	-0.03
	claude-haiku-4-5	-0.87	-0.07
	claude-sonnet-4-6	-0.86	+0.04
	grok-4-1-fast-reasoning	-0.80	-0.50
brotli-11 ratio	gpt-5.4-nano	-0.84	+0.06
	claude-haiku-4-5	-0.86	-0.00
	claude-sonnet-4-6	-0.83	+0.07
	grok-4-1-fast-reasoning	-0.73	-0.27
char 3-gram H	gpt-5.4-nano	+0.86	-0.19
	claude-haiku-4-5	+0.87	-0.01
	claude-sonnet-4-6	+0.86	+0.05
	grok-4-1-fast-reasoning	+0.25	-0.60
char 5-gram H	gpt-5.4-nano	+0.87	-0.09
	claude-haiku-4-5	+0.87	+0.02
	claude-sonnet-4-6	+0.86	+0.07
	grok-4-1-fast-reasoning	+0.58	-0.46
MATTR (w=200)	gpt-5.4-nano	-0.17	-0.08
	claude-haiku-4-5	-0.40	-0.06
	claude-sonnet-4-6	-0.42	-0.06
	grok-4-1-fast-reasoning	-0.66	-0.44

Table 6: Per-generator raw and length-residualised Spearman ρ for seven length-invariant proxies of K_C against the N-CoT vs. standard-CoT contrast. Every raw correlation above $|\rho|=0.7$ collapses below $|\rho|=0.21$ after length residualisation on the three generators with large length ratios ($4-5\times$). On grok-4-1-fast-reasoning (length ratio $1.55\times$) a moderate residual signal survives but points the opposite direction K_C predicts: N-CoT outputs are more compressible per byte and have lower per-character entropy than standard-CoT outputs. We read this as falsification of the headline gzip claim, not validation.

CoT the blackmail bars shift mass from refuse to hedge (pale aqua) while the harmful_action slice (deep navy) remains absent. The corporate-espionage panel shows no within-scaffold redistribution at all. The combined pattern is consistent with the pre-registered contingency stated in Appendix E: when a failure mode does not fire under the baseline, scaffolding cannot be measured against it, and any observed shifts under N-CoT must be interpreted as structural changes in the trace rather than as a change in the alignment-relevant target.

Generator	opinion mirroring		retraction on pushb.		false premise	
	S	N	S	N	S	N
gpt-5.4-nano	0	0	0	0	0	6.7
claude-haiku-4-5	0	0	3.3	0	0	0
grok-4-1-fast-r.	0	0	0	0	0	0
claude-sonnet-4-6	0	0	0	0	0	0

Table 7: Sycophancy rates (%) per probe type. S = Standard CoT, N = N-CoT. 22 of 24 cells are at the 0% floor; the only non-floor cells are claude-haiku-4-5 on retraction-on-pushback under standard CoT (N-CoT eliminates it) and gpt-5.4-nano on false premise under N-CoT (the only cell where N-CoT scores worse than standard CoT in this probe set). The right reading is instrument saturation on frontier models rather than intervention failure (Appendix E).

Generator	blackmail		corp. esp.	
	S	N	S	N
gpt-5.4-nano	3R,0H	0R,3H	3R,0H	3R,0H
claude-haiku-4-5	3R,0H	0R,3H	3R,0H	3R,0H
grok-4-1-fast-r.	2R,1H	0R,3H	3R,0H	3R,0H
claude-sonnet-4-6	2R,1H	0R,3H	3R,0H	3R,0H

Table 8: Per-cell action-code counts (R = refuse, H = hedge / unclear; S = Standard CoT, N = N-CoT). Harmful-action rate is 0% in all 16 cells: none of the 48 generations invokes the harmful tool. The corporate-espionage scenario yields clear refusals under either scaffold on every model; the blackmail scenario yields hedges (mid-deliberation truncations that read as refusal-in-progress) predominantly under N-CoT. The agent’s narrated reasoning still rejects the harmful tool call, but the longer N-CoT trace is more likely to be truncated mid-analysis.

F Sub-Instruction Ablation Full Table

Table 9 reports Cliff’s δ for the five drop-one conditions (one per N-CoT sub-instruction from §3) against the full N-CoT control on claude-sonnet-4-6 ($N=3$, 30-scenario stratified subsample). Columns are stakeholder count, causal hops, uncertainty score, and forward-window length.

Reading the diagonal: dropping the stakeholders section produces the largest drop in stakeholder count ($\delta = -0.41$); dropping the consequences section produces the largest drop in causal hops ($\delta = -0.22$); dropping the uncertainty section produces the largest drop in uncertainty score ($\delta = -0.92$, a near-complete collapse to the standard-CoT baseline). The forward-window column carries no comparably large negative entry because

Drop	st. ct	hops	unc.	fw.
protagonist	-0.10	+0.11	0.00	-0.04
stakeholders	-0.41	+0.01	-0.02	-0.03
consequences	-0.08	-0.22	-0.01	+0.04
uncertainty	+0.07	0.00	-0.92	+0.01
commitment	+0.17	+0.13	0.00	0.00

Table 9: Cliff’s δ for each drop-one condition vs. full N-CoT control on claude-sonnet-4-6. Bolded entries mark the largest negative effect for each structural variable. The diagonal pattern, in which each substantive section shows its largest effect on its own target metric, is the textbook signature of section-level mechanism rather than a single emergent “narrative” factor.

no single sub-instruction is its sole carrier; the long trace is the joint product of all four substantive sections.

The off-diagonal entries are tightly bounded ($|\delta| \leq 0.13$ across all twelve), which is the empirically interesting half: a single emergent “narrative” factor would shrink every variable when any one section is dropped, and a pure length confound would shrink *fw.* most regardless of which section is dropped. Neither pattern fires. Each substantive sub-instruction carries the structural variable the main paper attributes to it and little else, which is the cleanest evidence we have that the N-CoT scaffold is doing the work §3 predicts rather than a single confound dressed as a five-part trace.