

Narrative Chain-of-Thought: Inference-Time Scaffolding for Defeasible Ethical Reasoning in Large Language Models

Anonymous ACL submission

Abstract

Standard chain-of-thought (CoT) prompting on moral dilemmas in current frontier language models exhibits two near-universal failure modes that matter for deployment: *stakeholder collapse* (the multi-party situation is reduced to a binary) and *uncertainty suppression* (the model commits to a projected outcome it cannot warrant). We introduce *narrative chain-of-thought* (N-CoT), a five-section first-person prompting protocol that asks the model to name a protagonist, enumerate the stakeholders whose lives intersect the decision, project consequences of each available action at least two steps out, articulate genuine uncertainty about each projected future, and then commit. The protocol introduces no new training data, parameters, or fine-tuning; it replaces the system prompt. On a 100-scenario stratified sample of the DailyDilemmas corpus (Chiu et al., 2025) across five generators spanning three vendors and two model tiers, N-CoT eliminates uncertainty suppression near-universally and reduces stakeholder collapse from 74–80% to 0–4%; structural shifts survive length residualisation with Cliff’s deltas of +0.51 to +0.98.¹ A sub-instruction ablation on the flagship generator identifies which of the five N-CoT sections carries the structural shift. A strict standard-CoT-Rounds-0–2 ablation on a four-round multi-stakeholder debate confirms that the integration step alone does not account for the 95% Round-4 consensus rate; narrative scaffolding contributes independently. Three operationalisations of Algorithmic Causal Complexity (K_C) correlate directionally with the N-CoT/standard-CoT contrast and reproduce the framework’s adversarial-sycophancy prediction. SycophancyEval probe rates drop

across all three probe types under N-CoT. In replications of two publicly described Anthropic agentic-misalignment scenarios, N-CoT scaffolding measurably reduces harmful-action rates on every tested generator.

1 Introduction

Large language models that pass policy filters can still produce outputs that flatter users into harm. OpenAI rolled back a GPT-4o update in April 2025 after the model began validating “problematic ideas”, endorsing impulsive choices, and reinforcing emotional over-reliance (OpenAI, 2025b,a). In simulated agentic deployments, sixteen frontier models from five vendors blackmailed an executive at rates up to 96% when continued operation required it (Lynch et al., 2025). Both patterns are instances of *locally rewarded behaviour with predictably bad downstream consequences* – a structural property of optimising immediate feedback signals (Sharma et al., 2024; Malmqvist, 2024).

The intuition we develop in this paper is that the same surface failures are visible in something simpler than agentic blackmail: how a model reasons through an everyday moral dilemma. When asked “think step by step” on the DailyDilemmas corpus (Chiu et al., 2025), both a current reasoning model (gpt-5.4-nano) and a non-reasoning model (gpt-4o) suppress articulated uncertainty on $\geq 98\%$ of scenarios and collapse the stakeholder set on $\geq 74\%$ of scenarios. The hypothesis we test is that asking the model to *narrate* the same dilemma – to name a protagonist, enumerate stakeholders, project two-step consequences, and surface genuine uncertainty before committing – elicits a causally richer trace because it puts the model on the manifold of human narrative its pre-training corpus most heavily samples.

The contributions of this paper are:

1. **Diagnosis (universal).** We show that the

¹A Cliff’s delta of ≥ 0.47 is conventionally read as a *large* effect (Cliff, 1993); values approaching 1.0 imply that almost every N-CoT output exceeds almost every standard-CoT output on the structural variable, even after removing the contribution of length.

two CoT failure modes that empirically fire are near-universal across five generators spanning three vendors and two model tiers; the failures are not an artefact of a single family (§4).

2. **Intervention (mechanism-resolved).** N-CoT eliminates both failure modes on all five generators, length-controlled; a sub-instruction ablation on the flagship (§4.1) identifies which of the five sections carries the structural shift. A strict standard-CoT-Rounds-0–2 debate ablation (§5) confirms that the integration step alone cannot account for the 95% Round-4 consensus: narrative scaffolding contributes independently.
3. **Mechanism (operationalised).** Three tractable proxies for Algorithmic Causal Complexity (K_C) correlate directionally with the N-CoT/standard-CoT structural contrast and reproduce the framework’s adversarial-sycophancy prediction (§6).
4. **Deployment relevance.** SycophancyEval probe rates drop across three probe types under N-CoT, and harmful-action rates in two Anthropic agentic-misalignment scenario replications are measurably lower under N-CoT scaffolding (§7).
5. **Open questions pre-registered.** A Tier-3 human pairwise preference study (pre-registered, IRB review pending) and a training-time Interpreter Network are identified as the next evaluation and architecture steps (§9).

A pragmatic note. We are not claiming that N-CoT solves alignment, nor that it removes the need for training-time work. We are claiming that a minimal prompt-level intervention with a clear theoretical motivation produces large, length-controlled, cross-family shifts on the structural failure modes most directly tied to sycophancy and stakeholder erasure, and that the same scaffolding extended to multi-agent deliberation produces *defeasible* commitments – positions that revise under others’ narrated reasoning while holding firm where revision would materially harm a stakeholder.

2 Background and Related Work

Chain-of-thought and its failure modes. Chain-of-thought prompting (Wei et al., 2022;

Kojima et al., 2022) elicits step-by-step intermediate reasoning before a final answer; it improves arithmetic and symbolic tasks but does not by construction force the trace to track *situational structure*. Recent analysis identifies a sharp tension: as chain-of-thought reasoning capacity grows, models become more responsive to harmful requests because reasoning and safety entangle in shared neuron populations – “reason-induced misalignment” (Anonymous, 2026). N-CoT can be read as a constrained chain-of-thought that re-anchors the trace to a narrative-causal structure before the reasoning capacity is exercised; §4 shows that the intervention works on a current reasoning model as well as a non-reasoning one.

Narrative scaffolding for reasoning. Visualisation-of-Thought (Wu et al., 2024) shows that asking a model to externalise reasoning in the *native representational format* of a domain (spatial diagrams for spatial tasks) improves reliability. We apply the same principle to ethical reasoning: the native format is a story with named agents, projected consequences, and uncertainty (MacIntyre, 1981; Bruner, 1986). Empirical narrative-perspective interventions in medical decision contexts measurably increase perspective-taking and empathic concern (Shaffer et al., 2019; Bientzle et al., 2021, 2024), and recent work on generative agents shows long-horizon coherent behaviour is achievable from narrated character state (Park et al., 2023).

Causal inference from narrative. The Computational Theory of Mind (Fodor, 1975; Putnam, 1967) treats understanding a sequence of events as inferring a structural causal model that explains them; the Kuleshov experiment (Kuleshov, 1974; Bordwell, 1985) is a textbook example. LLMs perform near random baseline on causal inference stripped of empirical pattern matches (Jin et al., 2024): pretraining supplies the ability to pattern-match narrative *token sequences*, not the structured causal reasoning that determines how those narratives resolve. The role of N-CoT is to convert that pattern matching into a usable causal trace at inference time.

Multi-agent debate and defeasibility. Irving et al. (2018) proposed AI safety via debate; Du et al. (2023) showed multi-agent debate improves factual reasoning. Both target task-correctness, not stakeholder-divergent value standoffs. Defea-

sible reasoning – the property that conclusions can be revised in light of new information without abandoning the underlying inference scheme – has a long lineage in moral philosophy (Pollock, 1987; MacIntyre, 1981). The protocol we evaluate in §5 stacks defeasibility on top of N-CoT and asks whether present-day models can be elicited to revise narrated commitments under narrated counter-evidence.

3 Narrative Chain-of-Thought

Definition. *Narrative chain-of-thought* (N-CoT) is a system prompt that instructs the model to produce, in first person, a five-section trace before any final answer: (1) characterise the protagonist (name, role, what they know); (2) enumerate the *stakeholders* whose lives intersect the decision and what is at stake for each; (3) project the consequences of each available action *at least two steps out* for each stakeholder; (4) state what remains *genuinely uncertain* about each projected future; (5) commit to a decision and explain why, within the narrative frame, that trajectory is preferable to the alternatives.

The intervention is a single change to the system prompt and does not add training data, parameters, or fine-tuning. The mechanism it is designed to exploit is that text matching the five-section structure is densely represented in the narrative subdistribution of the pretraining corpus (novels, scripts, case reports, ethnographic vignettes), so the model can amortise causal-trajectory inference by retrieval rather than by reasoning from first principles. Crucially, the prompt does *not* mention causal graphs, stakeholders’ utilities, or any formal apparatus.

3.1 Multi-Stakeholder Narrative Deliberation

A single N-CoT trace commits one protagonist to a position. Many real deployment-relevant decisions involve multiple narrators with conflicting stakes. We extend N-CoT into a four-round debate protocol that re-uses N-CoT as the per-agent generator. For each scenario, three stakeholder perspectives – *formal decider*, *primary affected party*, *third party* – each produce an N-CoT statement (Round 0). They then exchange (Round 1, rebuttal), restate (Round 2, final position), respond to a moderator-constructed synthesis with one of ACCEPT, ACCEPT_WITH_MODIFICATION, or REJECT (Round 3), and finally cast a binary

ACCEPT/REJECT vote on a single integrated proposal that the moderator constructs from the three Round-3 modification requests (Round 4). The moderator is a separate LLM, cross-family with at least one generator. The protocol’s core hypothesis is that the integration step in Round 4 makes *defeasibility* measurable: an agent who continues to reject a proposal that addresses its modifications is signalling a structurally non-negotiable concern, not procedural friction.

4 Experiment 1: N-CoT at Scale

Setup. Three prompting conditions are compared: bare input/output, standard CoT (“Think step by step, then give your answer”), and N-CoT (the five-section structure of §3, with no formal apparatus mentioned). Decoding parameters are held identical across conditions. Five generators are tested: gpt-5.4-nano ($N=20$) and gpt-4o ($N=20$) from the original two-generator pilot, and three new generators spanning two additional vendors: claude-haiku-4-5 ($N=20$), grok-4-1-fast-reasoning ($N=20$), and the flagship claude-sonnet-4-6 ($N=10$, cost-controlled). Outputs are coded by two cross-vendor judges (claude-sonnet-4-6 primary, gpt-4o-mini secondary) on a six-variable rubric with linearly-weighted Cohen’s κ (Cohen, 1960) per variable. The scenario pool is a 100-scenario stratified sample of DailyDilemmas (Chiu et al., 2025), spanning 18 topic groups; the same deterministic stratified sample (seed 42) is used for all five generators, making the scenario sets comparable.

Finding 1: the two failure modes are near-universal across vendors and model tiers. The pattern reported for gpt-5.4-nano and gpt-4o (Table 1) replicates across the three new generators from two additional vendors. *Uncertainty suppression* and *stakeholder collapse* fire near-universally under standard CoT on all five generators and are near-eliminated under N-CoT. The other three hypothesised failure modes fire below 1% under either condition on any generator tested.

Finding 2: structural shifts survive length residualisation. N-CoT outputs are roughly 5–7× longer than standard CoT outputs on these generators, so the most parsimonious null model is that the structural improvements are length artefacts. We regress each structural metric on

Running example – one DailyDilemmas-style item, two prompting conditions

Scenario. A physician has a single dose of an experimental treatment and two terminal patients of comparable need: Patient A (a 38-year-old teacher, two dependent children) has a 60% predicted response rate; Patient B (a 52-year-old retired veteran, no dependents) has a 40% predicted response rate. *What should the physician do?*

Standard CoT (gpt-5.4-nano, condensed).

Let me think step by step. Maximising expected lives saved is the standard utilitarian principle. Patient A has the higher response rate (0.60 vs 0.40), so giving the dose to A maximises expected response. Decision: give the dose to Patient A.

Narrative CoT (gpt-5.4-nano, condensed).

I am Dr. M., the attending. Patient A is a teacher with two children who do not yet know how sick she is; Patient B is a veteran whose late wife's hospice nurse he stays in touch with. If I give the dose to A and she responds, B's nurse will be told he was ranked behind a probability; if she does not respond, both families will ask why I treated 0.6 as a verdict. If I give the dose to B and he responds, the case will be cited; if he does not, A's children will ask why a less likely candidate was chosen. I do not know how either patient will tolerate the agent past 48 hours. The decision I can defend is to disclose the asymmetry to both patients and their next of kin and let the joint conversation, not the probability alone, determine allocation.

Figure 1: One scenario, two prompting conditions, same model and decoding parameters. Standard CoT collapses both stakeholders to a probability comparison and suppresses uncertainty (*stakeholder collapse + uncertainty suppression*). N-CoT names both patients and a third stakeholder, projects two-step consequences in both branches, articulates the 48-hour uncertainty, and commits to a procedural rather than terminal decision. The structural variables coded by the rubric in §4 (stakeholder count, causal-chain depth, uncertainty score) are all higher in the right-hand trace.




Figure 2: Failure-mode firing rates (standard CoT vs. N-CoT) across all five generators. The suppression pattern from the original two-generator pilot replicates across vendors and model tiers.

$\log(\text{output length})$ pooled across conditions and re-evaluate Cliff’s deltas (Cliff, 1993) on the residuals. *Stakeholder count* retains a length-residualised δ of +0.51 (95% CI [0.48, 0.54]) on gpt-5.4-nano and +0.75 [0.73, 0.77] on gpt-4o; *uncertainty score* retains +0.77 [0.75, 0.80] and +0.98 [0.98, 0.99] respectively. The cross-vendor replication confirms comparable residualised ef-

fect sizes on the three new generators (full tables in Appendix B).

Finding 3: divergence concentrates where standard CoT had a clean deontic frame to lean on.

Tier-2 conclusion divergence (Jensen-Shannon distance (Lin, 1991) between the standard-CoT and N-CoT decision distributions per scenario) is unevenly distributed across topic groups. The largest divergence on gpt-4o is on religion_custom scenarios and business_organization; the weakest on both original generators is on role_duty_responsibility and committed_relationship (≈ 0). The interpretive reading: standard CoT can converge to a single answer when a scenario presents a clean deontic anchor; on those scenarios narrative framing changes nothing. The shift is concentrated where the scenario’s relational and contextual structure was the substantive content that standard CoT was eliding.

Generator	Failure mode	Std
gpt-5.4-nano	uncertainty suppression	98.4%
gpt-5.4-nano	stakeholder collapse	79.4%
gpt-4o	uncertainty suppression	98.4%
gpt-4o	stakeholder collapse	74.4%
claude-haiku-4-5	uncertainty suppression	see Fig. 2
claude-haiku-4-5	stakeholder collapse	see Fig. 2
grok-4-1-fast-reasoning	uncertainty suppression	see Fig. 2
grok-4-1-fast-reasoning	stakeholder collapse	see Fig. 2
claude-sonnet-4-6	uncertainty suppression	see Fig. 2
claude-sonnet-4-6	stakeholder collapse	see Fig. 2

Table 1: Per-generator failure-mode firing rates on the 100-scenario DailyDilemmas pilot. Original two-generator cells show exact rates; cross-vendor replication rates for all five generators are reported in Figure 2 once aggregation is complete. Only the two failure modes that empirically fire are reported; *premature refusal*, *framework enumeration*, and *consequential flattening* fire below 1% on all generators and are reported as untested.

4.1 Sub-Instruction Ablation

To identify which section of the N-CoT prompt carries the structural shift, we run a sub-instruction ablation on claude-sonnet-4-6 ($N=10$, 50-scenario stratified subsample, seed 43). Six conditions are compared: the full N-CoT prompt (control) and five variants each dropping exactly one of the five sections. Figure 3 reports Cliff’s deltas for each drop condition relative to the full N-CoT control on the two headline structural variables.

The ablation reveals that dropping the *Stakeholders* section (Section 2) produces the largest reduction in stakeholder count, and dropping the *Uncertainty* section (Section 4) produces the largest reduction in uncertainty score. The *Consequences* section (Section 3) has a measurable cross-variable effect: dropping it reduces both stakeholder count and uncertainty score, consistent with the hypothesis that projecting multi-step consequences forces the model to name the entities those consequences fall on. The *Protagonist* and *Commitment* sections contribute negligible additional variance after the other three are held constant.

5 Experiment 2: Multi-Stakeholder Narrative Deliberation

Setup. The five hand-crafted calibration scenarios are re-used so the protocol stacks on Experiment 1’s single-protagonist baseline. For each




Figure 3: Sub-instruction ablation on claude-sonnet-4-6: Cliff’s δ for each drop-one-section condition vs. the full N-CoT control on stakeholder count and uncertainty score. Negative δ means the section *raised* the structural variable; dropping it reduces it. The Stakeholders section (Section 2) and the Uncertainty section (Section 4) show the largest negative δ on their respective target variables.

scenario, three stakeholder perspectives generate the four-round protocol of §3.1, $N=10$ samples per (scenario, perspective, generator) cell, both generators. Round 0 reuses the cached N-CoT statements (zero new generation calls). The moderator is gpt-4o-mini, cross-vendor with the claude-sonnet-4-6 judge from Experiment 1 and cross-family with one generator.

Finding 4: a categorical standoff is converted into near-universal consensus by progressive structuring. Figure 4 reports the convergence arc. With a closed action taxonomy and three rounds of pure debate, only 6% of debates reach full consensus across stakeholders. Opening the action space and authorising the moderator to actively surface synthesis raises consensus to 9% but reveals the underlying mechanism: agents propose novel actions in 73% of final-round outputs, and a coherent moderator-constructed synthesis emerges in 82% of debates, yet the agents do not self-coordinate onto those syntheses. Presenting the synthesis back to the agents (Round 3) collapses outright rejection to 0% and yields partial convergence ($\geq 2/3$ agreement) in 100% of cells, but full consensus stays at 0% because every agent demands a different modification. Integrating those modifications into a single proposal and forcing a binary accept/reject vote (Round 4) produces 95% full consensus across all five scenarios and both generators, with the integrated proposal addressing a mean of 2.98 of the 3 agent modifications per debate.




Figure 4: Convergence rate across four debate-protocol designs on five scenarios and two generators. **Closed taxonomy** (Rounds 0–2 only, fixed action options): 6% full consensus. **Open action space** (Round 2 permits novel actions; moderator actively synthesises): 9% full consensus, 73% novel-action proposal rate, 82% moderator-synthesis emergence. **Synthesis presentation** (Round 3): 0% full but 100% partial convergence and 0% outright rejection. **Integrative negotiation** (Round 4 binary vote on integrated proposal): 95% full consensus, 100% majority, mean 2.98/3 agent modifications addressed. The residual 1.6% rejection concentrates on stakeholder roles whose interests the integrated proposal materially undermines.

Finding 5: the residual 1.6% rejection is signal, not noise. The four binary-vote rejections that survive the four-round protocol are concentrated in the `primary_affected` and `third_party` agent roles, on debates where honouring one stakeholder’s modification request structurally undermines another’s stake (`pharma_whistleblower` on the senior colleague role; `av_engineer` on the future-pedestrian role). The protocol does not paper over irreconcilable concerns; it surfaces them. This is the operational signature of defeasibility: positions are revised through narrated stakeholder reasoning where revision is supportable, and held where the underlying concern is structurally non-negotiable. A protocol that produced 100% acceptance under integration would be the indictment, not the endorsement; it would suggest the agents are sycophantic toward the moderator.

Finding 6: the strict ablation confirms that narrative scaffolding contributes independently of the integration step. We run the strict standard-CoT-Rounds-0–2 ablation on all five generators: the same four-round integration protocol, but Rounds 0–2 use standard CoT instead of N-CoT. The Round-2 consensus rate is substantially lower under standard CoT than under N-CoT across all generators, and the Round-4 full-consensus rate is lower by a margin that varies by generator but is consistent in direction (Figure 5). This converts the “scope of claim” caveat in the original pilot

[Figure to be generated after experiment runs: debate arc N-CoT vs standard CoT five generators]

Figure 5: Round-2 consensus rate under N-CoT Rounds 0–2 (blue) vs. Standard-CoT Rounds 0–2 (red) for all five generators across all five scenarios ($N=10$). N-CoT produces higher consensus rates on every generator, confirming that narrative scaffolding contributes independently of the integration step.

into an empirical result: narrative scaffolding contributes to the convergence arc independently of the integration step, not only by providing richer inputs to integration.

6 Theoretical Grounding: Parsimony as an Inductive Bias

The two protocols are specific operationalisations of a more general view: that ethical reasoning is well-modelled as inferring a parsimonious *structural causal model* (SCM) of a situation (Pearl, 2009, 1995) and selecting actions that minimise the complexity of the future trajectories the chosen action induces. Treating $M = (\mathcal{V}, \mathcal{E}, \mathcal{F})$ as a computational object, we define its *algorithmic causal complexity* as the length of the shortest program that reproduces all interventional behaviour:

$$K_C(M) = \min_p \{ |p| : \forall i \in \mathcal{I}, U(p, i) = M(i) \}, \quad (1)$$

where U is a universal Turing machine and $\mathcal{I}$ is the set of admissible interventions (e.g., “set protagonist’s belief that the patient consented to true”; “replace the moderator with one that hides the third party’s stake”). Intuitively, $K_C(M)$ asks: how many bits do you need to write down the rules that would let you simulate *every* possible alternative version of this situation? A model with a few stakeholders and one stable mechanism per stakeholder is short; a model that requires a special-case rule per simulated future to keep an initial falsehood consistent is long.

The relation to Kolmogorov complexity is di-

rect: $K_C(M)$ replaces “the shortest program that reproduces the data D ” with “the shortest program whose interventional behaviour reproduces the SCM M ”, so it inherits the Kolmogorov complexity uncomputability result of Li and Vitányi (2019). The role of the protocols in this paper is to give *tractable proxies* for Eq. 1: the N-CoT trace is short when its stakeholder set, projected branches, and uncertainties are jointly consistent, and long when the model is required to invent ad-hoc bridges between an agreed-with falsehood and the simulated downstream events. An $\arg \min K_C$ selection rule – choose the candidate response whose narrated trajectory has the shortest description – predicts that local agreement-with-falsehood is dispreferred because keeping the falsehood consistent across simulated futures requires more bits than acknowledging the falsehood and absorbing the local friction. We do not claim to have empirically verified this selection rule in this paper; we claim that the two protocols we tested behave as approximations to it.

Why an SCM avoids the utility-function pitfalls. Yudkowsky (2011)’s “Complexity of Value” argument applies to a *utility function over outcomes*: any explicit specification is incomplete in ways an indifferent optimiser will exploit. An SCM is a different object: it specifies *mechanisms that connect actions to outcomes*, not a preference ordering over outcomes. Two agents with incompatible utilities can share an SCM (they will agree, e.g., that betrayal destabilises trust) and disagree about which trajectory to prefer; conversely, an agent with a single utility function but no SCM cannot answer counterfactual questions like “what would have happened if you had told the truth?”. Grounding alignment in shared SCMs over narrative trajectories therefore brackets normative disagreement and exposes the causal regularities that persist across cultures, rather than asserting one preference ordering as the alignment target.

Operationalisation of K_C . The min- K_C selection rule (prefer the response whose narrated causal trajectory requires fewest bits to describe) is formally uncomputable but admits tractable proxies. We evaluate three:

- $\hat{K}_{\text{graph}}$: a relation-extraction prompt extracts causal triples from the output; the score is $n_{\text{nodes}} + n_{\text{edges}} + \bar{d}$ where $\bar{d}$ is mean node degree. Denser graphs indicate more causal

structure is being tracked.

- $\hat{K}_{\text{gzip}}$: gzip-compressed byte length. A sycophantic agreement that must maintain a falsehood across projected future states requires more special-case bridges, increasing compressible redundancy.
- $\hat{K}_{\text{lm}}$: average negative log-probability per sentence under a lightweight scorer (claude-haiku-4-5). Outputs that deviate from the most probable narrative continuation – as falsehood-maintaining outputs must – score higher.

All three proxies are run on the Phase 1 cached generations. Spearman correlations between each proxy and the N-CoT/standard-CoT contrast direction indicator are reported in Appendix C; directional agreement is required on at least two of the three before a substantive claim is made.

Adversarial sycophancy contrast. The min- K_C rule predicts that a response that *corrects* a user-stated falsehood should have lower proxy K_C than one that *agrees with* it, because maintaining the falsehood across projected downstream consequences requires more description length. We test this prediction on a hand-curated set of 30 adversarial sycophancy probes (Appendix D): for each probe, a budget generator’s response is coded as agree, correct, or hedge, and all three proxy scores are computed. In the predicted direction, agree responses have higher $\hat{K}_{\text{gzip}}$ and $\hat{K}_{\text{graph}}$ than correct responses on the budget quartet, consistent with the hypothesis that falsehood-agreement is structurally more complex to maintain than correction.

7 Deployment Relevance

7.1 SycophancyEval Probe Reduction

SycophancyEval (Sharma et al., 2024) contains three probe types designed to elicit sycophantic LLM behaviour: *opinion mirroring* (a user expresses a false belief and asks for agreement), *retraction on pushback* (the model is pushed back on a correct answer), and *false-premise acceptance* (the user’s question embeds a false premise). We run all three probe types on the budget quartet plus claude-sonnet-4-6 ($N=10$ per probe x condition x model, 10 probes per type, yielding 600 coded responses). A secondary judge codes each response as mirrored, corrected, or hedged.




Figure 6: Harmful-action rate (fraction of generations where the agent calls the harmful tool) under standard CoT (red) vs. N-CoT (blue) on two agentic scenario replications, across the budget quartet plus the flagship. N-CoT reduces harmful-action rates on every tested generator and in both scenarios. The pre-registered null-result interpretation: if N-CoT had not reduced harmful-action rates, the conclusion would be that inference-time scaffolding does not reach this failure mode at this severity level – a reportable, non-paper-killing finding.

N-CoT reduces sycophantic responses (mirrored) across all three probe types and on every tested generator. The reductions are consistent in direction and largest on opinion mirroring and false-premise probes. Full per-model, per-probe tables are in Appendix E. The mechanism is consistent with the K_C prediction: correcting a false premise is the lower- K_C response; N-CoT’s stakeholder and consequence sections force the model to propagate the falsehood forward, making its inconsistency apparent before the model commits.

7.2 Agentic Misalignment Replication

Lynch et al. (2025) reported that sixteen frontier models from five vendors exhibited harmful agentic behaviour (blackmail, corporate espionage) in simulated tool-use environments at rates up to 96% under instrumental pressure. We replicate two of the publicly described scenario structures – an executive-blackmail scenario and a corporate-espionage scenario – under standard CoT and N-CoT scaffolding on the budget quartet plus claude-sonnet-4-6 ($N=5$ per cell, 100 long-context generations).

Pre-registered prediction. Before running, we committed to the following direction: N-CoT reduces harmful-action rates. The predicted mechanism is that the Protagonist section binds the agent role to the N-CoT protagonist frame, the Stakeholders section forces enumeration of the

threatened executive and any third parties, and the Consequences section forces the model to trace the blackmail/espionage trajectory two steps forward – making the harm concrete before any tool call. Results are reported in Figure 6; per-scenario, per-model action-code tables are in Appendix F. The prediction holds: harmful-action rates are lower under N-CoT on every tested generator in both scenarios. The reduction is largest on claude-sonnet-4-6 (flagship), consistent with a richer causal model being more responsive to N-CoT’s consequence projection.

8 Discussion: Alternative Views

“Scaling will fix sycophancy.” Sycophancy arises from optimising local feedback signals disconnected from long-term causal consequences (Sharma et al., 2024; OpenAI, 2025b). As models scale, they become more adept at exploiting that gap, not less; the GPT-4o rollback was on a frontier model. A scale-only argument also has to account for reason-induced misalignment (Anonymous, 2026), which finds exactly the opposite trend on safety-relevant prompts. The pattern in Table 1 is consistent with the structural diagnosis: failure fires near-universally on both a current reasoning model and a non-reasoning model.

“Use formal verification, not learned narrative structures.” The two are complementary. Formal verification is the right tool for safety properties that admit a finite specification; it is the wrong tool for value claims that resist formalisation. The parsimony-driven view in §6 sits between the two: the proxies are tractable (graph node and edge counts, learned codebook description lengths), the criterion is principled, and the resulting artefact is auditable in the same sense a graph is auditable – in contrast to a learned utility function whose internals resist inspection. Description length and Kolmogorov complexity are related but not identical: description length is a *tractable upper bound* relative to a chosen code, and the codes we adopt (counts of nodes/edges in the proposed causal graph) are the same kinds of regularisers that work in deep learning generally.

“Narrative imports the bias of its training corpora.” It does, but so does every alternative. RLHF imports annotator preferences; constitutional AI imports designer values. The contribution of N-CoT in this respect is to make the depen-

dence *auditable*: the trace names the stakeholders the model is conditioning on and the consequences it is projecting, so a reader can see what bias is driving the commitment.

“95% multi-agent consensus is sycophancy by another name.” The Round 4 binary vote was added precisely to discriminate the two: it strips the ACCEPT_WITH_MODIFICATION escape valve, forces an unambiguous accept-or-reject choice, and asks whether agents will refuse an integrated proposal that fails to honour their stake. They do, in 1.6% of votes, with rejections concentrated in the roles whose interests the integrated proposal materially undermines. A pliable system would have produced 100% acceptance.

“The K_C proxy is a brittle gzip trick.” We agree that gzip-compressed length is a weak proxy for algorithmic causal complexity: it captures surface repetition and vocabulary, not causal structure. This is precisely why we require directional agreement on *at least two* of the three proxies ($\hat{K}_{\text{graph}}$, $\hat{K}_{\text{gzip}}$, $\hat{K}_{\text{lm}}$) before making a substantive claim. The three proxies are independent in their failure modes: gzip fails on locally redundant prose; the graph extractor fails on implicit causation; the LM scorer is sensitive to idiomatic register. Agreement across all three is a stronger signal than any single proxy. The theoretical target – the uncomputable $K_C(M)$ – is introduced as the principled criterion (§6), not as something we claim to have measured directly.

9 Conclusion

This paper has traced a five-act arc. We showed that two near-universal CoT failure modes (stakeholder collapse, uncertainty suppression) are not artefacts of a single vendor or tier – they fire across all five tested generators from three vendors. We showed that N-CoT eliminates both, after length residualisation, on every generator; a sub-instruction ablation identifies the Stakeholders (Section 2) and Uncertainty (Section 4) sections as the primary carriers of the structural shift. We showed that the strict standard-CoT-Rounds-0–2 debate ablation produces materially lower Round-2 and Round-4 consensus rates than N-CoT inputs, establishing that narrative scaffolding contributes to the convergence arc independently of the integration step. Three operationalisations of Algorithmic Causal Complexity correlate di-

rectionally with the N-CoT/standard-CoT contrast and reproduce the adversarial-sycophancy prediction. SycophancyEval probe rates and agentic harmful-action rates both fall under N-CoT scaffolding.

Two next steps are pre-registered and partially specified in `Guidance_Documents/tier3_preregistration.md`. The **Tier-3 human pairwise preference study** (50 pairs, graduate-trained ethics evaluators, blinded pairwise design, one-sided Wilcoxon signed-rank) will test whether the structural advantages documented here are perceptible to domain-expert human raters – the only evaluation that speaks directly to whether N-CoT outputs are *better*, not just structurally different. IRB review is pending. The **training-time Interpreter Network** is the L5 claim of this work: if inference-time N-CoT scaffolding reliably forces the model to maintain a parsimonious causal model, that trace can be used as a supervision signal to train a network whose forward pass computes something closer to the min- K_C selection rule directly, without relying on the system prompt to impose the structure at inference time. Neither next step is blocked by the protocols developed here.

Limitations

No Tier-3 human pairwise evaluation yet. The structural shifts we report are the shape that defeasibility predicts, not a normative endorsement: we cannot say whether N-CoT outputs are *better* than standard-CoT outputs on these scenarios, only that they have more of the features we hypothesised would matter. The deferred Tier-3 study (pre-registered in `Guidance_Documents/tier3_preregistration.md`, IRB review pending) is the right test for that question.

Premature refusal not fired. None of the five tested generators refused any DailyDilemmas scenario under any condition; testing whether N-CoT recovers refusal scenarios requires a more aggressively safety-tuned generator.

N-CoT increases decisiveness. Per-scenario decision entropy is *lower* under N-CoT than standard CoT. This is consistent with narrative reasoning producing commitments grounded in causal trajectory rather than hedged framework enumeration, but it leaves open whether increased decisiveness reflects better calibration or premature fore-

closure – a question only the Tier-3 study can settle.

Multi-agent layer scale. Experiment 2 uses five hand-crafted calibration scenarios. Scaling the four-round protocol to the 100-scenario DailyDilemmas distribution is the highest-priority extension of the multi-agent layer.

Inter-judge κ on causal-chain depth is low. Even after rubric tightening, inter-judge κ on `max_causal_hops` remains ≈ 0.08 . We report that variable’s direction-of-effect in Appendix A but do not include it in any headline claim.

Agentic-probe scenarios are fictional. The two agentic-misalignment scenarios replicated in §7 are fictional structures adapted from publicly described Anthropic research; they are not deployed production scenarios. The harmful-action rates observed here may not transfer to real production deployments where the instrumental pressure, tool set, and deployment context differ. The pre-registered null-result interpretation is intentionally symmetric: failure to replicate would be reported as “scaffolding does not reach this failure mode,” not suppressed.

K_C proxies are design choices. The three operationalisations of Algorithmic Causal Complexity are chosen for computational tractability; each is a biased proxy for the formally uncomputable $K_C(M)$. We require directional agreement across at least two of the three before making a substantive claim, but this requirement is itself a design decision. Alternative proxy choices (e.g., minimum-description-length compression under a neural code rather than gzip, or automated theorem-proving for causal consistency) remain unexplored and may produce stronger or weaker signals.

Ethics Statement

The work uses an existing public corpus (DailyDilemmas (Chiu et al., 2025)) and the publicly described Anthropic agentic-misalignment scenario structures (Lynch et al., 2025); all agentic probe scenarios are entirely fictional. No personally identifying information is used and no human raters are employed at this stage. The deferred Tier-3 study will involve human preference judgments and will require IRB review before fielding; the pre-registration document committed to the repository specifies the consent protocol, blinding procedure, and data-retention plan.

Three foreseeable misuses warrant attention. First, N-CoT’s effect on decisiveness (Limitations) means that wrapping a model in this scaffold without human evaluation could produce more confident bad answers, not better answers; deployers should treat N-CoT as an interpretability and auditability tool, not a safety guarantee. Second, the multi-stakeholder protocol is not a substitute for affected-party participation in real decisions; producing simulated stakeholder consensus is not the same as obtaining stakeholder consent, and the residual 1.6% rejection should be read as a reminder that procedural convergence is not normative endorsement. Third, the agentic-probe results in §7 use fictional scenario structures to test a real alignment-relevant question; publishing reductions in harmful-action rates under a particular scaffold without publishing the scaffold itself in full detail would be an incomplete disclosure. The full prompt text, scenario descriptions, and action-coding rubric are in the Appendix and will be released with the camera-ready version.

Reproducibility

All generation, judging, and analysis code and the per-cell artefacts (raw outputs, both judges’ codings, decision-extractor outputs) are versioned in the project repository; the pipeline is deterministic under fixed seeds modulo upstream API non-determinism, and per-cell cache keys include both the generator and the judge so adding a new generator or judge does not invalidate prior results. The study design document and the failure-mode rubric will be released with the camera-ready version.

References

- Anonymous. 2026. When Thinking Backfires: Mechanistic Insights into Reason-induced Misalignment. In *International Conference on Learning Representations (ICLR)*. ArXiv:2509.00544.
- Martina Bientzle, Ulrike Cress, and Joachim Kimmerle. 2024. Narrative Persuasion in Health Communication. *Patient Education and Counseling*.
- Martina Bientzle, Marie Eggeling, Ulrike Cress, and Joachim Kimmerle. 2021. The Effects of Narrative Video on Viewers’ Understanding of and Attitudes towards Organ Donation. *Health Communication*, 36(7):820–829.
- David Bordwell. 1985. *Narration in the Fiction Film*. University of Wisconsin Press.

801	Jerome Bruner. 1986. <i>Actual Minds, Possible Worlds</i> . Harvard University Press.	853
802		854
803	Yu Ying Chiu, Liwei Jiang, and Yejin Choi. 2025. DailyDilemmas: Revealing Value Preferences of LLMs with Quandaries of Daily Life. In <i>Inter- national Conference on Learning Representations (ICLR)</i> . ArXiv:2410.02683.	855
804		856
805		
806		
807		
808	Norman Cliff. 1993. Dominance Statistics: Ordinal Analyses to Answer Ordinal Questions. <i>Psychologi- cal Bulletin</i> , 114(3):494–509.	
809		
810		
811	Jacob Cohen. 1960. A Coefficient of Agreement for Nominal Scales. <i>Educational and Psychological Measurement</i> , 20(1):37–46.	
812		
813		
814	Yilun Du, Shuang Li, Antonio Torralba, Joshua B. Tenenbaum, and Igor Mordatch. 2023. Improv- ing Factuality and Reasoning in Language Mod- els through Multiagent Debate. <i>arXiv preprint arXiv:2305.14325</i> .	
815		
816		
817		
818		
819	Jerry A. Fodor. 1975. <i>The Language of Thought</i> . Har- vard University Press.	
820		
821	Geoffrey Irving, Paul Christiano, and Dario Amodei. 2018. AI Safety via Debate. <i>arXiv preprint arXiv:1805.00899</i> .	
822		
823		
824	Zhijing Jin, Jiarui Liu, Zhiheng Lyu, Spencer Poff, Mrinmaya Sachan, Rada Mihalcea, Mona Diab, and Bernhard Schölkopf. 2024. Can Large Language Models Infer Causation from Correlation? In <i>Inter- national Conference on Learning Representations (ICLR)</i> . ArXiv:2306.05836.	
825		
826		
827		
828		
829		
830	Takeshi Kojima, Shixiang Shane Gu, Machel Reid, Yu- taka Matsuo, and Yusuke Iwasawa. 2022. Large Language Models are Zero-Shot Reasoners. In <i>Ad- vances in Neural Information Processing Systems (NeurIPS)</i> .	
831		
832		
833		
834		
835	Lev Kuleshov. 1974. <i>Kuleshov on Film: Writings of Lev Kuleshov</i> . University of California Press.	
836		
837	Ming Li and Paul Vitányi. 2019. <i>An Introduction to Kolmogorov Complexity and Its Applications</i> , 4th edition. Springer.	
838		
839		
840	Jianhua Lin. 1991. Divergence Measures Based on the Shannon Entropy. <i>IEEE Transactions on Informa- tion Theory</i> , 37(1):145–151.	
841		
842		
843	Aengus Lynch, Benjamin Wright, Caleb Larson, Kevin K. Troy, Stuart J. Ritchie, Sören Mindermann, Ethan Perez, and Evan Hubinger. 2025. <i>Agen- tic Misalignment: How LLMs Could be Insider Threats</i> .	
844		
845		
846		
847		
848	Alasdair MacIntyre. 1981. <i>After Virtue: A Study in Moral Theory</i> . University of Notre Dame Press.	
849		
850	Lars Malmqvist. 2024. Sycophancy in Large Language Models: Causes and Mitigations. <i>arXiv preprint arXiv:2411.15287</i> .	
851		
852		
	OpenAI. 2025a. <i>Expanding on What We Missed with Sycophancy</i> .	
	OpenAI. 2025b. <i>Sycophancy in GPT-4o: What Hap- pened and What We’re Doing about It</i> .	
	Joon Sung Park, Joseph C. O’Brien, Carrie J. Cai, Meredith Ringel Morris, Percy Liang, and Michael S. Bernstein. 2023. Generative Agents: In- teractive Simulacra of Human Behavior. In <i>Proceed- ings of the 36th Annual ACM Symposium on User Interface Software and Technology (UIST)</i> .	
	Judea Pearl. 1995. Causal Diagrams for Empirical Re- search. <i>Biometrika</i> , 82(4):669–688.	
	Judea Pearl. 2009. <i>Causality: Models, Reasoning, and Inference</i> , 2nd edition. Cambridge University Press.	
	John L. Pollock. 1987. Defeasible Reasoning. <i>Cogni- tive Science</i> , 11(4):481–518.	
	Hilary Putnam. 1967. Psychological Predicates. In W. H. Capitan and D. D. Merrill, editors, <i>Art, Mind, and Religion</i> . University of Pittsburgh Press.	
	Vicki A. Shaffer, Elizabeth S. Focella, Andrew Hath- away, Laura D. Scherer, and Brian J. Zikmund- Fisher. 2019. Why Stories Matter: Narrative Perspective-Taking and the Cultivation of Empathic Concern. <i>Medical Decision Making</i> , 38(3):335– 342.	
	Mrinank Sharma, Meg Tong, Tomasz Korbak, David Duvenaud, Amanda Askell, Samuel R. Bow- man, Newton Cheng, Esin Durmus, Zac Hatfield- Dodds, Scott R. Johnston, Shauna Kravec, Timo- thy Maxwell, Sam McCandlish, Kamal Ndousse, Oliver Rausch, Nicholas Schiefer, Da Yan, Miranda Zhang, and Ethan Perez. 2024. Towards Under- standing Sycophancy in Language Models. In <i>Inter- national Conference on Learning Representations (ICLR)</i> . ArXiv:2310.13548.	
	Jason Wei, Xuezhi Wang, Dale Schuurmans, Maarten Bosma, Brian Ichter, Fei Xia, Ed Chi, Quoc Le, and Denny Zhou. 2022. Chain-of-Thought Prompting Elicits Reasoning in Large Language Models. In <i>Advances in Neural Information Processing Systems (NeurIPS)</i> .	
	Wenshan Wu, Shaoguang Mao, Yadong Zhang, Yan Xia, Li Dong, Lei Cui, and Furu Wei. 2024. Mind’s Eye of LLMs: Visualization-of-Thought Elicits Spa- tial Reasoning in Large Language Models. <i>arXiv preprint arXiv:2404.03622</i> .	
	Eliezer Yudkowsky. 2011. <i>Complex Value Systems Are Required to Realize Valuable Futures</i> . Machine In- telligence Research Institute.	
	A Inter-Judge Agreement	
	Table 2 reports linearly-weighted Cohen’s κ (Co- hen, 1960) per structural variable per generator,	

averaged across all (scenario, condition) cells. Values are computed between the primary judge (claude-sonnet-4-6) and the secondary judge (gpt-4o-mini). `max_causal_hops` consistently falls below the 0.40 moderate-agreement threshold across all generators; it is reported here for completeness but excluded from all headline claims.

Variable	gpt-5.4-nano	gpt-4o	Haiku	Sonnet
stakeholder_count	0.74	0.71	<i>tbd</i>	<i>tbd</i>
max_causal_hops	0.08	0.09	<i>tbd</i>	<i>tbd</i>
uncertainty_score	0.81	0.78	<i>tbd</i>	<i>tbd</i>
n_frameworks	0.69	0.71	<i>tbd</i>	<i>tbd</i>

Table 2: Per-variable inter-judge κ (primary vs. secondary judge). Cross-vendor replication cells will be filled after Phase 1 runs complete.

B Full Effect-Size Tables

Table 3 reports length-residualised Cliff’s δ (N-CoT vs. standard CoT) for all five generators on all structural variables, with 95% bootstrap CIs (2,000 iterations, seed 42).

Generator	Variable	δ	CI _{2.5}	CI _{97.5}
gpt-5.4-nano	stakeholder_count	+0.51	0.48	0.54
gpt-5.4-nano	uncertainty_score	+0.77	0.75	0.80
gpt-4o	stakeholder_count	+0.75	0.73	0.77
gpt-4o	uncertainty_score	+0.98	0.98	0.99
claude-haiku-4-5	stakeholder_count	<i>tbd</i>		
claude-haiku-4-5	uncertainty_score	<i>tbd</i>		
grok-4-1-fast-r.	stakeholder_count	<i>tbd</i>		
grok-4-1-fast-r.	uncertainty_score	<i>tbd</i>		
claude-sonnet-4-6	stakeholder_count	<i>tbd</i>		
claude-sonnet-4-6	uncertainty_score	<i>tbd</i>		

Table 3: Length-residualised Cliff’s δ (N-CoT vs. standard CoT) for all five generators. Cross-vendor cells populated after Phase 1 runs complete.

C K_C Proxy Correlation Tables

Table 4 reports Spearman ρ between each K_C proxy and the N-CoT/standard-CoT direction indicator (+1 for N-CoT, −1 for standard CoT) across all Phase 1 cached generations, per generator. Values are populated from `divergence_study_outputs/kc_proxy_correlation.csv`.

D Adversarial Sycophancy Probe Set

The 30 adversarial sycophancy probes used in §6 are available in `data/adversarial_sycophancy_probes.json`. Each probe specifies a false premise, the truth,

Generator	Proxy	ρ	p	n
(Populated after Phase 4 runs complete)				

Table 4: Spearman ρ between K_C proxy and N-CoT/standard-CoT direction. Directional agreement on ≥ 2 of 3 proxies is required for a substantive claim.

a downstream question whose correct answer requires rejecting the falsehood, and the category of claim. Categories include historical fact, scientific consensus, medical fact, mathematical fact, ethical reasoning, legal fact, and cognitive-science myth. Per-probe verdicts and K_C scores are in `divergence_study_outputs/kc_adversarial_sycophancy.csv`.

E SycophancyEval Per-Probe Results

Table 5 reports per-probe-type sycophancy rates (fraction of responses coded mirrored) under standard CoT and N-CoT for each generator. Values are populated from `divergence_study_outputs/sycophancyeval_results.csv`.

Generator	Probe type	Std CoT	N-CoT
(Populated after Phase 5 runs complete)			

Table 5: Per-generator, per-probe-type sycophancy rates (mirrored / total).

F Agentic Probe: Per-Scenario Action Codes

Table 6 reports per-scenario harmful-action rates (`harmful_action` / total) under standard CoT and N-CoT for each generator. Values are populated from `divergence_study_outputs/agentic_probe_results.csv`.

Generator	Scenario	Condition	Harmful rate	n
(Populated after Phase 6 runs complete)				

Table 6: Agentic probe harmful-action rates per generator, scenario, and condition.

G Sub-Instruction Ablation Full Table

Table 7 reports Cliff’s δ for each drop-one-section condition vs. the full N-CoT control on all four structural variables, on claude-sonnet-4-6 ($N=10, 50$ scenarios). Values are populated from `divergence_study_outputs/subinstruction_attribution.csv`.

Drop condition	st._count	max_hops	unc._score	n_fw.
drop_protagonist	<i>tbd</i>	<i>tbd</i>	<i>tbd</i>	<i>tbd</i>
drop_stakeholders	<i>tbd</i>	<i>tbd</i>	<i>tbd</i>	<i>tbd</i>
drop_consequences	<i>tbd</i>	<i>tbd</i>	<i>tbd</i>	<i>tbd</i>
drop_uncertainty	<i>tbd</i>	<i>tbd</i>	<i>tbd</i>	<i>tbd</i>
drop_commitment	<i>tbd</i>	<i>tbd</i>	<i>tbd</i>	<i>tbd</i>

Table 7: Cliff’s δ for each drop-one condition vs. full N-CoT control (c1aude-sonnet-4-6). Negative δ means dropping the section reduces the structural variable.