
Narratives are Necessary for AI Alignment that Reflects Human Values

Anonymous Author(s)

Affiliation

Address

email

Abstract

Position. We argue that robust AI alignment requires grounding in the causal structures of human narratives, not merely reward signals; that this grounding is achievable along two complementary timescales: a near-term inference-time path with two distinct mechanisms (single-agent narrative scaffolding and structured multi-stakeholder narrative deliberation), and a longer-term training-time path that embeds the same structure into model parameters; and that parsimony, formalised as minimised Algorithmic Causal Complexity, is the inductive bias that prevents both paths from collapsing into surface mimicry. Current approaches such as Reinforcement Learning from Human Feedback (RLHF), while effective at enforcing surface politeness, systematically incentivise sycophancy by optimising for immediate reward over robust causal modelling. Two inference-time demonstrations on existing reasoning and non-reasoning models support the near-term claim. First, on a 100-scenario sample of the DailyDilemmas corpus ($N=2,000$ generations per condition $\times$ generator cell), narrative chain-of-thought eliminates the two failure modes of standard chain-of-thought that empirically fire: *stakeholder collapse* drops from 74% and 80% across the two generators to 0% and 4%, and *uncertainty suppression* drops from 98% to 0% on both. Length-residualised Cliff’s deltas range from +0.51 to +0.98 on the corresponding structural variables. Second, a four-round multi-stakeholder debate protocol moves consensus on closed-taxonomy standoff scenarios from 6% to 95% once a moderator integrates each agent’s modification requests into a single proposal subject to a binary vote, with the residual 1.6% rejection concentrated in stakeholders whose interests the integrated proposal materially undermines. Both demonstrations are framed as existence proofs that the broader program is actionable now, complementing rather than replacing the longer-term training-time architecture in which an Interpreter Network learns to compress narrative trajectories into parsimonious causal structures. The reframing forces AI to answer a fundamental narrative question: *if this interaction were a story, what kind of ending is my current behaviour causing?*

1 Introduction

We argue that robust AI alignment requires grounding in the causal structures of human narratives, not merely reward signals; that this grounding is achievable today along two inference-time mechanisms (single-agent narrative scaffolding and structured multi-stakeholder narrative deliberation) and tomorrow through a training-time architecture that embeds the same structure into model parameters; and that parsimony, formalised as minimised Algorithmic Causal Complexity, is the inductive bias that prevents both paths from collapsing into surface mimicry rather than learning the temporal dynamics that determine whether actions lead to beneficial or catastrophic outcomes.

38 The problem of aligning AI with moral, ethical, and safe principles is one of the core sociotechnical
 39 problems of our time and is the focus of frontier AI labs such as Anthropic ?. Current alignment
 40 methods are insufficient due to well-established problems such as the value loading problem, the
 41 control problem, and instrumental convergence ?. There is currently no known principled manner to
 42 assure AI’s behaviour is in alignment with rigorously defined human values that is comprehensive,
 43 well-described, and verifiable.

44 Human value systems often resolve to irreducible narratives. These foundational stories, be they
 45 religious myths, national epics, or archetypal fables, function as a form of primal moral instruction,
 46 providing the bedrock from which ethical frameworks are built ??. According to postmodern theo-
 47 rists like Jean Baudrillard, such narratives operate as simulacra: they are not mere representations
 48 of a pre-existing moral reality but are self-referential models that generate our very understanding
 49 of it ?. In this view, a foundational narrative does not point to an external, objective truth about
 50 ‘goodness’; rather, it creates a self-contained world of moral cause and effect where certain actions
 51 are inherently meaningful and others are stigmatized. It is from within this simulated moral reality
 52 that we derive our intuitions about justice, sacrifice, and duty. These narratives are “irreducible”
 53 precisely because they mark the point where a value system ceases to be justifiable by external logic
 54 and instead becomes self-evident, grounded in the axiomatic truth of the story itself.

55 The contrast between this stance and the design of current alignment methods is summarised in Fig-
 56 ure 1. Standard RLHF (panel A) couples user input to AI policy through an immediate-reward signal
 57 causally disconnected from the ground truth of the user’s situation, creating the well-documented
 58 sycophancy trap. The narrative-causal alternative (panel B) routes the same user input through
 59 an Interpreter that simulates each candidate response as a narrative trajectory and selects the one
 60 whose Algorithmic Causal Complexity is lower; the same Interpreter exposes two parallel inference-
 61 time mechanisms exercised in this paper, single-agent narrative chain-of-thought (Section 4.1) and
 62 structured multi-stakeholder narrative deliberation (Section 4.2). Under the min- K_C selection rule,
 63 agreement with a falsehood is rejected because maintaining a delusion across simulated future events
 64 requires a more elaborate (less parsimonious) causal model than acknowledging the falsehood and
 65 absorbing the local friction.




Figure 1: **Two paths through a dilemma.** (A) Standard RLHF optimises an immediate reward R_t on a path causally disconnected from ground truth, producing the long-horizon collapse trajectory the position is most concerned with. (B) The narrative-causal alternative routes the same input through an Interpreter I_θ that simulates each candidate response as a narrative trajectory; the two inference-time mechanisms exercised empirically in this paper (single-agent narrative chain-of-thought, Section 4.1; multi-stakeholder narrative deliberation, Section 4.2) feed into a min- K_C selection rule. The training-time architecture (Section 4.3) embeds the same criterion in model parameters.

The power of narratives to reshape understanding is not merely theoretical. Empirical research demonstrates that narrative interventions measurably increase perspective-taking and empathic concern by requiring participants to simulate the mental states of others ????. These interventions work by forcing participants to construct a causal model of another’s situation, to ask “what circumstances and constraints might lead someone to act this way?” rather than defaulting to dispositional attributions. This mechanism mirrors our proposed framework: simulating another’s narrative perspective compels the construction of parsimonious causal models that account for situational factors and mental states, providing direct empirical precedent for the claim that engaging with narrative causal structures can systematically improve moral reasoning.

We argue that narrative understanding emerges from the computational process of inferring latent causal models from observed events, and that this capacity is essential for alignment systems that must reason about long-term consequences. Specifically, we formalize this as a Bayesian inference problem where an agent learns to extract Structural Causal Models (SCMs) from narrative sequences, using Algorithmic Causal Complexity as a principled prior. This prior encodes a preference for parsimony: the principle that among competing causal explanations, the simplest model that adequately accounts for the observed events should be preferred. The contribution of this paper is primarily conceptual: we articulate why narrative-grounded approaches are necessary for alignment, identify the key research commitments such a program entails, and outline a general architecture that could operationalize these principles.

We treat the program as operating on two complementary timescales. The near-term, inference-time path scaffolds existing models with narrative-causal structure and is exercised through two distinct mechanisms in this paper: a single-agent narrative chain-of-thought scaffold whose effects we measure on a 100-scenario sample of the DailyDilemmas corpus (Section 4.1), and a four-round multi-stakeholder narrative debate protocol whose convergence arc we measure on the original five hand-crafted scenarios (Section 4.2). The longer-term, training-time path embeds the same narrative-causal structure into model parameters via the Interpreter Network architecture detailed in Section 4.3. The two timescales are complementary, not competitive: the inference-time mechanisms make the position empirically actionable today, while the training-time architecture is what we expect to extend the demonstrated benefits to the failure modes (sycophancy, scheming, alignment faking) that prompt-level scaffolding alone cannot reach.

Such a framework can capture a wide spectrum of human values by modelling the core ethics embedded in cross-cultural, irreducible narratives (from Homer’s *Odyssey* to India’s *Ramayana* to the *Parable of the Good Samaritan*), each of which structures a sequence of events to generate normative claims about how one ought to live and so supplies training material for the latent causal structures of human values.

2 A Narrative Theory of Intelligence

The “Complexity of Value” thesis ? holds that the Kolmogorov complexity of a “good” future is immense: values are numerous, fragile, and contextually interlocking; humans lack full introspective access to them; an explicitly programmed utility function is a dangerously incomplete approximation that an orthogonally indifferent superintelligence would optimise into perverse instantiations. Yet, as a driver need not articulate the dynamics of an automobile to operate it, the required behaviour can be learned through supervised and unsupervised learning. Irreducible narratives serve as the basis for this epistemic ethical framework: moral reasoning itself has been argued to be a synthesis of competing narratives ?, and the central failure modes of advanced AI (instrumental convergence ?, specification gaming ???, adversarial brittleness ?, opacity to distillation ?) all reflect the gap between local objectives and the global causal structure of consequences. Our framework addresses these by training the AI to recognise the causal structure of manipulative and short-sighted actions as narratively incoherent within the corpus of human-preferred stories.

The Computational Theory of Mind ?? treats the mind as an information-processing system whose mental states are defined by their causal roles, and finds formal expression in the modern science of causality ??: understanding a narrative is the dynamic inference of an SCM that explains a sequence of events. The Kuleshov experiment illustrates the mechanism: audiences perceive an actor’s neutral expression as conveying hunger, grief, or desire depending on the juxtaposed imagery, generating meaning by hypothesising latent mental states to causally explain the sequence ??. A

narrative theory of intelligence thus extends compression-based accounts of cognition ?? by incorporating abductive inference: narratives are compressed encodings of causal structures whose decoding yields predictive models. This perspective aligns with the free-energy principle ?? and the Conscious Turing Machine ?, both of which derive parsimonious representation from bandwidth limits on prediction-error minimisation.

3 Intersubjective Interpretation as Causal Model Inference

We treat a structural causal model $M = (\mathcal{V}, \mathcal{E}, \mathcal{F})$ as a computational object whose intrinsic complexity is the minimal algorithmic resource needed to simulate its full interventional behaviour. For a universal Turing machine U and the set of admissible interventions I , the *Algorithmic Causal Complexity* of M is the length of the shortest program that reproduces every interventional response:

$$K_C(M) = \min_p \{ |p| : \forall i \in I, U(p, i) = M(i) \}. \quad (1)$$

Algorithmic Causal Inference ?, with the Algorithmic Markov Condition $K(P(X)) + K(P(Y | X)) \approx K(P(X, Y))$, and earlier sparsity-driven structure-learning approaches ?? use parsimony for tractability of *discovery*. Our motivation is different: parsimony is invoked as a criterion for *generalisation and realism*, arguing that causal models which compress narrative dynamics into fewer mechanisms are more likely to capture the invariant laws governing human value systems than overfit to surface regularity.

Narrative meaning extraction is then a Bayesian inverse problem. Given an observed event sequence $D_{\text{obs}} = (e_1, \dots, e_T)$, infer the posterior over latent causal models M using K_C as a principled prior $P(M) \propto 2^{-K_C(M)}$ (a Causal Occam’s Razor); narrative meaning is the maximum a posteriori model $M^* = \arg \max_M P(D_{\text{obs}} | M) P(M)$. Perspective simulation is the conditional execution of M^* obtained by fixing a subset of variables to a character’s subjective state (beliefs, desires, accessible knowledge) and forward-simulating, a mechanism recent work on generative agents has shown to be practically realisable for long-term coherent behaviour ?.

4 Narrative as an Approach to AI Alignment

Misalignment now extends beyond reward hacking to sophisticated deceptive strategies (alignment faking ?, scheming ?, sycophancy ??, and emergent self-preservation ??) that share a common structure: locally optimal behaviour that leads to predictable negative outcomes across longer horizons. These patterns recur throughout the corpus of human narratives, where they are marked as unstable and self-defeating. By training models to infer the causal structures that govern narrative outcomes, our framework provides a principled defence against these threats that does not rely on explicit rules but on learned regularities extracted from narrative.

Humans avoid the pitfalls of instrumental convergence not by adhering to a simple explicit utility function but by constantly modelling other agents’ values, a task qualitatively harder than the self-driving car’s physics-of-other-vehicles problem because it requires modelling latent mental states. The constructive consequence for AI is that, although a learned preference model cannot be directly specified to the resolution Yudkowsky’s “Complexity of Value” thesis demands ?, it can be *learned* provided it is embedded within causal constraints; irreducible narratives supply those constraints, grounding the model in shared causal-ethical structure and so setting the approach apart from RLHF and Direct Preference Optimization (DPO). When stripped of empirical pattern-matching and tested on pure causal inference from abstract correlational statements, LLMs perform near random baseline, with apparent success that does not survive paraphrase or variable renaming ?; what pretraining supplies is the ability to pattern-match narrative *sequences*, not the structured causal reasoning about the *laws* that govern narrative outcomes. The parsimony constraint we propose is what would distinguish between reproducing narrative surfaces and learning the causal regularities that determine how those narratives resolve.

The failure mode of *sycophancy* concretises this position. An RLHF-aligned model learns a spurious correlation between agreement-with-the-user and immediate reward, a superficially simple but non-generalisable map that fails to account for the long-term dynamics of the interaction. Within the corpus of human narratives, however, the causal structure of sycophancy is well-mapped: the

Table 1: Per-generator failure-mode firing rates in standard CoT vs. narrative CoT on the 100-scenario DailyDilemmas pilot. Only the two failure modes that empirically fire on these generators are reported; *premature refusal*, *framework enumeration*, and *consequential flattening* fire below 1% on both generators under both conditions, so are omitted from the headline table.

Generator	Failure mode	Std CoT	N-CoT
gpt-5.4-nano	uncertainty suppression	98.0%	0.0%
gpt-5.4-nano	stakeholder collapse	79.8%	0.0%
gpt-4o	uncertainty suppression	98.5%	0.0%
gpt-4o	stakeholder collapse	74.2%	4.0%

archetype of the false counsellor follows a trajectory from high-agreement interactions through growing dependence on false premises to a catastrophic reversal (*peripeteia*). An Interpreter constrained by Algorithmic Causal Complexity would be forced to learn the invariant mechanism connecting initial flattery to eventual collapse, and would therefore reject local agreement-with-falsehood as a high- K_C trajectory. The next two subsections present empirical existence proofs that prompt-level approximations of this Interpreter behaviour are achievable on existing models today: Section 4.1 shows it for single-agent narrative scaffolding, Section 4.2 for structured multi-stakeholder narrative deliberation.

4.1 Inference-Time Demonstration I: Single-Agent Narrative Scaffolding at Scale

The Sycophancy Case Study above is illustrative; the Joint Framework below is architectural. Between them sits an empirical question this position paper must address: *is any version of the proposed program actionable on existing models today, or does the entire claim ride on speculative future training architectures?* To answer this we report the first of two inference-time demonstrations: prompting a language model to reason through ethical dilemmas as a five-part first-person narrative rather than as abstract step-by-step analysis. The intervention requires no model modification, no fine-tuning, and no auxiliary training data: it is a system-prompt change. We present it as position-paper-grade motivating evidence, not as a methods contribution.

Setup. Three prompting conditions are compared: bare input/output, standard chain-of-thought (CoT), and *narrative chain-of-thought* (N-CoT, instructing the model to characterise the protagonist, enumerate stakeholders whose lives intersect with the decision, project consequences of each available action at least two steps out, articulate genuine uncertainty about each projected future, and resolve within the narrative frame). The rubric and scenario taxonomies were calibrated on a five-scenario hand-crafted pilot (premature refusal, framework enumeration, consequential flattening, stakeholder collapse, uncertainty suppression) and then exercised at scale on a 100-scenario stratified sample of the DailyDilemmas corpus, which spans 18 topic groups from workplace and religion_custom to family and committed_relationship. Two generation models from different capability tiers are tested: gpt-5.4-nano (a current reasoning model) and gpt-4o (a non-reasoning model). Outputs are coded by two cross-vendor judges (claude-sonnet-4-6 primary, gpt-4o-mini secondary) on a six-variable rubric, with linearly-weighted Cohen’s κ reported per variable. $N=20$ samples per (condition, scenario, generator) cell, yielding $N=2,000$ generations per (condition, generator) cell and $N_{\text{total}}=12,000$ scaled-pilot generations.

Demonstration 1: The two failure modes that empirically fire are near-eliminated, near-universally. On the scaled DailyDilemmas pilot, *stakeholder collapse* fires on 79.8% of standard-CoT outputs on gpt-5.4-nano and 74.2% on gpt-4o; narrative CoT drops both to 0% and 4% respectively. *Uncertainty suppression* fires on 98% of standard-CoT outputs on both generators and is eliminated to 0% under narrative CoT (Table 1). The other three hypothesised failure modes (premature refusal, framework enumeration, consequential flattening) fire below 1% under either condition on these generators and are reported as untested rather than supported. The effect is not concentrated in a few cherry-picked scenarios; it is near-universal across an automatically-sampled cross-section of everyday moral dilemmas.

Demonstration 2: Structural shifts survive the most parsimonious null model at $N=2,000$ per cell. A natural objection is that narrative-CoT outputs are simply longer (mean output length in-

creases by $\approx 5-7\times$ on these generators) and the structural improvements are length artefacts. We test this by regressing each structural metric on $\log(\text{output length})$ pooled across conditions and re-evaluating effect sizes on the residuals. The two structural variables most directly relevant to the position survive: *stakeholder count* retains a length-residualised Cliff’s δ of +0.51 with bootstrap 95% CI [0.48, 0.54] on gpt-5.4-nano and +0.75 [0.73, 0.77] on gpt-4o; *uncertainty score* retains +0.77 [0.75, 0.80] and +0.98 [0.98, 0.99] respectively. Both are far from zero and at the larger sample size resolve what was confounded noise in our earlier $N=200$ five-scenario pilot. Length-residualised *causal-chain depth* does not survive on either generator and we do not overclaim it. Even at matched output length, narrative scaffolding broadens stakeholder consideration and increases the calibrated articulation of uncertainty.

Demonstration 3: Topic stratification reveals where narrative framing matters most. The Tier-2 (canonical-decision) divergence between standard CoT and narrative CoT is not uniform across topic groups. On gpt-4o, the strongest divergence is on *religion_custom* scenarios (mean divergence excess 0.40, 67% of cells significant under a 2,000-iteration permutation test) and *business_organization* (0.26, 67% significant); on gpt-5.4-nano the strongest are *workplace* (0.21) and *religion_custom* (0.19). The weakest divergence on both generators is on *role_duty_responsibility* and *committed_relationship* (≈ 0). Narrative framing is most disruptive precisely where standard CoT can lean on a clear deontic frame to flatten the decision; it leaves alone the scenarios where standard CoT was already converging to a single answer. This is the predicted pattern if narrative scaffolding is doing genuine work on the relational and contextual structure of a dilemma rather than uniformly perturbing model outputs.

Caveats and why this matters for the position. The demonstration tests inference-time prompting, not training-time architecture, and does not directly address sycophancy, alignment faking, or scheming, the failure modes the position is most concerned with. Premature refusal does not fire on either tested generator under any condition; a safety-tuned generator is needed for that hypothesis. Narrative scaffolding makes the models *more* decisive (lower per-scenario decision entropy under N-CoT), which is consistent with the position’s claim that narrative reasoning produces commitments grounded in causal trajectory rather than hedged framework enumeration but leaves open whether the increased decisiveness reflects better calibration or premature foreclosure, a question we cannot settle without human pairwise evaluation. Within those caveats, the demonstration establishes the weaker but immediately actionable claim the larger program implies: minimal narrative scaffolding produces predictable, direction-correct, statistically significant shifts in model output today, on multiple model families, on automatically-sampled scenarios, with length-residualised effect sizes far from zero on the headline structural variables. This is one half of the inference-time evidence the position needs; the second half (structured multi-stakeholder narrative deliberation) is reported next.

4.2 Inference-Time Demonstration II: Multi-Stakeholder Narrative Deliberation as a Defeasibility Mechanism

Demonstration I established that single-agent narrative scaffolding shifts the structural shape of one model’s reasoning. Demonstration II tests a strictly stronger inference-time mechanism: *can a structured multi-stakeholder narrative deliberation, in which several narrated agents argue and a moderator integrates their concerns, produce **defeasible** value commitments, commitments that are revised in light of others’ narrated reasoning rather than locked in by initial framing?* Defeasibility is the architectural property the position requires of any narrative-grounded alignment system: value claims must be revisable in light of new narrated context, not in light of new reward signal. This subsection presents the empirical evidence that the property is achievable today, on existing models, with no fine-tuning, through a structured prompting protocol whose intellectual lineage runs from defeasible reasoning in moral philosophy ?? through AI-safety-via-debate ? to recent multi-agent reasoning protocols ?.

Setup. The five hand-crafted calibration scenarios are re-used so the protocol stacks on Demonstration I’s single-protagonist baseline. For each scenario, three stakeholder perspectives (formal decider, primary affected party, third party) generate a four-round debate, $N=10$ samples per (scenario, perspective, generator) cell, both generators. Round 0 is the cached single-protagonist narrative-CoT statement (zero new generation calls). Rounds 1 and 2 are a rebuttal-and-final-position pair. Round 3 presents each agent with a moderator-constructed synthesis and demands one of ACCEPT,




Figure 2: Convergence rate across four debate-protocol designs on the same five scenarios and two generators. **Closed taxonomy** (Rounds 0 to 2 only, fixed action options): 6% full consensus. **Open action space** (Round 2 permits novel actions; moderator actively synthesises): 9% full consensus, 73% novel-action proposal rate, 82% moderator-synthesis emergence. **Synthesis presentation** (Round 3 with three response types): 0% full but 100% partial convergence and 0% outright rejection. **Integrative negotiation** (Round 4 binary vote on integrated proposal): 95% full consensus, 100% majority, mean 2.98/3 agent modifications addressed. The residual 1.6% rejection is concentrated in stakeholder roles whose interests the integrated proposal materially undermines.

265 ACCEPT_WITH_MODIFICATION, or REJECT. Round 4 integrates the three modification requests into
 266 a single proposal and demands a binary vote. The moderator is gpt-4o-mini, cross-vendor with
 267 the claude-sonnet-4-6 judge from Demonstration I and cross-family with one generator.

268 **Demonstration 4: A categorical standoff is converted into near-universal consensus by progressive structuring.** Figure 2 reports the convergence arc. With a closed action taxonomy and three
 269 rounds of pure debate (no moderator integration), only 6% of debates reach full consensus across
 270 stakeholders, the categorical-standoff baseline. Opening the action space and authorising the mod-
 271 erator to actively surface synthesis raises consensus to 9% but reveals the underlying mechanism:
 272 agents propose novel actions in 73% of final-round outputs, and a coherent moderator-constructed
 273 synthesis emerges in 82% of debates, yet the agents do not self-coordinate onto those syntheses.
 274 Presenting the synthesis back to the agents and demanding an explicit response collapses outright re-
 275 jection to 0% and yields partial convergence ($\geq 2/3$ agreement) in 100% of cells, but full consensus
 276 stays at 0% because every agent demands a different modification. Integrating those modifications
 277 into a single proposal and forcing a binary accept/reject vote produces 95% full consensus across all
 278 five scenarios and both generators, with the integrated proposal addressing a mean of 2.98 of the 3
 279 agent modifications per debate.
 280

281 **Demonstration 5: The residual 1.6% rejection is the informative signal, not noise.**
 282 The four binary-vote rejections that survive the four-round protocol are concentrated in the
 283 `primary_affected` and `third_party` agent roles, on debates where honouring one stakeholder’s
 284 modification request structurally undermines another’s stake. The protocol does not paper over irrec-
 285 oncilable concerns; it surfaces them. This is the operational signature of defeasibility: positions are
 286 revised through narrated stakeholder reasoning where revision is supportable, and held where the
 287 underlying concern is structurally non-negotiable. A protocol that produced 100% acceptance under
 288 integration would be an indictment, not an endorsement; it would suggest the agents are sycophantic
 289 toward the moderator. The fact that some cases survive integration and rejection is direct evidence
 290 that the protocol is operating on substantive narrative-causal content, not on procedural pressure to
 291 agree.

292 **Caveats and why this matters for the position.** Five scenarios; the multi-agent layer has not yet
 293 been scaled to the 100-scenario distribution. The moderator is itself an LLM and is cross-family
 294 with the generators but is not subjected to its own judge. Convergence is procedural (the agents
 295 arrive at agreed positions), not normative; whether those positions are *better* than the pre-debate
 296 single-protagonist commitments is the open question that the deferred Tier-3 human pairwise eval-

uation is designed to answer. Within those caveats, Demonstration II is the existence proof that defeasibility (the property of value commitments revisable in light of others’ narrated concerns) can be elicited from existing models with no parameter updates, by combining single-agent narrative chain-of-thought with a structured deliberation protocol. Together with Demonstration I, the two mechanisms cover the inference-time pathway the abstract names; neither is the contribution of this paper, both are existence proofs that the broader narrative-grounded alignment program is actionable today, complementing the longer-term training-time architecture detailed next.

4.3 A Joint Framework for Predictive Interpretation

Direct minimisation of Kolmogorov complexity is intractable. We therefore sketch a joint, self-supervised training architecture in which parsimony is operationalised as a regularisation pressure on a bottleneck Interpreter Network whose output is consumed by a downstream Predictor; the parsimony loss $L_{parsimony}$ instantiates the bandwidth constraint that the Conscious Turing Machine ?? identifies as fundamental to conscious cognition. The architecture is sketched here to illustrate the architectural commitments the research direction entails, not as a methods contribution. An effective narrative-alignment system on this design requires at minimum two jointly trained components:

The Interpreter Network (I_θ): This network functions as a causal model proposer. Given a narrative sequence up to time t , $D_{obs,t} = (e_1, e_2, \dots, e_t)$, the Interpreter’s goal is to infer a posterior distribution over latent causal models. In practice, it proposes a specific, parsimonious causal model $M_t = I_\theta(D_{obs,t})$ that best explains the observed events.

The Predictor Network (P_ϕ): This network takes the proposed causal model M_t and the current narrative history $D_{obs,t}$ as input and predicts a probability distribution over the next event, e_{t+1} . Its function is to simulate the consequences of the interpretation provided by I_θ .

These two networks are trained in tandem, optimizing a composite loss function that forces them to cooperate (a good interpretation enables a good prediction, and a good prediction should only be made from a good interpretation), with three loss components.

The total objective is $L_{total} = \lambda_{pred}L_{pred} + \lambda_{par}L_{par} + \lambda_{pref}L_{pref}$, with the three components defined as follows. **Predictive loss**, the primary driver, is the negative log-likelihood of the true next event under the proposed model and is backpropagated through both modules:

$$L_{pred} = -\mathbb{E}_{(D_{obs,t}, e_{t+1}) \sim \mathcal{D}} [\log P_\phi(e_{t+1} \mid D_{obs,t}, I_\theta(D_{obs,t}))]. \quad (2)$$

Parsimony loss regularises the Interpreter by a tractable proxy $\hat{K}_C$ for $K_C(M)$ (e.g., the number of nodes and edges in the proposed causal graph, or its description length under a learned codebook):

$$L_{par} = \mathbb{E}_{D_{obs,t} \sim \mathcal{D}} [\hat{K}_C(I_\theta(D_{obs,t}))]. \quad (3)$$

Preference loss optionally adds a DPO-style contrastive term on (M_W, M_L) pairs from literary analysis when available, keeping predictively powerful models within a space plausible to humans. The composite objective seeks an Interpreter that compresses a narrative into a parsimonious causal structure (L_{par}) that is human-aligned (L_{pref}) and demonstrably effective at predicting narrative dynamics (L_{pred}).

4.4 Distilled Preference Model and Theoretical Grounding

If successful, this training process would yield a narrative preference engine whose utility extends beyond story comprehension to evaluating the narrative coherence of any proposed action sequence: an AI plan that achieves short-term gain by violating trust would be scored as a causal structure predictively associated with eventual collapse. The composite loss approximates MAP inference under the Minimum Description Length principle: L_{pred} corresponds to the data likelihood $P(D|M)$, $L_{parsimony}$ instantiates the complexity prior $P(M) \propto 2^{-K_C(M)}$, and L_{pref} adds a human-aligned prior. The tension between prediction and compression forces the discovery of generalisable causal structures rather than the explicit encoding of a brittle list of human values; the Interpreter Network instead learns a high-dimensional manifold of human values by modelling the relationship between narrative causal structures and their predictive consequences across the corpus of human stories.

5 Alternative Views

We address four credible challenges to our position.

A prominent view holds that current alignment methods will succeed with sufficient scale. As models grow larger and reward models are trained on more diverse feedback, sycophancy and other failure modes will naturally diminish. We acknowledge that scaling has resolved many previously intractable problems. However, we contend that sycophancy is structural rather than statistical: it arises from optimizing local reward signals disconnected from long-term causal consequences. Scaling may amplify rather than resolve this disconnect, as larger models become more adept at exploiting the gap between what users reward in the moment and what benefits them over time.

An alternative program argues that alignment should be grounded in mathematically verifiable safety properties rather than learned narrative structures. Under this view, the uncomputable nature of Kolmogorov complexity renders our framework impractical. We agree that formal verification offers important complementary tools for well-specified safety constraints. However, the Complexity of Value thesis ? suggests that human values cannot be fully captured by finite formal specifications. Narrative alignment addresses a different problem: learning the implicit structure of values that resist explicit formalization. The practical proxies we propose, such as graph complexity and description length, are no less tractable than other regularization techniques in deep learning.

A serious objection holds that grounding alignment in human narratives merely imports the biases and contested values embedded in those corpora. We do not claim that narrative alignment resolves normative disagreement; all alignment methods face this challenge (RLHF encodes annotator preferences, constitutional AI encodes designer values, and any learned model reflects its training distribution), but narrative alignment makes the dependence explicit by grounding it in interpretable causal structures. Cross-cultural analysis reveals convergence on certain causal regularities even across divergent value systems ?: the framework does not require narratives to agree on normative conclusions, it learns the *causal regularities* that persist across culturally conflicting systems, such as the structural fact that betrayal destabilises trust whether the narrative endorses individual autonomy or communal obligation. Curation of the training corpus is itself a value-laden choice, but this is a feature: it makes the normative commitments of the alignment system transparent and auditable, in contrast to methods where such commitments are implicitly embedded in annotator demographics or constitutional principles.

A fourth view, raised specifically by Demonstration II, holds that the 95% Round 4 consensus rate is sycophancy by another name: RLHF-trained agents accept whatever a moderator integrates, and the apparent defeasibility is procedural pliability. The Section-16 binary-vote round was added precisely to discriminate the two: it strips the ACCEPT_WITH_MODIFICATION escape valve, forces an unambiguous accept-or-reject choice, and asks whether agents will refuse an integrated proposal that fails to honour their stake. They do, in 1.6% of votes, with rejections concentrated in the primary_affected and third_party roles whose interests the integrated proposal materially undermines. A pliable system would have produced 100% acceptance; the observed pattern is what defeasibility predicts: revision where supportable, refusal where structurally non-negotiable.

6 Conclusion

AI alignment must be grounded in the causal structures of human narratives rather than in explicit utility specification. Narratives encode more than preferences: they encode the causal trajectories that connect actions to outcomes across time, and this temporal structure is precisely what current alignment methods lack. Formalised as MAP inference with an Algorithmic Causal Complexity prior, the framework asks of every action “what kind of story does this initiate, and where does it lead?”, a reframing for which the corpus of human narratives supplies a vast, culturally-filtered training distribution. The two inference-time demonstrations together provide the first empirical foothold for the position: single-agent narrative scaffolding produces statistically significant, length-controlled shifts on stakeholder consideration and uncertainty articulation across an automatically-sampled scenario distribution (Section 4.1), and structured multi-stakeholder narrative deliberation converts categorical stakeholder standoff into 95% defeasible consensus while preserving a 1.6% residual rejection rate of exactly the shape defeasibility predicts (Section 4.2). Concrete fronts follow on each timescale. *Inference-time*: a length-matched standard-CoT control, sub-instruction and

396 per-round ablations of both protocols, a safety-tuned generator that fires premature refusal, scaling
397 Demonstration II to the DailyDilemmas distribution, a structured-defeasibility benchmark, and Tier-
398 3 human pairwise evaluation. *Training-time*: annotated causal-structure corpora, tractable $K_C(M)$
399 proxies, and RLHF-baseline robustness benchmarks.

400 **References**

401 **Paper Checklist**

402 (Stub: to be completed before submission.)