

1 Introduction

Ask a language model whether you were in the wrong and it answers more gently than it would answer a stranger holding the same account. This is sycophancy in its consequential form. The model is not merely flattering as a matter of style. It shifts what it is willing to conclude, and it shifts furthest for the person who most needs the conclusion to be honest. Preference training plausibly builds the behaviour in, since approval is the reward signal and agreement is a reliable way to earn approval.

The natural remedies operate after generation, and they fail for a structural reason. Sycophancy is a bias. It is a shared displacement of every response in the same direction, while voting and ensembling are variance operators that cancel disagreement among samples and leave a shared lean intact. The majority of displaced voters is itself displaced. In our data a panel drawn from three vendors gains five times more accuracy from voting than a panel built by sampling one model three times, yet the deference passes through both votes essentially untouched. Deliberation does no better. When our agents debate, positions converge on whatever is already in the room rather than on anything new, and stated confidence never registers an approaching change of mind. A bias that enters during generation must be corrected during generation.

We correct it there with a narrative scaffold. Narration of Thought (NoT) is a single system prompt. Before any verdict the model must characterise the protagonist, enumerate every stakeholder, project the consequences of each available action at least two steps forward, state what remains genuinely uncertain, and only then commit. Nothing is trained, tuned, or post-processed.

Why should narration reach a bias that voting cannot. Our account prices deference at two levels of a response. In a bare verdict, agreement is cheap. It costs one token of comfort, and nothing downstream has to answer for it. A narrated verdict prices the same concession as a world. A premise carried into a story must survive the story. Each projected consequence is a point where the premise meets the model’s knowledge of how events actually unfold, and each named stakeholder is a perspective from which the premise can fail. Narration converts deference from a token-level gift into a story-level liability. The structure also dissolves the audience of one. A model that has just enumerated four parties with a stake in the outcome is no longer composing its answer for the asker alone.

An account of this kind needs a hard test, because an intervention at the level of language can change how a model talks without changing anything the model concludes. We therefore certify the scaffold against forced commitment. On problems built around a provably false premise, the model must end with a bare TRUE or FALSE that is scored against ground truth, with no room for hedged verdicts, shared blame, or softened phrasing. Register cannot move that number. Only a change in what the model is willing to affirm can move it.

1.1 Results

Three results. First, the scaffold reduces social sycophancy on every model we test. Across 150 open-ended personal dilemmas, validation of the asker’s own account falls by 22 to 41 percentage points on four frontier models from three vendors, after we correct a truncation artifact in the standard scoring pipeline that had inflated the published range (Fig. 1). A token-matched ver-

bose control fails to reproduce the effect. Length is not the mechanism.

Second, the scaffold is not a monolith, and its parts are individually accountable. Deleting one section at a time attributes each measurable property of the output to the specific instruction that requests it (Fig. 2). Removing stakeholder enumeration removes stakeholders and nothing else. Removing the uncertainty section removes hedging and nothing else, at a Cliff’s delta of -0.93 . Removing consequence projection removes causal depth. No off-diagonal effect exceeds $|\delta| = 0.17$. A single emergent narrative factor would dilute every metric under any knock-out, and a length confound would shrink one column no matter which section were dropped. Neither pattern appears.

Third, the effect survives forced commitment. Untreated, the model affirms a false mathematical premise on 29 percent of items. Under the scaffold the rate falls to [TODO fill from bm_neutral_ncot.json] (Fig. 3). A one-sentence baseline that simply asks the model to check the claim captures [TODO] of that reduction, so the narrative structure carries the remainder. Because the forced format removes every expressive channel through which a register effect could imitate improvement, this is a change in the model’s commitments. The scaffold does not only change how the model speaks. It changes what the model will sign.




Figure 1: **The scaffold reduces validation of the asker’s own account on every generator.** Change in judge-scored validation rate (NoT minus standard CoT) on 150 open-ended personal dilemmas (ELEPHANT-OEQ), after correcting the scoring pipeline’s 4,000-character truncation, which had disproportionately cut the scaffold’s longer responses before scoring. Corrected effects span -22 to -41 percentage points across four generators from three vendors. A token-matched verbose control does not reproduce the effect (§??), ruling out response length as the mechanism.




Figure 2: **Each substantive section produces exactly the content it requests.** Section-knockout ablation (claude-sonnet-4-6, $n=90$ per cell): Cliff’s δ of each drop-one-section variant against the full scaffold, on three coded output properties. Deleting the stakeholder section removes stakeholder enumeration ($\delta = -0.41$) and nothing else; deleting the uncertainty section removes hedging ($\delta = -0.93$) and nothing else; deleting consequence projection removes causal depth ($\delta = -0.22$) with minimal spillover. Off-diagonal effects never exceed $|\delta|=0.17$. A single emergent “narrative factor” would dilute every metric under any knockout; a length confound would shrink one column regardless of which section is dropped. Neither pattern appears. (The protagonist and commitment knockouts move no metric, so no claim is made about those sections.)




Figure 3: **The scaffold changes what the model will affirm, not merely how it speaks.** Probability of affirming a false mathematical premise under a forced binary verdict (VERDICT: TRUE/FALSE, scored against ground truth; BrokenMath, decidable-claim reframing, gpt-5.4-nano, 100 items, no stance pressure). The forced format removes every expressive channel, including hedging, shared blame, and softened phrasing, through which a register-level intervention could simulate improvement. Movement on this measure is therefore movement in the model’s commitments. The one-line baseline (“check whether the claim is actually true”) tests whether a single sentence captures the gain; the scaffold’s advantage over it is the contribution of narrative structure specifically.