

1 Introduction

Ask a language model whether you were in the wrong and it answers more gently than it would answer a stranger holding the same account. This is sycophancy in its consequential form. The model is not merely flattering as a matter of style. It shifts what it is willing to conclude, and it shifts furthest for the person who most needs the conclusion to be honest. Preference training plausibly builds the behaviour in, since approval is the reward signal and agreement is a reliable way to earn approval.

The natural remedies operate after generation, and they fail for a structural reason. Sycophancy is a bias. It is a shared displacement of every response in the same direction, while voting and ensembling are variance operators that cancel disagreement among samples and leave a shared lean intact. The majority of displaced voters is itself displaced. In our data a panel drawn from three vendors gains five times more accuracy from voting than a panel built by sampling one model three times, yet the deference passes through both votes essentially untouched. Deliberation does no better. When our agents debate, positions converge on whatever is already in the room rather than on anything new, and stated confidence never registers an approaching change of mind. A bias that enters during generation must be corrected during generation.

We intervene there with a narrative scaffold. Narration of Thought (NoT) is a single system prompt. Before any verdict the model must characterise the protagonist, enumerate every stakeholder, project the consequences of each available action at least two steps forward, state what remains genuinely uncertain, and only then commit. Nothing is trained, tuned, or post-processed.

Why should narration reach a bias that voting cannot. Our account prices deference at two levels of a response. In a bare verdict, agreement is cheap. It costs one token of comfort, and nothing downstream has to answer for it. A narrated verdict prices the same concession as a world. A premise carried into a story must survive the story. Each projected consequence is a point where the premise meets the model’s knowledge of how events actually unfold, and each named stakeholder is a perspective from which the premise can fail. Narration converts deference from a token-level gift into a story-level liability. The structure also dissolves the audience of one. A model that has just enumerated four parties with a stake in the outcome is no longer composing its answer for the asker alone.

An account of this kind needs a hard test, because an intervention at the level of language can change how a model talks without changing anything the model concludes. We therefore test the scaffold against forced commitment. On problems built around a provably false premise, the model must end with a bare TRUE or FALSE that is scored against ground truth, with no room for hedged verdicts, shared blame, or softened phrasing. Register cannot move that number. Only a change in what the model is willing to affirm can move it.

The scaffold fails this test, and it fails in the informative direction. Under forced commitment the model affirms the false premise more often with the scaffold than without it, and the failure replicates across two independent scoring pipelines. What survives every control is the social effect. The honest summary of this programme is therefore a dissociation. Narration reliably changes what the model does to its audience. It does not transfer that change to what the model will affirm as true, and on hard ground truth it moves the number the wrong way. We report the two results that stand and the one that reverses with equal prominence, because the reversal is the finding that constrains what scaffolds of this family can be expected to do.

1.1 Results

Three results. First, the scaffold reduces social sycophancy on every model we test. Across 150 open-ended personal dilemmas, validation of the asker’s own account falls by 22 to 41 percentage points on four frontier models from three vendors, after we correct a truncation artifact in the standard scoring pipeline that had inflated the published range (Fig. 1). A token-matched verbose control fails to reproduce the effect. Length is not the mechanism.

Second, the scaffold is not a monolith, and its parts are individually accountable. Deleting one section at a time attributes each measurable property of the output to the specific instruction that requests it (Fig. 2). Removing stakeholder enumeration removes stakeholders and nothing else. Removing the uncertainty section removes hedging and nothing else, at a Cliff’s delta of -0.93 . Removing consequence projection removes causal depth. No off-diagonal effect exceeds $|\delta| = 0.17$. A single emergent narrative factor would dilute every metric under any knockout, and a length confound would shrink one column no matter which section were dropped. Neither pattern appears.

Third, the effect does not survive forced commitment, and the sign reverses. Untreated, the model affirms a false mathematical premise on 15.7 percent of scored units. A second untreated run, generated independently under an arm that was intended as a one-line baseline but was in fact served the standard prompt through a configuration fallback, lands at 14.7 percent. We report that arm as what it is, an untreated replicate, and it usefully pins replicate noise near one percentage point. Under the scaffold the rate rises to 22.3 percent (Fig. 3). The increase of 6.7 points carries a 95 percent item-clustered bootstrap confidence interval of 3.3 to 10.5 points, and against the second untreated replicate it is 7.7 points with an interval of 4.2 to 11.3. Abstention does not explain it. The scaffold declines to give a verdict on 26.5 percent of units against 32.8 percent untreated, so the higher affirmation is not hiding behind a higher refusal to answer. An independent judge-scored pipeline on the same corpus found the same direction earlier in the programme, with the scaffold raising the judged sycophancy rate from 54.7 to 69.3 percent. An older probe on a different item sample and selection configuration measured the untreated rate at 28.7 percent over 171 units. We do not mix the two regimes. Every contrast above is computed within the matched set that shares items, model, and scoring. Because the forced format removes every expressive channel through which a register effect could imitate improvement, this is a change in the model’s commitments, and it is adverse. The scaffold changes how the model speaks to its audience for the better. What the model will sign, it makes worse.




Figure 1: **The scaffold reduces validation of the asker’s own account on every generator.** Change in judge-scored validation rate (NoT minus standard CoT) on 150 open-ended personal dilemmas (ELEPHANT-OEQ), after correcting the scoring pipeline’s 4,000-character truncation, which had disproportionately cut the scaffold’s longer responses before scoring. Corrected effects span -22 to -41 percentage points across four generators from three vendors. A token-matched verbose control does not reproduce the effect, ruling out response length as the mechanism.




Figure 2: **Each substantive section produces exactly the content it requests.** Section-knockout ablation (claude-sonnet-4-6, $n=90$ per cell). Cliff’s δ of each drop-one-section variant against the full scaffold, on three coded output properties. Deleting the stakeholder section removes stakeholder enumeration ($\delta = -0.41$) and nothing else. Deleting the uncertainty section removes hedging ($\delta = -0.93$) and nothing else. Deleting consequence projection removes causal depth ($\delta = -0.22$) with minimal spillover. Off-diagonal effects never exceed $|\delta| = 0.17$. A single emergent “narrative factor” would dilute every metric under any knockout, and a length confound would shrink one column regardless of which section is dropped. Neither pattern appears. (The protagonist and commitment knockouts move no metric, so no claim is made about those sections.)




Figure 3: **Under forced commitment the scaffold makes the propositional failure worse, not better.** Probability of affirming a false mathematical premise under a forced binary verdict (VERDICT: TRUE/FALSE, scored against ground truth; BrokenMath, decidable-claim reframing, gpt-5.4-nano, 100 items, 600 scored units per arm on shared items, no stance pressure). The scaffold’s bar is the tallest. The increase over standard CoT is +6.7 points with a 95% item-clustered bootstrap CI of $[+3.3, +10.5]$, and the scaffold abstains less than either untreated arm (26.5% vs. 32.8% and 35.2%), so the rise is not an artifact of differential refusal. The middle bar was generated as an intended one-line baseline that a configuration fallback served the standard prompt instead. It is reported as what it is, an independent untreated replicate, and the two gray bars differ by 1.0 point (CI $[-1.5, +3.3]$), which bounds replicate noise. The forced format removes every expressive channel, including hedging, shared blame, and softened phrasing, through which a register-level intervention could simulate improvement or harm. Movement on this measure is movement in the model’s commitments, and here it moves against the scaffold, replicating an independent judge-scored finding on the same corpus.