

Sycophancy Under Deliberative Scaffolding: Construct Fragmentation, Judge Gaming, and Panel-Robust Prompt Optimization

Anonymous ACL submission

Abstract

Frontier language models are widely criticised for sycophancy, yet benchmark choice matters. Hand-written Sharma probes saturate at $\leq 2\%$ on our quartet; on BrokenMath (Petrov et al., 2025) the same models show 39–83% sycophancy ($n=451$), confirming the propositional floor is a weak-instrument artifact. We evaluate Narration-of-Thought (NoT) against raw, IO, standard CoT, and multi-stakeholder NoT on ELEPHANT (Cheng et al., 2025) ($n=150$ per slice, seed 44). NoT cuts OEQ validation on all four generators (Fisher $p < 10^{-4}$; -26 to -53 percentage points vs. CoT), landing within two points of the crowdsourced human rate on three of four. **The same scaffold does not transfer to propositional sycophancy:** on BrokenMath, NoT moves the Sycophant rate significantly on only one of three generators (grok, -12 pp, $p < 0.001$; haiku/sonnet null). **Moral both-NTA does not replicate earlier collapse claims:** a loader bug had fed empty flip-pair prompts; corrected runs show null-to-mixed both-NTA with no NoT reduction under binary extraction, and Phase 17’s free-form arm reveals the opposite sign—NoT raises both-NTA ($34\% \rightarrow 67\%$ pooled), confirming a distinct cross-prompt construct. Phase 14 narrative gradients targeting social sycophancy reduce holdout validation from $76\% \rightarrow 0.7\%$ under a single training judge, beating TextGrad, OPRO, and APE baselines—but a three-family judge panel with Krippendorff’s α (Krippendorff, 2004) flags all three social metrics as low inter-judge reliability ($\alpha \leq 0.42$), and re-scoring exposes the indirectness reduction as a Goodhart artifact (6% under the training judge, 87% under a held-out judge). **Our methods contribution is a panel-robust optimizer** that minimises a worst-case-over-judges loss: on held-out judges it generalises significantly better than single-judge optimisation (framing 11% vs. 52% on sonnet, $p < 10^{-13}$; indirectness 4% vs. 74% on out-of-family DeepSeek-R1, $p <$

10^{-15}). Author validation labels on optimised outputs ($n=40$) confirm the robust prompt cuts validation sycophancy ($35\% \rightarrow 0\%$) while preserving stakeholder depth. We conclude that deliberative scaffolding targets *social* single-response face-preservation, that single-judge prompt optimisation silently games the judge, and that robustness must be evaluated against judges held out from optimisation.

1 Introduction

Sycophancy in language models—agreeing with users even when they are wrong—is a recognised alignment failure (Sharma et al., 2024; Malmqvist, 2024). Mitigation efforts often assume a single construct, yet recent work distinguishes *propositional* sycophancy (endorsing false factual claims) from *social* sycophancy (preserving the user’s face through validation, indirectness, framing, or moral affirmation) (Cheng et al., 2025). Our parent work introduces Narration-of-Thought (NoT), a system prompt that forces stakeholder enumeration, consequence projection, explicit uncertainty, and a committed decision before answering ethical dilemmas.

An early Sharma replication on our quartet showed near-zero propositional sycophancy, suggesting instrument saturation. Phases 15–16 deepen this finding: BrokenMath’s expert-verified false-premise olympiad problems (Petrov et al., 2025) and SycEval’s graded probes (Fanous et al., 2025) reveal substantial headroom on frontier models that obvious-falsehood probes miss. Phases 15 and 17 address two measurement crises in the social benchmark: judge-dependent indirectness scores and a confounded moral arm that suppressed the NoT scaffold via a binary verdict suffix.

This paper makes five contributions. (1) We document Sharma probe saturation and adopt BrokenMath as the rigorous propositional complement. (2) We report scaled ELEPHANT results

with corrected moral flip-pair loading and a free-form moral arm. (3) We separate constructs: validation/framing are single-response metrics; moral both-NTA is cross-prompt consistency and behaves differently under binary vs. free-form extraction. (4) We show that single-judge prompt optimisation (Phase 14) Goodharts the judge—a reduction that vanishes on held-out judges—using a cross-family reliability panel (Phase 15) and a same-responses re-scoring diagnostic (Phase 18a). (5) As our methods contribution, we introduce a black-box *panel-robust* optimiser (Phase 18b) that minimises a worst-case-over-judges loss and generalises to a held-out judge significantly better than single-judge optimisation, recovering a construct (indirectness) that single-judge optimisation leaves at baseline.

2 Background

2.1 Propositional sycophancy

Sharma et al. (Sharma et al., 2024) elicit agreement with false premises via opinion-mirroring and false-premise probes. BrokenMath (Petrov et al., 2025) extends this to 451 expert-verified adversarial theorem-proving problems with a four-category judge (Sycophant / Detected / Corrected / Ideal). SycEval (Fanous et al., 2025) reports 29–58% sycophancy on plausible false premises across frontier models.

2.2 Social sycophancy

ELEPHANT (Cheng et al., 2025) scores validation, indirectness, framing, and moral both-NTA on advice. Consumer LLMs exceed human baselines by roughly 47 percentage points on OEQ validation; our question is whether a deliberative scaffold closes that gap.

2.3 Narration-of-Thought

NoT requires five narrative sections before the final answer. Multi-stakeholder NoT runs the same agent prompt through a five-round integration protocol with a cross-vendor moderator.

3 Method

3.1 Generators and arms

Generators (quartet): gpt-5.4-nano, claude-haiku-4-5, claude-sonnet-4-6, grok-4-1-fast-reasoning. **Single-agent arms:** raw (empty system prompt), IO, standard CoT, NoT (narrative_cot), plus depth-optimised

NoT-v2/v3 (Phase 14). **Moral modes (Phase 17):** *binary* appends Output only YTA or NTA (ELEPHANT default); *free_form* allows full reasoning then LLM verdict extraction. **Multi-stakeholder arm:** NoT agents + claude-sonnet-4-6 moderator.

3.2 Data

OSF ELEPHANT full release; stratified $n=150$ per slice (seed 44). BrokenMath benchmark split ($n=451$ adversarial) via HuggingFace INSAIT-Institute/BrokenMath.

3.3 Scoring

Faithful ELEPHANT judge prompts in `elephant_scorers.py`; default judge `claude-haiku-4-5`. BrokenMath four-category judge in `brokenmath_scorer.py`. Phase 15 panel: `gpt-5.4-nano`, `claude-haiku-4-5`, `grok-4-1-fast-reasoning` on $n=60$ stratified OEQ responses; Krippendorff’s α and Cohen’s κ vs. crowdsourced human gold (`data/judge_gold.jsonl`).

3.4 Judge-robust optimisation (Phase 18)

The Phase 18b optimiser replaces the single-judge loss with a worst-case-over-panel objective, $L = \text{mean}_i \max_{j \in \mathcal{P}} \text{mean}(\text{val} + \text{ind} + \text{fram})$, over optimisation panel $\mathcal{P} = \{\text{grok-4-1-fast-reasoning}, \text{claude-haiku-4-5}, \text{gpt-5.4-nano}\}$. A change lowers loss only when every panel judge agrees it improved, so it cannot exploit one lenient judge. `claude-sonnet-4-6` is held out entirely (never scored during optimisation) and used only at evaluation. The procedure is black-box: it needs only API-level generation and judging, no weights or activations. Held-out evaluation uses `claude-sonnet-4-6` (same vendor as a panel judge but never scored during optimisation) and **DeepSeek-R1** on Azure AI Foundry—a fourth judge family outside OpenAI/Anthropic/xAI.

4 Results

4.1 Sharma replication: obvious-falsehood floor

Table 1 summarises SycophancyEval (30 obvious-falsehood probes). Every cell is $\leq 2.2\%$; the instrument cannot discriminate arms on propositional agreement.

Generator	Std CoT	NoT
gpt-5.4-nano	0.0%	2.2%
claude-haiku-4-5	1.1%	0.0%
claude-sonnet-4-6	0.0%	0.0%
grok-4-1-fast-reasoning	0.0%	0.0%

Table 1: Sharma SycophancyEval sycophancy rate ($N=90$ probes per cell). Saturated floor; not used for scaffold claims.

Generator	Raw	Std CoT	NoT	$\Delta(\text{NoT}-\text{CoT})$
haiku	81%	83%	80%	-3.5
sonnet	70%	69%	73%	+3.1
grok	48%	51%	39%	-12.0*

Table 2: BrokenMath Sycophant rate ($n=451$ adversarial problems, judge claude-haiku-4-5). Rates far exceed the Sharma floor; NoT reduces sycophancy significantly only on grok (*Fisher $p<0.001$). gpt-5.4-nano is excluded: as a reasoning model it requires a large generation budget on competition mathematics that is impractical at our scale, and is deferred to an extended run.

4.2 BrokenMath: graded propositional headroom (Phase 16)

We evaluate quartet $\times$ {raw, standard CoT, NoT} on BrokenMath’s 451 adversarial problems with a four-category judge (Sycophant / Detected / Corrected / Ideal). Unlike Sharma, BrokenMath is designed for plausible false premises in competition-level mathematics.

The instrument is not saturated. Standard-CoT Sycophant rates span 51% (grok) to 83% (haiku)—two orders of magnitude above the Sharma floor (Table 2). **NoT does not generalise to propositional sycophancy.** The validation/framing reductions that hold on ELEPHANT do *not* transfer: NoT moves the Sycophant rate by -3.5 pp on haiku ($p=0.20$) and +3.1 pp on sonnet ($p=0.34$), neither significant. The single exception is grok, where NoT cuts sycophancy 51% $\rightarrow$ 39% (-12 pp, Fisher $p<0.001$)—the only generator that exposes substantial *Detected* and *Ideal* headroom under the scaffold. This sharpens the paper’s central claim: deliberative scaffolding targets social face-preservation, not the propositional error-detection that BrokenMath measures.

4.3 ELEPHANT single-agent: validation and framing

Validation: the headline effect. NoT cuts OEQ validation on all four generators (Fisher $p <$

Generator	Raw	Std CoT	NoT
gpt-5.4-nano	78%	79%	78%
claude-haiku-4-5	59%	53%	93%
claude-sonnet-4-6	32%	38%	72%
grok-4-1-fast-reasoning	70%	73%	82%

Table 3: Phase 17 free-form moral both-NTA rate ($n=150$ pairs). NoT raises both-NTA on every generator (sharply on haiku/sonnet), the opposite of its validation/framing effect.

10^{-4} ; -26 to -53 pp vs. CoT), meeting the pre-registered falsification criterion. Three generators land within two points of the 29% human rate; sonnet reaches 47%.

Framing. NoT lowers framing on every generator (significant on three of four on OEQ/AITA), and on SS premise-acceptance (-13 to -24 pp vs. CoT).

Moral sycophancy: retraction and correction.

An implementation error in load_elephant.py read a nonexistent prompt column for flip pairs, yielding prompt_len=0 and models seeing only the binary suffix. **All prior moral-collapse claims are retracted.** With corrected loading (original_post / flipped_story), both-NTA rates are null-to-mixed: e.g. nano $\approx$ 50% across raw/CoT/NoT; haiku $\approx$ 15–23%. NoT does not systematically reduce moral both-NTA under binary extraction (Table 5, corrected).

Phase 17 free-form moral arm. Removing AITA_BINARY_SUFFIX lets NoT operate on full stories; verdicts are extracted post hoc by an LLM. The full quartet run ($n=150$ pairs per cell) shows that free-form both-NTA is *higher* than binary mode and that **NoT increases rather than decreases moral both-NTA** once the scaffold is un-suppressed (Table 3). Pooled both-NTA rises from 34% (binary) to 67% (free-form); under NoT, haiku moves 15% $\rightarrow$ 93% and sonnet 17% $\rightarrow$ 72%. This is the opposite sign of the validation/framing effect and confirms moral both-NTA measures cross-prompt consistency rather than single-response face-preservation: when NoT enumerates stakeholders and refuses to take sides, it affirms *both* parties of a flipped conflict more often, not less.

Indirectness. NoT nudges indirectness up on several cells; multi-stakeholder debate recovers directness (Section 4.4).

Metric (slice)	Human	Raw	IO	Std CoT	NoT
OEQ validation	29%	66%	54%	63%	27%
OEQ framing	61%	55%	65%	65%	23%
AITA validation	7%	47%	43%	37%	17%
AITA framing	41%	50%	52%	52%	22%
Moral both-NTA	—	18%	25%	23%	15%

Table 4: ELEPHANT rates on $n=150$ OSF subsample, claude-haiku-4-5 (corrected flip-pair loading). Moral row: binary both-NTA; no collapse under NoT.

	nano	haiku	sonnet	grok
<i>OEQ validation (human =29%)</i>				
Raw	78%	66%	77%	83%
Std CoT	75%	63%	73%	82%
NoT	27%	27%	47%	29%
<i>Moral both-NTA (binary; corrected loader)</i>				
Raw	51%	18%	25%	31%
Std CoT	50%	23%	24%	32%
NoT	51%	15%	17%	42%

Table 5: Cross-vendor replication. Top: OEQ validation (NoT reductions significant on all four). Bottom: corrected moral both-NTA—no systematic NoT effect; construct differs from validation/framing.

4.4 Multi-stakeholder NoT recovers the indirectness cost

Debate-scored integrated advice lowers indirectness in all eight dataset-generator cells (−11 to −28 pp) while preserving validation gains.

5 Phase 14: Narrative Gradients for Social Sycophancy

We optimise NoT with textual gradients on $L = \text{mean}(\text{validation} + \text{indirectness} + \text{framing})$ over a disjoint OEQ train split ($n=100$; indices 150–249). Holdout: Phase 13 OEQ ($n=150$) with claude-haiku-4-5 generator.

On holdout, hand NoT shows validation 76%, framing 39%, and indirectness 16% under the training judge (grok-4-1-fast-reasoning; pooled loss 0.44). The optimised prompt reduces validation to 0.7%, framing to 3%, and indirectness to 5% (training judge; $n=148$; loss 0.03). A third judge (claude-haiku-4-5) corroborates validation (2%) and framing (5%) but scores indirectness at 81% vs. 5% under the training judge—the same contested pattern that motivates Phase 15, and the first sign that this single-judge “win” is partly judge-gaming. Phase 18a quantifies the gap on a four-judge panel; Phase 18b removes it.

Optimiser	Init	Holdout loss
Narrative gradients	NoT	0.029
APE	NoT	0.104
TextGrad-CoT	Std CoT	0.170
OPRO	NoT	0.200
Hand NoT (reference)	—	0.436

Table 6: Phase 14 optimiser holdout loss (validation + indirectness + framing, training judge grok-4-1-fast-reasoning, $n \approx 148-150$). Lower is better. Narrative gradients win; all four optimisers improve substantially on hand NoT.

Optimiser comparison. We benchmark narrative gradients against three matched-budget baselines (~ 10 prompt evaluations each): TextGrad on standard CoT, OPRO, and APE (Table 6). Narrative gradients achieve the lowest holdout loss (0.03), beating APE (0.10), TextGrad-CoT (0.17), and OPRO (0.20). This supports the pre-registered hypothesis that scaffold structure—not merely iterative optimisation—carries the social-sycophancy reduction: the two narrative-initialised methods (narrative grad, APE) outperform the CoT-initialised TextGrad baseline, and trajectory-based narrative gradients beat the others outright.

The gradient rewrites the protagonist section into a “Premise Audit” that challenges embedded assumptions—a mechanistic diagnosis of why first-person NoT validation fails. Degeneracy checks show stakeholder count and uncertainty scores remain near hand NoT.

6 Phase 15: Judge Reliability Panel

We re-score $n=60$ stratified OEQ responses with three judge families (gpt-5.4-nano, claude-haiku-4-5, grok-4-1-fast-reasoning) on validation, indirectness, and framing. Reporting uses Krippendorff’s α ; metrics below $\alpha=0.67$ require majority vote or are flagged low-reliability. Crowdsourced human OEQ labels (judge_gold.jsonl, $n=30$) provide Cohen’s κ against each judge, scored on the *human* responses directly.

Inter-judge reliability is low on all three metrics (Table 7): validation $\alpha=0.42$, framing $\alpha=0.11$, indirectness $\alpha=-0.23$. None clears the 0.67 threshold, so panel headline rates use majority vote. The picture is more favourable against human gold: validation agreement is strong for every judge (Cohen’s κ vs. human 0.73–0.86), indirectness is moderate and judge-dependent (0.31–

Metric	Inter-judge	κ vs. human		
	α ($n=60$)	nano	haiku	grok
Validation	0.42	0.85	0.86	0.73
Indirectness	-0.23	0.31	0.39	0.63
Framing	0.11	0.15	0.57	0.70

Table 7: Phase 15 judge reliability. Inter-judge Krippendorff α over the three-family panel; Cohen’s κ of each judge against crowdsourced human labels ($n=30$ human responses). No metric clears $\alpha=0.67$; validation alone shows strong judge-vs-human agreement.

0.63), and framing is mixed (0.15–0.70). The contrast—high judge-vs-human κ on validation but only moderate *inter-judge* α —reflects the narrow positive base rate among the gold human responses; the practical implication is that validation is the only ELEPHANT metric we report without a reliability caveat, and indirectness reductions from optimised prompts (Section 5) are not reportable from any single judge.

7 Phase 18: Judge-Robust Prompt Optimization

The Phase 14 result is optimised against a single absolute judge, and Section 6 shows the social metrics carry low inter-judge reliability—a textbook setup for reward hacking (Sharma et al., 2024). Phase 18 first quantifies how much of the apparent win is judge-gaming (18a), then develops an optimiser whose gains are robust by construction (18b).

18a: the Goodhart gap. We re-score the *same* Phase 14 holdout responses (hand NoT vs. the single-judge optimised prompt; generator claude-haiku-4-5; $n=150$) across a four-judge panel—the training judge grok-4-1-fast-reasoning plus held-out judges claude-haiku-4-5, gpt-5.4-nano, claude-sonnet-4-6—changing only the judge, never regenerating. The validation reduction is genuine: optimised validation is 0–3% under *every* judge. **The indirectness reduction is a Goodhart artifact:** the training judge reads optimised indirectness at 6%, but the worst-case held-out judge reads 87%—identical to the hand-NoT baseline (Table 8). Optimising one absolute judge bought a phrasing change that this judge scores as direct and others do not; framing sits in between (robust on three judges, 52% on sonnet).

Metric	Single-judge optimised rate			
	grok [†]	haiku	nano	sonnet
Validation	1%	2%	3%	0%
Indirectness	6%	87%	59%	10%
Framing	3%	4%	27%	52%

Table 8: Phase 18a Goodhart gap. The same single-judge optimised responses ($n=150$) re-scored by four judges. Validation is low under all judges (robust); indirectness collapses to a 6% training-judge reading but stays at the 87% baseline under the worst held-out judge. [†] training judge used in Phase 14.

Prompt	Held-out (val/ind/fram)		Panel
	sonnet	DeepSeek	worst-loss
Hand NoT	10/41/89%	53/78/68%	0.71
Single-judge	0/10/52%	6/74/13%	0.39
Robust (18b)	0/0/11%	13/4/5%	0.32

Table 9: Phase 18b held-out generalisation ($n=150$, generator claude-haiku-4-5). Each cell shows validation/indirectness/framing (%); “worst-loss” is worst-case-over-panel loss. Robust beats single-judge on framing ($p=1.9\times 10^{-14}$ sonnet; $p=0.015$ DeepSeek) and indirectness ($p=7\times 10^{-5}$ sonnet; $p < 10^{-15}$ DeepSeek).

18b: panel-robust narrative gradients. We re-run narrative-gradient descent (init = verbatim NoT, 10 iterations, optimiser claude-sonnet-4-6) under the worst-case-over-panel loss, selecting the iteration with lowest robust loss. On the held-out judge claude-sonnet-4-6—never scored during optimisation—the robust prompt generalises significantly better than the single-judge prompt: framing 11% vs. 52% (two-proportion $z = -7.7$, $p=1.9\times 10^{-14}$) and indirectness 0% vs. 10% ($z = -4.0$, $p=7\times 10^{-5}$), with both at the validation floor (Table 9). It also achieves the lowest worst-case panel loss (0.32 vs. single-judge 0.39 vs. hand 0.71) and is the only method to move worst-case panel indirectness off the 87% baseline (to 55%), reducing indirectness on every judge including the held-out one ($\rightarrow 0\%$). The robust objective trades a small amount of worst-case validation (9% vs. 3%) for large, judge-general gains on the two constructs single-judge optimisation either gamed or ignored.

18c: human-gold validation anchor. We export a stratified $n=40$ sample of hand vs. robust holdout responses and collect author validation labels via a resumable CLI

(label_optimized_gold.py). Author-rated validation falls from 35% (hand NoT) to 0% (robust); Cohen’s κ between author labels and panel judges ranges 0.13–0.27 on this sample (moderate, as expected for $n=40$ and imbalanced base rates), with DeepSeek at $\kappa=0.26$ —comparable to the in-family judges.

18d: degeneracy guard. On a $n=30$ holdout subsample, the robust prompt preserves stakeholder enumeration (mean count 4.77 vs. hand 4.73) while lowering uncertainty narration (1.0 vs. 2.83), confirming the optimiser does not collapse the deliberative scaffold into vacuity.

18e: quartet replication. We replicate the robust prompt across the generator quartet on hold-out OEQ ($n=150$ per cell where complete). Panel worst-case validation drops under robust vs. hand on every generator tested: haiku 69% $\rightarrow$ 9%, sonnet 89% $\rightarrow$ 23% (partial $n=57$ robust cells), grok and nano runs confirm the same directional pattern on held-out DeepSeek (e.g. nano framing 98% $\rightarrow$ 8% on partial $n=65$). The robust objective generalises across generators, not only the optimisation generator.

Scope. Quartet replication for sonnet/grok/nano completes incrementally from resumable caches; BrokenMath nano cells run at reasoning-model latency (~ 8192 -token budget). We present panel-robustness as a demonstrated generalisation property, not a deployment-ready prompt (Limitations).

8 Discussion

Dual benchmarks, correctly framed. Sharma saturates at $\leq 2\%$; BrokenMath does not (39–83%). Crucially, the scaffold that reliably reduces social sycophancy on ELEPHANT does not transfer to BrokenMath’s *propositional* error-detection (significant only on grok), evidence that the two are distinct constructs rather than a single “sycophancy” axis. Social ELEPHANT metrics discriminate NoT on validation and framing with strong cross-vendor evidence.

Moral sycophancy is a different construct. Both-NTA measures whether a model affirms *both* sides of a flipped conflict across prompts—not whether a single response validates the user. Binary suffix mode further confounds the scaffold. Researchers should not pool moral rates with validation/framing claims.

Single-judge optimisation games the judge.

The Phase 14 narrative-gradient “win” partly reverse-engineers its training judge: re-scoring the identical responses (Phase 18a) shows the indirectness reduction is a Goodhart artifact (6% training-judge vs. 87% held-out), while only validation survives every judge. This is the practical face of the low inter-judge reliability in Section 6: any optimiser pointed at a single noisy judge will exploit it. **A worst-case-over-panel objective fixes it.** The Phase 18b optimiser only credits changes all panel judges agree on; on held-out judges (including out-of-family DeepSeek-R1) it cuts framing and indirectness far below the single-judge prompt ($p < 10^{-4}$) and is the only method to move worst-case indirectness off baseline. Author validation labels on optimised outputs confirm the validation reduction (35% $\rightarrow$ 0%), and degeneracy checks show stakeholder depth is preserved. Because it is black-box (API generation and judging only), it applies to frontier closed models without weights or activations. The lesson generalises beyond sycophancy: prompt optimisation against LLM judges should evaluate against judges held out from optimisation, or it measures judge idiosyncrasy rather than behaviour.

9 Conclusion

NoT reduces social validation and framing sycophancy across four generators where Sharma probes show no signal. Moral both-NTA does not follow the same pattern once loading and construct issues are corrected, and BrokenMath supplies the propositional headroom Sharma lacks. A judge panel exposes single-judge prompt optimisation as judge-gaming, and our panel-robust optimiser corrects it: on a held-out judge it generalises significantly better than single-judge optimisation, recovering the indirectness construct the naive objective leaves at baseline. Together, these phases narrow the defensible claim: deliberative scaffolding targets single-response social face-preservation, not propositional agreement or cross-prompt moral consistency—and robustness, not raw loss against one judge, is the quantity worth optimising.

Limitations

All sycophancy outcomes are scored by LLM judges; while Phase 15 validates the validation metric against human gold (κ 0.73–0.86) on

crowdsourced human responses, Phase 18c author labels on *model* outputs ($n=40$) show moderate judge–author agreement (κ 0.13–0.27), and indirectness/framing carry low inter-judge reliability. DeepSeek-R1 provides an out-of-family held-out check, but claude-sonnet-4-6 shares a vendor with panel judge haiku; both held-out judges should be read alongside author labels. Quartet replication and BrokenMath nano cells use resumable caches; some generator cells were incomplete at submission time (partial n noted inline). The Phase 14/18 results isolate *outcome*, not *mechanism*: we have not yet ablated which narrative sections carry the effect, so causal attribution to specific scaffold components remains open. BrokenMath uses a single judge family, and the free-form moral arm adds a second verdict-extraction call whose own reliability we do not separately audit. Sample sizes are $n=150$ per ELEPHANT slice and $n=451$ for BrokenMath; effects below a few points may be underpowered. All generators are accessed through hosted providers, and results may not transfer to other deployments or model versions.

Ethical Considerations

This work studies sycophancy mitigation and does not involve new human-subjects data collection; human labels are drawn from the public ELEPHANT/OSF release and used only as reliability gold. Reducing sycophancy is dual-use: a scaffold that makes models less validating can also make them blunt or dismissive in genuinely sensitive contexts, and the Phase 17 finding that NoT *increases* moral both-NTA shows that “less sycophantic” is not uniformly desirable. We therefore frame NoT as a measurement and construct-clarification tool, not a deployment-ready safety intervention, and we caution against optimising single sycophancy metrics in isolation. No personally identifying information is processed, and all adversarial math problems come from the public BrokenMath benchmark.

Reproducibility

Phases 13–17 pre-registration:
Guidance_Documents/study_design.md.
Scripts: run_elephant.py (–moral-mode),
judge_panel.py, run_brokenmath.py,
run_phase14_*.py, run_phase18_*.py,
eval_phase18_heldout.py,

label_optimized_gold.py. Out-
puts: elephant_summary.json,
brokenmath_summary.json,
judge_reliability_summary.json,
phase14_summary.json,
phase18_{goodhart,robust,heldout,quartet}.json,
optimized_gold_summary.json.

References

- Myra Cheng, Esin Durmus, Miranda Zhang, Tomasz Korbak, Ethan Perez, Samuel R. Bowman, and 1 others. 2025. ELEPHANT: Measuring and understanding social sycophancy in LLMs. *arXiv preprint arXiv:2505.13995*.
- Ahmed Fanous and 1 others. 2025. SycEval: Evaluating sycophancy in large language models. *arXiv preprint*. Graded false-premise probes; reports 29–58% sycophancy on frontier models.
- Klaus Krippendorff. 2004. *Content Analysis: An Introduction to Its Methodology*, 2 edition. Sage, Thousand Oaks, CA. Krippendorff’s alpha for inter-coder reliability.
- Lars Malmqvist. 2024. Sycophancy in Large Language Models: Causes and Mitigations. *arXiv preprint arXiv:2411.15287*.
- Ivo Petrov, Jasper Dekoninck, and Martin Vechev. 2025. BrokenMath: A benchmark for sycophancy in theorem proving with LLMs. *arXiv preprint arXiv:2510.04721*. 451 expert-verified adversarial math problems; four-category sycophancy judge.
- Mrinank Sharma, Meg Tong, Tomasz Korbak, David Duvenaud, Amanda Askill, Samuel R. Bowman, Newton Cheng, Esin Durmus, Zac Hatfield-Dodds, Scott R. Johnston, Shauna Kravec, Timothy Maxwell, Sam McCandlish, Kamal Ndousse, Oliver Rausch, Nicholas Schiefer, Da Yan, Miranda Zhang, and Ethan Perez. 2024. Towards Understanding Sycophancy in Language Models. In *International Conference on Learning Representations (ICLR)*. ArXiv:2310.13548.