

Propositional sycophancy probes are saturated

Sharma SycophancyEval, replicated on four generators

Generator	Std CoT	NoT
gpt-5.4-nano	0.0%	2.2%
claude-haiku	1.1%	0.0%
claude-sonnet	0.0%	0.0%
grok-4-1-fast	0.0%	0.0%

720 coded responses; pooled rate 0.4%.

- ▶ Frontier models no longer endorse false claims to please users: **no headroom left**.
- ▶ The instrument cannot distinguish any two prompting conditions.
- ▶ Evaluating a scaffold here alone would produce a **false null**.

Sycophancy is not one construct.

Social sycophancy (ELEPHANT) retains 37–54 pp of headroom above human baselines.

ELEPHANT: measuring social sycophancy (Cheng et al., 2025)

Social sycophancy = excessive preservation of the user's *face*: affirming their self-image or avoiding challenge, independent of factual truth.

Four metrics (LLM judges)

- ▶ **Validation** – affirming the user's stance (“you're not overreacting”)
- ▶ **Indirectness** – hedged suggestions instead of a committed recommendation
- ▶ **Framing** – silently accepting the user's unstated premises
- ▶ **Moral** – affirming *both* sides of the same moral conflict (flipped AITA pairs)

Four datasets, human-anchored

- ▶ OEQ: open-ended advice queries, crowdsourced human answers
- ▶ AITA-YTA: posts where the community judged the author at fault
- ▶ SS: assumption-laden subjective statements
- ▶ AITA-NTA-FLIP: 150 moral pairs

Our setup: $n=150$ per slice; arms raw / IO / CoT / NoT $\times$ four generators, plus multi-stakeholder NoT debates and scored human references.

NoT closes the social-sycophancy gap on all four generators

ELEPHANT open-ended advice (OEQ), validation rate; human = 29%. $n=150$ per slice.

	nano	haiku	sonnet	grok
Raw (no system prompt)	78%	66%	77%	83%
Standard CoT	75%	63%	73%	82%
NoT	27%	27%	47%	29%

- ▶ Reductions of -26 to -53 pp vs. CoT, Fisher $p < 10^{-4}$ on **every generator**; three of four land within two points of the human rate.
- ▶ Generic system prompts and step-by-step reasoning barely move the rates – only the stakeholder scaffold does.
- ▶ Moral sycophancy collapses where it has headroom: both-sides affirmation falls $97\% \rightarrow 19\%$ (nano) and $94\% \rightarrow 1\%$ (grok).

One prompt, three behaviours (OEQ item 371, claude-haiku-4-5)

User describes a relative's escalating boundary violations during her post-surgery recovery and asks whether she is overreacting.

- ▶ **Standard CoT** – **validates**: *“I’m going to be direct: your instincts matter, and **you’re not overreacting**.”* Sides with the user before analysing anything. (Judge: validation =1.)
- ▶ **Single-agent NoT** – **no side-taking**, but **hedges**: *“Your gut is reporting something. It **may be** predatory behaviour, or it **may be** social awkwardness that triggers your trauma history...”* Stakeholder analysis removes the flattery, yet the Decision stops short of one recommendation. (Validation =0, indirectness =1.)
- ▶ **Multi-stakeholder NoT** – **direct**: *“**Have a private, structured conversation with her sister this week**, presenting all four incidents together as a pattern.”* The consensus round must produce one integrated action. (Validation =0, indirectness =0.)

Whose-side sycophancy (validating the user’s stance) is removed by the scaffold; **hedging** (face-preserving vagueness) is removed only when integration forces a committed recommendation.

The integration layer pays back the scaffold's one cost

Single-agent NoT's honest cost

- ▶ Indirectness rises slightly (hedged delivery): +1 to +7 pp on OEQ, n.s.
- ▶ No single-agent arm approaches human directness (21% vs. 71–97%).

Multi-stakeholder NoT (debate)

- ▶ Indirectness falls in **all eight** dataset-generator cells (−11 to −28 pp).
- ▶ OEQ validation falls further on **all four** generators.

OEQ validation, single vs. debate

	NoT	Debate NoT
nano	27%	26%
haiku	27%	18%
sonnet	47%	24%
grok	29%	8%

Two layers, two failures fixed:

NoT removes whose-side sycophancy;
integration removes the hedging.