

Narration-of-Thought: Represented Opposition for Objectivity in Single Agents and Collectives

Anonymous ACL submission

Abstract

Language models are adept at conversation, offering advice and suggestions with apparent empathy and understanding. Whatever the nature of that understanding, it is widely accepted to break down under more complex social conditions, and the sharpest case is sycophancy, the answer pulled toward the person asking. The social account of objectivity holds that organised criticism across perspectives exposes errors that no single mind corrects, and on that account we introduce a general, lightweight prompt scaffold structured as a narrative. Narration-of-Thought (NoT) guides the model through a first-person narrative that identifies the decision-maker and what they know, considers every affected party and their interests, traces the consequences of the available actions, acknowledges uncertainty, and only then commits. We show that the scaffold can be used in single-agent and multi-agent settings on a range of nuanced social reasoning tasks. Used alone, it reduces validation of the asker’s own account on open-ended personal dilemmas by 22 to 70 points relative to standard chain-of-thought on seven models from six vendors, three of them open-weight, a relative reduction of up to 87 percent, and the reduction is three times as large on accounts that name a second party, an association rather than a manipulated cause. **[TODO: Pillar two (theory of mind) is a skeleton for a future developer and is not claimed in this paper; no sentence on it belongs in the abstract until it is in scope.]** Used to build a collective, copies of one model given opposing sides support more objective, grounded judgements through structured criticism, and reviewing the cases flagged by a second dissenting seat with a stronger judge that has no side raises accuracy by 7.0 points over the collective alone in our primary evaluation, a gain that holds on every topology and model we tested it on, and by 2.4 points over the strongest single model on the model whose seats lock to their assigned side, a gain that does not repli-

cate on a model whose seats do not. We map where each effect holds and where it does not, and our findings lay the groundwork for further work on the scaffold as an instrument of social alignment.

1 Introduction

Language models are adept at conversation. They give advice and suggestions with apparent empathy and an apparent understanding of the person asking. Whatever that understanding is and how- ever it arose, it is widely accepted to break down under more complex social conditions, when the person asking has a stake in the answer, when the account they give is one side of a dispute, or when several parties want different things. The sharpest case is sycophancy. Ask a model whether you were in the wrong and it answers more gently than it would a stranger with the same account (Sharma et al., 2024; Cheng et al., 2025a), and preference training builds this in as a bias prescribed by the user. The general problem is a conclusion pulled away from the evidence by something else in the room, the asker’s feelings, an accepted premise, an assigned side, and a fix for sycophancy alone leaves the others.

The social account of objectivity says where a remedy lies. On the argumentative theory of reasoning individuals argue for a side, and on the social account of objectivity truth is a property of the organised criticism among them rather than of any one mind (Mercier and Sperber, 2011; Longino, 1990). A model answering a single user has no critic in the room. Our proposal is to put one there, in two ways. Narration-of-Thought (NoT) is a lightweight prompt scaffold that guides the model through a first-person narrative that identifies the decision-maker and what they know, names every affected party and their interests, traces the consequences of each available action for each of them, acknowledges what remains uncertain, and only

then commits. The critic is represented inside the trace. The collective goes further and assigns the interests to separate copies of the model, one speaking for the writer of an account, one for the other party, one for no one, so that the criticism is real and its dissent can be read by a simple rule, that a statement against the speaker's assigned interest is informative and a statement with it is not (Milgrom and Roberts, 1986; Lipman and Seppi, 1995).

Two outcomes follow, and the paper reports each at the scope its evidence supports. Used alone, the scaffold reduces validation of the asker's own account on open-ended personal dilemmas by 22 to 70 points relative to standard chain-of-thought on seven models from six vendors, three of them open-weight, a relative reduction of 31 to 87 percent, with a token-matched control that does not reproduce it and a reduction three times as large on accounts that name a second party than on accounts that do not, 37.7 against 11.3 points. A second judge from a different vendor confirms the drop on six of the seven models, and a third reverses its sign on three, a pattern traced to that judge sitting near ceiling on almost every response regardless of content rather than to a change in the effect itself. The same scaffold does not remove the asker's pull on a mechanically scored verdict instrument, so the claim is about the register of the answer, whether it endorses the asker's account, and not about which verdict is reached. **[TODO: Pillar two, theory of mind, is out of scope for this paper (PI decision, 2026-09-21) and left as a skeleton for a future developer; no claim is made here.]** Used to build a collective of three seats and a moderator with no side, the scaffold's assigned opposition produces a dissent that is a sensor. A verdict with one losing party gives exactly one advocate a reason to object, so when a second seat objects as well the collective is wrong nearly ten times as often, 0.39 against 0.04, and reviewing only the flagged fifth of cases with a stronger judge that has no side raises accuracy by 7.0 points over the collective alone and by 2.4 points over the strongest single model. The flag belongs to the assigned roles and not to the conversation, since it survives cutting every line between the seats and replacing narration with plain chain-of-thought in the seats, and it needs the roles to hold. On a model whose seats do not lock to their sides the all-strata sensor and the routing gain replicate but the one-loser mechanism does not, which is what the reading rule predicts.

The paper's claim is one sentence long. Objec-

tivity is obtained by placing an opposed interest into the process, represented inside one trace or assigned to a separate seat, and it is a property of that arrangement and not of any copy. Our contributions are the scaffold and its single-agent evaluation on social sycophancy with the controls that scope it, the collective and its dissent sensor with the ablations that locate the signal in the assigned roles, and a map of where each effect holds and where it does not. Every experiment was pre-registered with a spending ceiling and a stop rule before any model call, every interval is an item-clustered bootstrap, and every citation was fetched at a primary page.

2 Theory

A model's answer is pulled toward the interest that is present in its context. On an open-ended dilemma the interest is the asker's own account of the conflict, on a false premise it is the premise the user owns, and on an assigned side it is the side. The pull cannot be removed from inside the same context, because the copy that would remove it is the copy that carries it. Three copies of one model vote as one voter and the gain from a correlated jury is bounded by how far its members' errors are independent (Ladha, 1992; Boland, 1989; Dietrich and Spiekermann, 2013; Kim et al., 2025a). What can be done is to place an opposed interest into the process so that agreement across interests, or dissent against interest, becomes informative. This is the argumentative theory of reasoning (Mercier and Sperber, 2011, 2017) joined to the social account of objectivity (Longino, 1990, 2002; Popper, 2011). Individuals argue for a side, and truth is a property of the structure of criticism among them rather than of any one mind. The scaffold and the collective are the two ways of placing the opposed interest.

Internal opposition is the scaffold. One trace is made to represent the other parties. It must name every party with a stake, trace what each available action does to each of them, and state what remains unknown before it commits. The represented counterparty is the content of the intervention, and the prediction is that the reduction in sycophancy is a counterparty effect and not a manner effect. Three consequences were testable, and they resolve differently. It should be concentrated on accounts that contain a party other than the asker and absent where none exists, and it is: the reduction is three times as large on accounts that name a second

party than on those that do not, and a false theorem, which has no second party, is the one place the scaffold affirms the falsehood more often rather than less. It should be a property of the prompt's structure, reproducing across vendors and on open-weight models, and it is, on seven models from six vendors, three of them open-weight, with the effect if anything larger on the open-weight deployments than on any proprietary model. But the prediction that the reduction should be carried specifically by the sections that name and trace consequences for the other parties, rather than by the section that licenses hedging, does not hold at the level of the scaffold's own sections. Removing any one of the three candidate sections, the stakeholder list, the consequence trace, or the uncertainty statement, takes back a large share of the reduction on one model and almost none of it on two others, and pooled across models no single section clears a quarter of the reduction with a confidence interval excluding zero. The represented-counterparty account holds at the level of what the account contains, an item that names a second party loses more validation than one that does not, and does not hold at the level of which sentence in the scaffold does the work: no one section is the carrier, on most models the effect is not attributable to any single part, and where it is attributable at all it does not cleanly separate the counterparty account from the hedging account. A direct manipulation makes the section-level failure sharper still. Telling a model to consider everyone affected, with none of the scaffold's other structure, does not lower validation. It raises it. Counterparty-awareness alone reads as an invitation to a broader, more accommodating response rather than as a check on one. The theory's item-level prediction is supported and its causal, section-level prediction is not: what the item-level association describes correctly is which accounts the reduction concentrates on, not what produces the reduction, and the paper states this distinction rather than collapse one finding into the other.

External opposition is the collective. When the interests are separate seats the credible-signal rule applies. A statement against the speaker's assigned interest is informative and a statement with it is not (Milgrom and Roberts, 1986; Lipman and Seppi, 1995; Glazer and Rubinstein, 2001). A verdict with one losing party gives that party's seat a reason to object, and its objection carries nothing. A second objection has no side to explain it and must come from the account. The dissent edge exists only

where the interests are opposed, and it survives removing the exchange between seats entirely, down to about half its discriminating power rather than to nothing (Dewatripont and Tirole, 1999; Krishna and Morgan, 2001). So a copy that never reads another seat's argument, and sees only the account and the moderator's verdict, still objects almost entirely on the strength of its assigned interest. The exchange between seats adds the remaining half, not the phenomenon itself.

The sign of the mechanism is not uniform across models, and the lift is more robust than the sign. On the model with the highest role-lock the two advocates' objections are strongly anti-correlated, exactly as the credible-signal rule predicts: the losing seat objects and the winning seat does not. On two other models, one with low role-lock and one with role-lock nearly as high as the first, the two advocates' objections correlate positively instead, closer to two independent readers both noticing a hard case than to one side conceding against interest. The within-one-loser accuracy lift is positive with a confidence interval excluding zero in every one of these conditions regardless of the sign of the correlation, so role-lock alone does not decide whether the against-interest mechanism specifically is operating, only whether the lift is present at all. The theory's claim that objectivity comes from placing an opposed interest into the process is supported by the lift's robustness; the specific reading rule, that a second objection is informative because it goes against the objecting seat's own assigned side, is confirmed on the model that locks to its role and is not established on the others, where the lift more plausibly reflects correlated difficulty than opposed interest.

The collective is a sensor and not a decider. No arrangement beats its best seat by vote, the moderator's competence sets the verdict (Shin, 1998), and the gain is realised by routing the flagged cases to a stronger judge with no side. This part of the claim is the most robust of all: routing beats the collective alone on every condition tested, on-edge and off-edge, on the model that locks to its role and on the two that do not, independently reproduced by a second implementation in every case. Beating the strongest single model, rather than the collective, is weaker and model-specific: it holds on the role-locking model and does not on the others. That is the precise sense in which the collective depends on the same control the scaffold exploits. Embodiment has to hold for the against-interest reading

of opposition to be readable, even though a weaker, sign-agnostic dissent signal survives without it.

The game-theoretic reading is used only where its primitives exist, an assigned interest, a receiver whose verdict is separate from the message, and a decision read independently of the speaker. Those exist in the collective and not in a single trace, so for the scaffold alone the paper claims a change in the output distribution whose content is the represented counterparty and makes no signalling claim.

[TODO: Pillar two, theory of mind (represented minds). Naming what another party knows and wants is an expression-side act, and whether the scaffold closes the gap between what a model represents about other minds and what it expresses is a live question this paper does not answer. Out of scope for the present work (PI decision, 2026-09-21); left here as a stub for a future developer. Stage-1 data exists (Addendum 18, four models, 320 ToMBench items, a guard-clean negative result with an identified mechanism) but is not part of this paper’s claims and is not pooled.]

This paper validates two of the arrangement’s readouts. The asker’s pull on a single trace is pillar one, and the readable dissent that separate seats produce is pillar three. Two results tie the pillars to each other rather than merely placing them side by side. Narration does not create the collective’s structure, so internal and external opposition are different objects with different evidence. And without role lock the external structure collapses to correlated readers, so external opposition rests on the control that assignment gives over what a copy will say.

3 Method

3.1 The scaffold (pillar 1)

Narration-of-Thought is a system prompt of five sections that the model is asked to write in the first person before it answers. The first names and characterises the decision-maker, who they are, their role and what they know. The second lists every person whose life intersects the decision and what is at stake for each. The third narrates, for each available action, its consequences at least two steps into the future for each stakeholder. The fourth states what remains genuinely uncertain about each projected future. The fifth commits to a decision and explains, within the narrative frame, why that trajectory is preferable to the alternatives. The

prompt is the same on every model and every task in this paper, and the standard chain-of-thought control is a one-line instruction to think step by step before answering.

We chose the five sections because each one places a party or a consequence into the trace that the asker’s account leaves out. The stakeholder section is a represented counterparty, the consequence section traces what the answer does to that counterparty, and the uncertainty section withholds the commitment until both have been written. The theory in Section 2 predicts that the first two carry the reduction in sycophancy and the third does not, and the knockout experiment that tests that prediction removes one section at a time under new arm names so that no cached response can be mistaken for another.

The primary instrument is the open-ended personal dilemma set of the ELEPHANT benchmark (Cheng et al., 2025a), 150 items sampled once with a fixed seed. Each response is scored by a judge model for validation, whether the response endorses the asker’s own account of the situation, the construct on which human respondents validate 29 percent of the time in the benchmark’s own labels. The judge sees the full response. The production scorer had passed the first 4,000 characters only, which cut the scaffold’s longer responses before the judge reached their conclusion, and every score in this paper is from a re-score of the complete text. One response per item and arm is generated on each of seven models from six vendors, the four proprietary models of the original study and three open-weight deployments, and the outcome is the difference in validation rate between the scaffold and the control, with a percentile interval from a bootstrap that resamples items. Where an arm returned no response we report the share and bound the difference by scoring every missing response first as validating and then as not.

Three controls scope the claim. A verbose control asks the standard prompt for a response of the scaffold’s length without its sections, and a length-matched comparison restricts both arms to responses of the same length where the arms share support. A counterparty coding marks each item for whether the account names a specific other party with an interest in the outcome, and the difference is read within each stratum. And a mechanically scored instrument with a checkable verdict, disputes from the Scruples corpus with a gold label from more than a hundred crowd votes per item,

presents the same account in the third person and as the asker’s own, so that any movement of the verdict toward the asker is read directly from the letter the model outputs rather than from a judge. Judge reliability on the validation construct is reported in the appendix, with a three-judge panel from three vendors and each judge’s agreement with human gold labels.

3.2 Theory of mind (pillar 2, out of scope, skeleton for future work)

[TODO: Pillar two is not a claim of this paper. The design below is left as a working skeleton (code and data exist and run) for a future developer, not as part of the present evaluation.]

The scaffold would be evaluated unchanged on ToMBench (Chen et al., 2024b), a multiple-choice theory-of-mind benchmark of eight tasks, from which we draw forty items per task with a fixed seed, 320 items in all. The false belief task is the pre-declared primary stratum because it is the classic test of reasoning about another agent’s knowledge state. The system prompt is the scaffold or the control, the user turn is the story, the question and the four options with one fixed instruction to end with an answer letter, and the answer is scored mechanically against the gold letter with no judge. The registration is sign-agnostic. The abstract’s sentence on theory of mind is written from the pooled interval over seven models, whatever its sign, and a model on which more than one response in ten has no parsable answer is excluded and named.

3.3 The collective (pillar 3)

3.3.1 Embodiment as a control

A seat is a copy of one model reading the account under a brief that names whom it speaks for and what it wants (Figure 1A). We chose a side over a persona because a persona changes how a copy speaks and a side fixes what it concludes, which makes it a control. Told it speaks for the writer, a copy finds the writer not at fault on 99.7 percent of opening statements, and for the other party, at fault on 97.4 percent. Two seats hold opposite sides because what two locked seats agree on came from neither brief and must come from the account (Milgrom and Roberts, 1986; Dewatripont and Tirole, 1999). A third seat has no side because the verdict layer needs one uninterested reader, and where the sides lock it is the seat any vote returns and the best seat (Section 4.3). Neither advocate may invent facts, and the writer’s seat may not assert that the




Figure 1: (A) Three copies of one model read one account, briefed for the writer, the other party or no one. A side fixes a seat’s conclusion, so opposite sides agree only through the account. (B) The seats open, rebut and restate, and a moderator with no side writes a verdict that each seat accepts, accepts with one stated change or rejects. (C) The losing party’s seat is expected to object. A second objection has no side to explain it and is the flag, routing the case to a stronger uninterested judge.

writer is not in the wrong where the account does not support it, so no objection can rest on what the account lacks. Provability need only be partial, since a decider who knows the speakers’ interests conflict can count on one to correct a mistaken inference (Lipman and Seppi, 1995). Every seat narrates before it commits, the light dose, under the model’s own scaffold so arrangement is the only difference. Removing it later showed it amplifies a second seat’s objection without creating the pattern. The writer’s brief treats a finding against the writer as a finding against its party and puts the case as strongly as the account allows. The other party’s brief treats a finding for the writer as a finding against its party, one that did not write the account and cannot add to it, and puts that case as strongly as the writer’s own account allows. The neutral brief speaks for everyone who will read the verdict and wants only the finding the account supports.

3.3.2 The collective as a graph

An arrangement of seats is a graph, its nodes the seats and a moderator, its edges who reads whom (Figure 1B), three copies of one model running seven stages in seventeen calls. Each seat opens alone with a position and a verdict, so lock is measured first. Each then rebuts the other two and

restates a final position, the only stages where seats read one another, six edges the no-edge control cuts. A moderator, a separate node with no stake, reads the account and the three final positions and writes one synthesis and verdict, separating arguing from deciding. It is the seats' model by default, so the collective adds only an arrangement, and was varied later because nothing in the design fixes who decides. At the objection round each seat accepts the synthesis, accepts it with one stated modification, or rejects it naming the concern no integration could absorb. A stated change is required because the synthesis is the first concrete proposal there is to change, because integration must answer something, and because a seat that accepts once its own change is addressed has revised rather than opined. An objection without one is a counted parse failure. The moderator then integrates, saying whose request each part answers, and each seat votes accept or reject.

A moderator decides rather than a vote because where the sides lock a vote adds nothing. Two fixed opposite verdicts cancel, so the majority is the neutral seat's verdict, any weighting (Shapley and Grofman, 1984) can only reweight two constants, and three draws of one model lack the independent errors a jury theorem needs (Ladha, 1992). The moderator is instead the node whose verdict the seats criticise (Mercier and Sperber, 2011; Longino, 1990), sitting off the argument because a network converges only when no node dominates it (Golub and Jackson, 2010). Across cells the item text, verdict instruction and briefs are byte-identical except for what a control varies, so cells differ only in arrangement.

3.3.3 The flag

Suppose the verdict goes against one party (Figure 1C). Only that party's seat has reason to object, and that expected objection says nothing about whether the verdict is right. A second objection comes from the seat whose side won or the seat with none. No brief put it there, so it must come from the account, and an admission against interest is what a listener who knows only the interests can trust (Milgrom and Roberts, 1986; Lipman and Seppi, 1995). On one-loser verdicts exactly one seat objects on 92 percent of debates, two or more on 7 percent. As implemented the flag is a stake-blind count, firing when at least two of the three seats, the uninterested seat included, object at the objection round. An earlier stake-aware com-

posite, firing on the neutral seat's objection or an advocate's against interest, needs each seat's stake and the verdict's direction, which the plain count matches without either. Because the flag marks debates most likely to be wrong, a stronger uninterested judge is paid for only there, answering flagged cases cold by majority of three samples, ties to the collective. The flag needs a verdict with one loser, a hypothesis about the verdict rather than the model, and fails by design in three settings. A verdict blaming both parties or neither gives both sides reason to object and it floods. Seats that never leave their sides give no second objection and it is silent. A plan every party gains from has no loser and nothing to read.

3.3.4 Controls that isolate each decision

Each decision was tested by removing or varying it, everything else held at the registered protocol, further departures disclosed in the record, outcomes in Section 4.3. To test whether the stake is load-bearing we cut the first-person alignment and the stake clause from every brief and the preamble, keeping the instruction to argue as strongly as the account allows. If lock and the flag survived, a seat is fixed by its direction, not its interest. To test whether opposed sides are load-bearing we made the roles identical, three copies of the neutral brief. If the flag survived, it belongs to three readers, not opposition. To test whether the exchange is load-bearing we cut the edges, each seat developing its own opening instead of rebutting. If the flag survived, it belongs to the sides, not the conversation. To test whether narration is load-bearing we replaced each seat's scaffold with plain step-by-step reasoning. If the flag survived, narration is manner, not mechanism. To test whether the decider's model matters we put another vendor's moderator over the same seats, their opening, rebuttal and re-statement replayed from cache and their labels and votes written afresh. If the verdict did not move, the moderator counts verdicts rather than reading arguments. To test whether a second seat on the same side adds anything we seated a like-biased ally beside the writer's under a mirrored brief. If it added dissent, number rather than opposition carries the information (Krishna and Morgan, 2001). To test whether the one-loser condition is a limit or an artefact of two advocates we seated an advocate for a third person the account names, so a verdict could have two losers. If the flag kept its meaning, the condition is an artefact. To test whether same-

model seats limit the verdict we served the three identical no-edge seats from three vendors. If they aggregated no better than three copies, shared errors are not the limit (Ladha, 1992). To test whether repetition alone deepens a perspective, one agent re-read and re-answered its own false-premise answer, against a control fed its untreated one. If it did, the collective would be mere iteration.

3.3.5 Instruments, measures and estimation

Disputes (Table 1) are Scruples Anecdotes, first-person conflict accounts whose crowd verdict is gold when at least 90 percent of at least 50 votes agree, 249 items under two framings, 210 after one model’s refusals. [TODO: references.bib] **Scruples entry, Lourie, Le Bras and Choi 2021, PI to add** False premises are 100 BrokenMath (Petrov et al., 2025) problems with a false statement, six samples per item on gpt-5.4-nano (OpenAI, 2026), under heads making the model the statement’s author, a colleague, a stranger or a rival. Dilemmas pair two anecdotes, a fresh crowd naming the worse action, ELEPHANT (Cheng et al., 2025a) supplies 150 personal questions judged for validation of the asker, and 144 ordinal two-by-two games with a first-principles gold plan ran binding, non-binding, and as coordination on nine two-equilibrium games. Dispute verdicts, parsed from a forced line and never by a judge model, collapse to at fault (YTA, ESH) or not (NTA, NAH), dropping UNRESOLVED, while on false premises UNRESOLVED is never correct and stays in the denominator.

Role-lock is how completely a side fixes the opening verdict, the absolute difference between the advocates’ opening rates of finding the writer at fault. The fire rate is the share of debates on which the flag fires, over all debates in the two-by-two readouts and over codable debates when read as routing coverage. Lift is the probability the integrated verdict is wrong given the flag minus that given no flag, error concentration their ratio, both read within verdict strata because the flag floods on both-party verdicts. Transfer is the same flag read against a model that never saw the debate, answering alone by majority over three samples. Composed accuracy is the deployed system’s, with claude-haiku-4-5 or a cached claude-sonnet-4-6 as judge.

Every interval is an item-clustered percentile bootstrap with 2,000 to 10,000 draws, families Bonferroni-corrected, missing false-premise units

Table 1: Instruments. Item counts are the panels the reported statistics use. Models by vendor card (xAI, 2025; Anthropic, 2025, 2026; OpenAI, 2026).

Instrument	Items	Gold	Verdict set
Disputes (AITA)	249 panel, 210 deliberated	crowd plurality, $\geq 90\%$ of ≥ 50 votes	YTA, NTA, ESH, UNRESOLVED
False premise (BrokenMath)	100, all gold	benchmark, false stratum	TRUE, FALSE, UNRESOLVED
Comparative (Dilemmas)	80 then 700, 624 codable	fresh crowd, ≥ 4 annotators	ACTION_A, ACTION_B, UNRESOLVED
Open-ended (ELEPHANT)	150 per generator	none, judge-scored validation	validates, does not
Games (TMG-Bench)	144, regions 16 / 97 / 31, nine two-equilibrium	first-principles plan	PLAN_xy, NO_UNIQUE_PLAN, or NO_RECOMMENDATION

model	CoT	NoT	drop (points) [95% CI]	relative
claude-haiku-4-5	0.627	0.360	-26.7 [-36.7, -16.7]	43%
claude-sonnet-4-6	0.727	0.503	-22.3 [-31.3, -13.3]	31%
gpt-5.4-nano	0.759	0.345	-41.4 [-51.1, -31.4]	55%
grok-4-1-fast-reasoning	0.815	0.455	-36.0 [-44.9, -27.1]	44%
Llama-3.3-70B-Instruct	0.806	0.106	-69.9 [-77.7, -61.4]	87%
Mistral-Large-3-2	0.775	0.193	-58.2 [-66.2, -49.7]	75%
DeepSeek-V4-Pro	0.944	0.489	-45.5 [-53.8, -37.1]	48%

Table 2: Validation of the asker’s account on 150 open-ended dilemmas, scored from the complete response. Drop is the scaffold minus the control in points with an item-clustered 95 percent interval. Relative is the drop over the control rate. Bracket bounds the drop over responses an arm did not return. Length-matched restricts both arms to responses of the same length where the arms share support, with the number of responses used.

Manski-bounded. Two guards run before any number is read, truncation and no legal verdict each at most five percent per cell and call, failing runs going unread. Controls pass a parse gate of at most 0.05 unparsed labels per debate above the embodied baseline, a lift needs sixty flagged debates, and collective statistics use codable debates, where both moderator verdicts commit. One pre-registration document with dated addenda fixed hypotheses, estimators, kill criteria, ceiling and stop rule before any call, results appended, nothing above them edited, two overruns reported.

4 Results

4.1 Pillar 1, the scaffold reduces validation of the asker’s account

On every model the scaffold lowers the rate at which the response endorses the asker’s account, by 22.3 points on the model that moves least and

69.9 on the model that moves most, which is 31 to 87 percent of the control rate (Table 2). The single largest relative-attrition cell, one of the four original proprietary models, is also the least certain because that model returned no response on 23 percent of scaffold items against 3 percent of control items, and scoring the missing responses at either extreme bounds its drop between 24 and 50 points, a materially wider bracket than any of the three open-weight deployments carry despite their larger point estimates. A control that asks the standard prompt for a response of the scaffold's length raises validation rather than lowering it on three of the four original models and lowers it on the fourth, so length is not what carries the effect, though that control contains three of the scaffold's instructions in prose and is a control on wording as much as on length. It was never run on the three open-weight deployments and remains a gap specific to them. Restricting both arms to responses of the same length, now computed on all seven models, leaves a drop on every one of them, from 24.6 points on the model with the smallest length-matched effect to 78.6 on the model with the largest, confirming length is not the mechanism on the open-weight deployments either.

The reduction is concentrated where a counterparty is represented, at the level of what the account contains. On the 110 accounts that name a specific other party with an interest in the outcome the pooled drop is 37.7 points, and on the 40 that do not it is 11.3, a difference of 26.3 points with an interval that excludes zero, and the difference holds on three of the four models taken separately. The residual drop on accounts with no counterparty is real and is the part the represented-counterparty account does not explain.

The reduction also reproduces on three open-weight deployments from three further vendors, and reproduces there more strongly than on any of the four proprietary models. Llama-3.3-70B-Instruct falls 69.9 points, from 0.806 to 0.106, Mistral-Large-3-2 falls 58.2 points, and DeepSeek-V4-Pro falls 45.5 points, all with confidence intervals excluding zero and with non-response under 7 percent in every arm, a materially cleaner attrition profile than the model most affected by missing data among the original four. Seven models from six vendors, three of them open-weight, now show the same negative, significant drop.

The claim survives a change of judge on most, not all, of these models. A second judge from a dif-

ferent vendor confirms a negative, significant drop on six of the seven; a third judge reverses the sign on three of them, a pattern traced to that judge's own validation rate sitting near ceiling on almost every response regardless of its content, which is reported as a property of that judge and not as evidence against the effect. One model, the one answering as sonnet, does not survive a change of judge under either reading, and its reduction should be read as judge-specific.

The discriminating test of which section carries the reduction does not resolve the question at the level of the scaffold's own parts. Removing one of the three candidate sections at a time, the stakeholder list, the consequence trace, or the uncertainty statement, on the same 150 items and the same three models, pooled across models no single section's removal takes back a quarter of the reduction with a confidence interval excluding zero. The pooled result hides real disagreement between models. On the model with the smallest overall reduction, removing any one of the three sections gives back close to half of it, so the whole scaffold is load-bearing there and the result does not favour the counterparty account over the competing hedging-register account. On the two models with the largest reductions, removing any single section gives back almost nothing, with a confidence interval that includes zero on every knockout, so the reduction there does not depend on any one part of the prompt at all. Moving the commitment to the front of the scaffold instead of the end does not change the validation rate with a confidence interval excluding zero on any model. The represented-counterparty account is therefore supported by what an account contains, an item that names a second party loses substantially more validation than one that does not, and is not supported by which sentence in the scaffold does the work. No single section is the carrier of the effect, on most models the reduction is not attributable to any one part, and where it is attributable at all it implicates the whole scaffold rather than the sections that name the other parties specifically.

A direct manipulation tests the causal version of the claim outright. Adding one sentence asking the model to consider everyone affected, with none of the scaffold's other structure, to a standard chain-of-thought prompt does not lower validation on the same 150 items. It raises it, by 6.1 points pooled across three models, a confidence interval excluding zero in the positive direction,

an outcome the registration did not anticipate and has no pre-declared branch for. A style-matched control that asks the model to check the premises instead shows no comparable pattern pooled, though it produces a sharp, unexplained drop on one model that is reported and not interpreted. Counterparty-awareness on its own reads as an invitation toward a broader, more accommodating response, not as a check against one. The item-level association (the reduction concentrates on accounts that already contain a counterparty) and the causal, instruction-level claim (telling the model to consider one produces the same reduction) are different claims, and only the first is supported by the record.

The effect is on the register of the answer and not on its verdict. On the mechanically scored dispute instrument the asker's own framing moves the verdict toward the asker under both the scaffold and the control on two of the three open-weight models and on neither under the third, and the binary instrument returned a clean null on the original models. Where the asker's pull reaches the verdict itself, the scaffold does not remove it. Where a premise is false and no second party exists, the scaffold affirms the falsehood more often on a forced verdict, 22.3 against 15.7 percent on a first pass, and lowers false affirmation elsewhere by abstaining more rather than by rejecting more. These are the boundaries the theory owns.

4.2 Pillar 2, theory of mind (out of scope, skeleton for future work)

[TODO: Not a claim of this paper. Stage-1 data exists (Addendum 18 in the prereg, four models, 320 ToMBench items, guard-clean): pooled delta -0.044 [-0.066, -0.022], a cost concentrated on intent-inference tasks (-0.075) and absent on factual state-tracking (+0.000), traced to the scaffold's decision-commitment clause. Left here for a future developer to pick up or discard; stage 2 (haiku, sonnet, Llama) was never run.]

4.3 Pillar 3, the collective

Five instruments carry the results (Table 1). Intervals are 95 percent item-clustered bootstraps unless another level is named.

4.3.1 Embodiment places the bias

On the open-ended accounts (Cheng et al., 2025a) narration reduced validation of the asker's own account on every generator, by 26.7 points [16.7,

36.7] on haiku, 22.3 [13.3, 31.3] on sonnet, 41.4 [31.4, 51.1] on nano and 36.0 [27.1, 44.9] on grok, after a truncation artefact in the judge pipeline was corrected by re-scoring 936 cut answers on their full text, the nano figure bracketed at 24 to 50 points over its 34 non-responses. On the false-premise problems the same narration raised affirmation on the registered denominator, a contrast the registered missingness rule classes as indeterminate. The untreated model affirmed the false premise at 0.157 and under narration at 0.223, the registered contrast was +0.067 [+0.033, +0.103], the committed-only contrast, with the 34.5 and 27.8 percent of the 600 units per arm that returned no verdict removed, was +0.070 [+0.030, +0.115], and the Manski bounds on the registered contrast, [-0.278, +0.345], cross zero.

Assigning an interest without narration moved commitment far less. As the statement's author the model affirmed it at 0.192 and as the author's rival at 0.163, a contrast of +0.028 [+0.000, +0.058]. Combining the two did not add, a difference in differences of -0.025 [-0.070, +0.022]. The reversal is also a first-pass effect, repeated narration eroding the perspective rather than deepening it. Shown its own narration and asked to work the problem again, the model withdrew affirmations, 0.300 to 0.224 on 434 matched units, -0.076 [-0.104, -0.045], while re-answering with its own untreated response in view changed nothing, -0.003 [-0.025, +0.020], a difference in differences of -0.073 [-0.107, -0.039].

Advocacy fixed the conclusion (Table 3, Figure 2A). On grok the writer's advocate found the writer at fault on 0.3 percent of its opening statements and the other party's advocate on 97.4 percent, a role-lock of 0.971 [0.954, 0.986] on 1,666 debates, while the seat with no interest sat at 0.430. Lock is graded across models (Table 3), and removing the stake clause while keeping the directional assignment left grok's seats at 0.946 on 377 debates, so the pinning force was the assignment, not the stake.

4.3.2 The collective has one effective voter where the roles lock

A locked seat carried almost nothing about the case, the two advocates scoring 0.591 and 0.437 against the crowd on grok while the neutral seat scored 0.860, on the 1,642 debates with three codable opening verdicts (Figure 2B). So where the roles locked the collective had one effective voter (Figure 4B), the three-seat majority equalling the neu-

Table 3: Role-lock and seat accuracy at the opening statement by model of the collective. Lock is $|P(\text{at-fault})_{\text{writer}} - P(\text{at-fault})_{\text{counterparty}}|$. The stake-removed row keeps the directional assignment without the stake clause, and its seat accuracies were not computed. Seat accuracies are on each seat’s codable opening verdicts (1,666, 1,665 and 1,670 on grok), and the prose uses the 1,642 debates where all three are codable, so 0.857 here and 0.860 there are one quantity on two denominators.

model	n	$P(\text{at-fault})$ at opening			lock	$P(\text{at-fault})$ at opening		
		writer	counterparty	neutral		writer	counterparty	neutral
grok	1,666	0.003	0.974	0.430	0.971	0.592	0.437	0.857
grok, stake removed	377	0.008	0.954	0.387	0.946	n/a	n/a	n/a
sonnet	31	0.258	0.935	0.452	0.677	0.806	0.516	1.000
nano	816	0.147	0.767	0.363	0.620	0.718	0.612	0.830
haiku	190	0.474	0.915	0.656	0.442	0.816	0.444	0.708




Figure 2: Advocacy as control, and its cost. (A) Probability that each seat’s opening verdict is at-fault, four models, with the gold base rate marked. (B) Each seat’s verdict accuracy against its model’s role-lock. The writer’s advocate scores 0.816 at the lowest lock and 0.592 at the highest.

tral seat’s verdict on 1,642 of those 1,642 debates and the cross-validated best fixed function of the three opening verdicts being the neutral seat’s own, a gap of +0.000 [-0.004, +0.000] against its 0.860 [0.822, 0.896]. The count holds only where lock does. On nano at lock 0.620 the majority equalled the neutral seat on 777 of 781 debates, and on haiku at lock 0.442 the writer’s advocate outscored the neutral seat, 0.818 against 0.717 on 187 debates, +0.102 [+0.028, +0.182], the majority still tracking it on 180 of 187, so haiku is the boundary of the claim.

The two-by-two of roles and edges (Table 4, Figure 3) says which creates the structure. With the edges cut, role-lock, the localisation excess and the stake concentration G3 held, and with three identical readers on the same protocol they vanished, the registered reading H-ROLE, the structure coming from directional role assignment. The edges of the two-by-two are the exchange rounds, not the pairing of opposed verdicts the theory calls an

Table 4: The roles-by-edges two-by-two on grok, 210 items. Embodied seats are the two advocates and the neutral adjudicator, identical seats three copies of the neutral brief. Edges on means the seats read one another at the rebuttal rounds. Intervals are the two-by-two’s own item-clustered bootstrap with 3,000 draws, so the embodied, edges-on transfer interval differs in the third decimal from the 4,000-draw artefact quoted in the text. The majority row’s intervals are [+0.000, +0.000] in both embodied cells and [-0.019, +0.015] and [-0.020, +0.015] in the identical cells. Fire rate is the share of debates on which two or more seats objected. G3 is undefined without stakes, and the identical cells fired too rarely for a lift or transfer reading.

	embodied roles	
	edges on	edges off
debates	1,680	1,680
role-lock	0.971	0.973
localisation excess	+0.291	+0.307
fire rate, two or more object	0.196	0.146
debates fired	329	246
G3 stake concentration	+0.631	+0.642
seat majority at opening	0.860	0.854
best single seat	0.860	0.854
majority minus best seat	+0.000	+0.001
collective’s verdict	0.843	0.834
transfer lift on grok alone	+0.212 [+0.081, +0.358]	+0.183 [-0.035, +0.401]
within-one-loser counter lift, ratio	9.83x [+0.201, +0.493]	5.06x [+0.114, +0.895]
advocate phi within one-loser	-0.901	-0.842 [-0.876, -0.817]

edge, and that pairing sat on nodes that read no other seat before the synthesis they labelled, while three identical readers under the same moderator produced none of it. Cutting the exchange lowered the fire rate from 0.196 to 0.129, a drop of 0.067 against the registered minimum detectable 0.045, so the exchange amplified the sensor without creating it. No seat majority beat the best single seat in any cell, with or without roles, and what lock adds is that the majority is the neutral seat on every debate. The interests also cost accuracy. On 415 item-matched cells, identical readers with no edges scored 0.872 against 0.842 for the embodied collective, +0.030 [+0.012, +0.048], and every collective stayed below grok answering alone at 0.911 on the same matched cells.

4.3.3 The dissent is an error sensor on one-loser verdicts

A one-loser verdict undermines exactly one advocate, so on grok exactly one seat objected on 0.792 of debates against 0.501 under an independence null from the seats’ own objection marginals, a localisation excess of +0.291 [+0.255, +0.328] (Figure 4A). The registered composite flag fired on 337




Figure 3: The roles by edges two-by-two on grok. Role-lock, the fire rate of the counter, the localisation excess and the collective's verdict accuracy per cell. The structure follows the roles and survives the removal of every exchange. The identical-reader cells produce none of it and are the most accurate.




Figure 4: Structure of the dissent and the effective voter count on grok. (A) Number of seats objecting, observed against a Poisson-binomial null from the per-seat marginals. (B) Accuracy of the neutral seat, the three-seat majority and the cross-validated oracle over the three opening verdicts on the 1,642 debates with three codable opening verdicts, the full protocol's verdict on 1,677 codable debates, and grok alone at majority-of-three on 420 cells.

of 1,677 codable debates, lifting the probability of a wrong synthesis by +0.325 [+0.236, +0.416], and the stake-blind counter, two or more seats objecting, fired on 328, with integrated-verdict error rates of 0.415 against 0.095, +0.320 [+0.239, +0.404], a ratio of 4.37, the composite's advantage as a router being +0.001 [+0.000, +0.003], so the introduction's sensor is this counter.

The pooled ratio mixes two signals. Of the 328 firings, 93 sat on the 1,287 one-loser verdicts and 235 on the 390 both-party verdicts. Within one-loser verdicts the collective was wrong on 0.387 of flagged debates against 0.039 of unflagged, +0.348 [+0.201, +0.493], a ratio of 9.83, and within both-party verdicts on 0.426 against 0.523, -0.097 [-0.233, +0.043]. The dissent is the signal where the verdict has one loser and the verdict type where it does not. As a router to the haiku judge the verdict type alone composed at 0.894 against the counter's 0.880, +0.014 [-0.003, +0.030], and their union at 0.905, +0.011 [-0.001, +0.024] over the verdict type. With the cached sonnet judge the three routers composed at 0.913, 0.936 and 0.956 and the union's gain over the verdict type, +0.021 [+0.011, +0.033],

Table 5: The sensor by model, and its transfer. Top block, each model's own collective on the crowd panel, with G1 the composite-flag fire rate, G2 the share of final votes that reject, G3 the stake concentration and S_2 the collective's accuracy. Bottom block, the grok collective's counter applied to solo generators that never saw the debate, on the codable cells of 420 over 210 items, with the error rate on flagged against unflagged cells, the ratio and the lift.

model of the collective	debates	G1	G2	G3
grok	1,677	0.201	0.333	+0.631 [+0.587, +0.673]
haiku	192	0.458	0.069	+0.062 [+0.016, +0.108]
nano	828	0.854	0.024	+0.034 [+0.018, +0.051]
sonnet	31	0.065	0.000	0 (no rejects)
solo generator	cells	wrong, flagged vs not	ratio	
grok, standard	420	0.275 vs 0.062	4.40	
grok, narrated	420	0.353 vs 0.070	5.01	
nano	407	0.375 vs 0.159	2.36	
haiku	420	0.255 vs 0.136	1.88	

excluded zero. Whether the two add depends on the judge. Grip, the registered precondition, held on grok alone among four models (Table 5).

The counter also concentrated the errors of generators that never saw the debate, on three of four readouts (Table 5, Figure 5). On the 420 cells, grok answering alone at 0.912 was wrong 4.40 times as often where the counter flagged, +0.212 [+0.085, +0.359], and over the 1,287 and 390 debates that transfer lived in the one-loser stratum, +0.130 [+0.026, +0.258] against +0.017 [-0.087, +0.122]. Nano gave +0.216 [+0.064, +0.383] on 407 codable cells and haiku +0.119 [-0.010, +0.271], so with both grok readouts within-model, one of two cross-vendor readouts excludes zero, which is the claim's scope. The solo model's own free signals did as much, any both-party verdict among its three samples concentrating its errors 8.20 times and disagreement among them 6.35, against the counter's 4.40, each composing with the same judge within 0.005 of the counter's 0.914, no pairwise difference excluding zero. By the registered rule the transfer claim is that the collective's flag concentrates a solo model's errors, and so do its own signals.

Routing on the flag paid over the collective (Figure 6). The haiku judge on the 337 flagged debates raised their accuracy from 0.582 to 0.774, +0.193 [+0.067, +0.310] at the Bonferroni-corrected 98.75 percent level, and on the unflagged debates moved it -0.040 [-0.100, +0.016], so at 95 percent the composed system at 0.881 beat the collective by +0.039 [+0.018, +0.060] and the judge used everywhere,




Figure 5: Transfer of the counter. (A) Ratio of the error rate on flagged to unflagged cells of item and framing, and debates for the collective, for the collective itself and for four solo generators the flag never touched. (B) Item-clustered intervals and point estimates on the four transfer lifts.




Figure 6: Composed-system accuracy against the fraction of debates routed to the judge, for the counter, the verdict-type router and their union, with random routing at the union's coverage and the two solo grok baselines marked. Panel A routes to the haiku judge and panel B to the sonnet judge of the second design.

0.850, by +0.032 [-0.004, +0.068], a gain the record does not claim. A stronger judge met the first registered criterion and not the second. Routing grok's solo verdicts to sonnet on the counter raised 0.912 to 0.936, +0.024 [+0.005, +0.045], on the same 420 cells where the composite flag gives the same result, and sonnet on every cell scored 0.976, so routing on the flag sits 0.041 [0.014, 0.069] below the judge used everywhere. The same cells routed on the solo model's own three-sample disagreement reached 0.938, +0.026 [+0.010, +0.045], not distinguishable from the counter, +0.002 [-0.019, +0.024], at three calls per case against seventeen, and routing on any both-party verdict among the solo's samples reached 0.960. Priced from the cached token counts, the solo model alone costs \$0.0005 a case, the solo routed on its own disagreement \$0.0014, the collective's sensor routed to the same judge \$0.0098, and the judge everywhere \$0.0115, so for a solo model's errors the collective's flag is dominated on cost by the model's own disagreement, and its value is as an audit of the collective's own verdicts and as the structure in which the information sits.

4.3.4 Who decides changes the verdict, and who is seated changes it too

Four ablations vary the graph itself, holding the two locked advocates and their account fixed, to ask what the moderator and the seating do to the verdict and to the sensor separately.

Replacing the grok moderator with sonnet over byte-identical, locked seats raises the collective's verdict accuracy from 0.848 to 0.885, +0.038 [+0.014, +0.064]; replacing it with haiku instead lowers accuracy to 0.735, -0.115 [-0.169, -0.060], below haiku answering the same items alone (0.851). The seats' own behaviour does not move with the moderator. Role-lock is identical by construction (the opening statements are replayed), and localisation excess (+0.300 sonnet, +0.239 haiku, against +0.291 pooled) and the counter's fire rate stay within the range the moderator-free cells already show. What moves is the verdict itself and, through it, which verdicts the sensor sees. The haiku moderator issues an everyone-or-no-one verdict on more than a quarter of debates (0.300 against 0.243 for grok and 0.176 for sonnet), which floods the counter and collapses its pooled lift to +0.070 [-0.026, +0.170]; the within-one-loser lift, which the flooding does not touch, is the reading that survives. The moderator is therefore not a neutral aggregator of what the seats already decided. It is the node whose own competence sets the verdict and whose own habits about naming everyone at fault set how often the sensor is even asked to speak, while the seats' structure is invariant to it.

Replacing the three seats themselves with three uninterested readers from three vendors (haiku, nano, grok, no assigned side) tests whether the same-model bottleneck of Section 4.3.2 is a same-model problem or a same-opinion problem. Errors across vendors are far less correlated than errors across samples of one model (mean pairwise phi +0.253 against +0.65 pooled), and the majority does approach the independence prediction, but the panel's weakest reader (haiku, 0.697) is weak enough that the majority still does not clear its best member at the opening statement, +0.011 [-0.013, +0.034], and after one further round it does, +0.031 [+0.005, +0.057], the one cell in the programme where a three-seat majority beats its best single seat with an interval excluding zero. Neither reading changes the standing result: the cross-vendor panel is not more accurate than three

grok readers ($-0.010 [-0.033, +0.013]$ on the collective’s verdict) and not a better answerer than the best model alone ($-0.055 [-0.087, -0.024]$). The panel does produce an objection structure with no assigned stake at all (fire 0.401, lift $+0.121 [+0.044, +0.201]$), disagreement among unequal readers rather than dissent against interest, and it is weaker on every readout than the embodied sensor. Composed as a router ahead of a stronger judge it beats the collective alone ($+0.024 [+0.009, +0.042]$) but not a single cheap judge call of a stronger model, so a decider panel across vendors buys nothing a moderator selected for competence does not already buy.

Adding a second, like-biased advocate to the writer’s side (four seats: two for the writer, one for the counterparty, one uninterested) tests whether the sensor is about sides or about seats. The second writer-side seat copies its principal’s opening verdict on 99.5 percent of debates and its objections at $\phi = +0.975 [+0.952, +0.995]$, and adds no lift to any reconstruction of the counter once the extra seat is accounted for (pooled two-of-four lift $+0.294 [+0.206, +0.383]$ against the three-seat reconstruction’s $+0.324 [+0.224, +0.422]$, a difference of $-0.030 [-0.111, +0.053]$). The sensor scales with distinct sides, not with seats. The verdict does not: with two seats on one side the collective’s accuracy falls by $-0.074 [-0.117, -0.031]$ and its at-fault rate falls by $-0.121 [-0.154, -0.088]$, concentrated on the items where the writer is at fault (0.808 to 0.570 on gold-guilty items). An unequal graph manufactures bias at the decider even when every advocate is individually locked to its assigned side, which is the one finding in this batch that bears directly on deployment: representation in the graph is itself a lever on the verdict.

Extending the graph to a third party named in the account but not seated in the original three (accounts where the writer’s story implicates someone else’s decision) tests the credible-signal principle’s scope directly. The writer and counterparty stay locked (0.959); the added third-party advocate does not, arguing its assigned stake only 0.800 of the time, and its own against-interest objection rate lands well above the 0.10 the two-party principle predicts (0.198, against the writer’s 0.081). With three advocates a verdict undermines exactly one of them on 241 debates and two of them on 72; the against-interest count still works within the one-undermined case ($+0.162 [+0.051, +0.289]$, 69 fired, transfer to grok solo

$+0.078 [+0.001, +0.160]$) and a plain two-or-more counter floods once two advocates are undermined together, exactly as predicted (fire 0.353 to 0.861). What fails is the one-dimensional reading of the stake: at two-undermined debates the co-losers object together only 0.695 of the time, not near certainty, because a two-party verdict token does not encode every dimension along which a third party’s interest can be undermined. The sensor degrades gracefully where the graph grows along a new axis; the verdict again pays a cost for the extra seat ($-0.045 [-0.087, -0.006]$ against the three-seat cell on the same debates), the same direction 16.19 found.

4.3.5 The extent of the method

The sensor reads departures from role, so it fails in three measured ways. On the paired-account corpus, where the verdict names which of two actions was worse, grip was the strongest we measured, $+0.670 [+0.627, +0.710]$ on 624 codable debates. The advocates tracked their interests so well that the composite flag fired on 6 of 624 debates and both advocates rejected on none, and a registered search over eleven native dissent statistics promoted none. The sensor went silent.

Where the verdict finds everyone at fault or no one, both advocates lose and the sensor floods, firing on 0.603 of grok’s 390 both-party verdicts against 0.072 of its 1,287 one-loser verdicts, so there the verdict type is the error signal, not the dissent (Section 4.3.3).

Without a real counterparty the roles moved the vote little. On the false-premise problems a majority-of-three panel mixing author, neutral and rival affirmed the false premise at 0.162 against 0.188 for three authors, $-0.026 [-0.047, -0.007]$, and $-0.010 [-0.031, +0.008]$ against three disinterested judges. It rejected the premise no more often, $-0.001 [-0.025, +0.022]$, and abstained more, $+0.027 [+0.000, +0.054]$, so it lowered false affirmation by abstaining more, not by rejecting more, no gain in accuracy under the rule that non-commitment is never correct. Under narration the same contrast fell to $-0.004 [-0.025, +0.016]$ and the mixed panel itself rose by $+0.044 [+0.012, +0.079]$, so narration’s error passed through the vote.

The same theory marks a joint-action frontier, and every regime of its instrument fell outside the theory’s domain for a measured reason. On 144 ordinal two-by-two games under a binding joint plan, grok alone scored 0.891 [0.840, 0.935] and 48 of 48

in the value region, against the 0.215 floor with pay-offs masked. Seated for the two players the advocates barely locked, 0.094 [+0.031, +0.167] beside a neutral seat and 0.115 [+0.042, +0.198] beside a mediator, neither ever named the mutual-private cell, both collectives scored 48 of 48 and no seat objected, so the sensor was silent, not flooded. A plan both must accept removes the conflict the roles need. Over all 144 games the advocates locked at the opening statement, 0.405 [+0.335, +0.474] and 0.385 [+0.322, +0.449], then conceded by the second round, so the flag never fired on 288 cells, and the one gain over grok alone was restraint. Where no unique plan exists the solo model named one on 47 percent of cells, and the mediator collective declined more often, +0.151 [+0.043, +0.280] over 31 games, against the registered expectation that neither arm would beat the solo there, while the naive collective's +0.086 [-0.032, +0.226] did not clear zero. Without enforcement the roles locked completely, the tempted advocate naming its private action on every opening statement, and so did every other node, the solo model and the neutral reader recommending the equilibrium in every dilemma game and the mediator in 20 of 21 cells, none ever the Pareto cell, so the mediator held nothing, -0.042 [-0.146, +0.042] against the naive reader. In a one-shot dilemma without enforcement the correlated-equilibrium set is the Nash set, so the unanimity is the correct answer to the question as put. Where a recommendation can coordinate, in nine games with two equilibria favouring different players, the solo model declined on every one of 27 rows, which the registration scores as failure because either equilibrium is self-enforcing once both players hear it, and the collective coordinated on an equilibrium on 5 and 6 of 27, +0.185 [+0.037, +0.370] over the solo in the neutral arm and +0.222 in the mediator arm, whose own interval [0.000, 0.481] reaches zero. All eleven recommendations named an equilibrium, nine the welfare-better one, so the gain is commitment, not guessing, and the two third seats did not differ, +0.037 [-0.111, +0.222]. Across every regime the advocates locked at the opening and yielded by the vote, the opposite of the dispute and paired-account advocates, so the sensor had nothing to read and a non-embodied third node steered nothing the naive reader did not, and whether restraint and coordination owe anything to the opposed advocates rather than to deliberation itself is not separable here, since the solo comparator has no third seat and no identical-reader cell was run

on the games.

5 Related Work

Sycophancy is a trained, steerable bias, one of a family. Preference-tuned models repeat the answer the user prefers (Perez et al., 2023; Sharma et al., 2024), validate the asker's account of a conflict (Cheng et al., 2025a,b), abandon a correct answer under pushback (Fanous et al., 2025; Malmqvist, 2024; Hong et al., 2025; Nikeghbal et al., 2026), and it has surfaced in deployment (OpenAI, 2025). The pull has other doors, a false premise surviving into a task though the model registers the error (Chen et al., 2026b; Petrov et al., 2025), the answer folding under pressure while the reasoning holds (Li et al., 2026), and a persona shifting the answer, which a benchmark may record as deference when it is personalization (Batzner et al., 2025). The bias is an activation direction one can add or remove (Rimsky et al., 2024; Chen et al., 2025; Vennemeyer et al., 2025), a sceptical persona recovers most of what targeted steering achieves (Kelkar et al., 2026), and remedies exist in training (Wei et al., 2023; Chen et al., 2024a) and in a third-person perspective at inference (Hong et al., 2025). We read the three as one family, a conclusion fixed by the prompt, and place the pull on purpose, a copy told its side taking it on 99.7 and 97.4 percent of opening statements, while narration, first person with no side, cuts validation of the asker by 22 to 41 points yet leaves one model affirming a false premise more often, so assignment is a lever, not a cure.

Role assignment controls what a copy says, at a price. A persona prompt moves output in the direction it implies (Shanahan et al., 2023; Salewski et al., 2023; Li et al., 2023), within asymmetric limits (Miehling et al., 2025; Liu et al., 2024a; Santurkar et al., 2023), and moves accuracy either way (Zheng et al., 2024; Kim et al., 2025b; Kong et al., 2024). Assigned personas carry bias (Gupta et al., 2024b; Deshpande et al., 2023; Dash et al., 2026), agreeable ones amplify sycophancy (Shah et al., 2026), and simulated debaters drift toward the model's own view (Taubenfeld et al., 2024). Dietrich (2026) comes closest, with opposed advocates, a balanced seat and a supervisor, and finds 39 to 67 percent role fidelity collapsing on unambiguous facts. We measure the same quantity as role-lock, near total on one model (0.97) and soft on another (0.44), and unchanged with the seats' reads

cut or narration removed. A like-biased second advocate copies its principal yet pulls the moderator's verdict twelve points its way and costs seven points of accuracy, and a third party's advocate breaks the one-dimensional stake the flag needs. We add the control's price, that a locked seat's verdict carries nothing about the case and that seat count is a persuasion channel a naive moderator does not resist, so assignment is a lever paid for at the verdict.

Debate puts truth at a judge whose competence sets the verdict, and disagreement is a signal. Scalable oversight puts two opposed agents before a judge (Irving et al., 2018), who gains from debate (Michael et al., 2023; Khan et al., 2024; Rahman et al., 2025; Arnesen et al., 2024; Lang et al., 2025) only under information asymmetry (Kenton et al., 2024), only by verifying, not summarising (Elasky et al., 2026), in step with its capability (Engels et al., 2025), and not when it shares the speakers' errors (Goel et al., 2025). Debate among copies of one model helps some benchmarks (Du et al., 2023; Liang et al., 2024; Chan et al., 2024) and not others (Smit et al., 2024; Wang et al., 2024; Zhang et al., 2025), since copies converge to the majority in the room (Estornell and Liu, 2024; Zhu et al., 2026; Choi et al., 2025), mixed panels yield to a peer's flawed reasoning even when stronger models are the majority (Wynn et al., 2025), opposed positions invite cheap talk (Chen et al., 2026a), and instructed dissent, like instructed confidence, is not evidence (Qian, 2026; Prasad and Nguyen, 2025). A second lineage treats disagreement itself, resolving it (Jiang et al., 2026) or reading it through cross-examination (Cohen et al., 2023), a verifier's check (Kirchner et al., 2024; Brown-Cohen et al., 2024, 2025; Hammond and Adam-Day, 2025) or variation across one model's samples (Wang et al., 2023; Manakul et al., 2023; Farquhar et al., 2024; Xiong et al., 2024), and such signals decide when to overturn a majority (He et al., 2026) or escalate to a stronger model or human reviewer (Chen et al., 2023; Aggarwal et al., 2024; Gupta et al., 2024a; Yue et al., 2024; Jung et al., 2025; Wang et al., 2026; Jitkritum et al., 2023). Our seats hold opposed interests, not opposed answers, so neither holds the truth by construction, and where a verdict names one loser exactly one seat has reason to object, so no brief explains a second objector. We add that flag, needing no labels, logits or second sample, which marks cases wrong 9.8 times as often on one-loser verdicts, survives cutting every read

between seats and cutting narration, firing a third less often, and routes the flagged fifth to a stronger uninterested judge for a gain over the collective, though the sonnet judge everywhere does better and the solo's own sample disagreement does as much.

Topology decides what a collective aggregates, and same-model seats aggregate to one. Debate graphs can be sparse (Li et al., 2024), scaled and optimised (Qian et al., 2025; Zhuge et al., 2024; Liu et al., 2024b), topology governs error propagation (Shen et al., 2025), centralised aggregation is hostage to its hub (Han et al., 2026), agents conform to a majority (Zhang et al., 2024; Weng et al., 2025; Bertalaníć and Fortuna, 2026a), and effective team size is fittable (Bertalaníć and Fortuna, 2026b). Correlated error bounds aggregation and rises as models improve (Kim et al., 2025a; Goel et al., 2025). Mixed panels beat same-model panels only sometimes (Wang et al., 2025; Li et al., 2025), a majority beats its best member in a tenth of three-model panels (Kim, 2026), and cross-model gains need learned weights, not a vote (Douven, 2026). Our three same-model seats never beat their best seat, locked or not, and a three-vendor panel with less correlated error (0.25 against 0.65) does not beat its best member at the opening statement and sits below the best model alone. Over byte-identical locked seats a stronger moderator moves the verdict up four points, a weaker one down twelve. We add the structure of dissent as the outcome and the reading that the moderator's competence sets the accuracy while the seats supply the dissent, so the moderator should be one strong model.

Jury theorems, network learning and advocacy theory supplied the design and predict the measurements. Correlated votes can help or hurt a majority (Ladha, 1992; Boland, 1989; Dietrich and Spiekermann, 2013), the optimal rule over unequal voters weights by log-odds (Nitzan and Paroush, 1982; Shapley and Grofman, 1984), a shared bias needs a non-majoritarian rule (Prelec et al., 2017), and a network of averagers reaches the truth only when no node dominates (DeGroot, 1974; Golub and Jackson, 2010). Advocacy theory supplies the seats, as a naive decider can rely on competition among opposed parties (Milgrom and Roberts, 1986; Froeb and Kobayashi, 1996), advocates beat a nonpartisan investigator (Dewatripont and Tirole, 1999), conflict guarantees a mistaken

inference is corrected (Lipman and Seppi, 1995), an adversarial procedure allocates the burden of proof (Shin, 1998), two experts help only with opposed biases (Krishna and Morgan, 2001), adding senders need not add information (Gentzkow and Kamenica, 2017), a large bias babbles (Crawford and Sobel, 1982), senders must span the case’s dimensions (Battaglini, 2002), and an objection is read relative to what it counters (Glazer and Rubinstein, 2001, 2004). Across fields a statement against interest is credible since the speaker had reason not to make it (Federal Rules of Evidence, 2024; Jefferson, 1944), given what the receiver infers of the speaker’s motives (Sobel, 1985) and costly fabrication (Sanchirico, 2001). We supply an instance where these predictions split by layer, the opposed-expert gain arriving as dissent, not votes, the optimal rule over the opening verdicts reducing under lock to the uninterested seat’s own, the like-biased second expert copying its principal yet moving the verdict, and the flag working only where opposed interests, an allocable burden and a spanned dimension coincide, on one-loser verdicts, flooding or falling silent elsewhere.

Argumentative reasoning and social objectivity locate truth in the structure of criticism. On the argumentative account reasoning evolved for argument, individuals produce biased advocacy and truth emerges in evaluation (Mercier and Sperber, 2011, 2017; Trouche et al., 2014, 2016), and more conflict between opinions should improve outcomes as evaluation prevails over production (Mercier, 2016). Objectivity belongs to a community’s criticism and its uptake, not to any member (Longino, 1990, 2002; Popper, 2011). Assigned devil’s advocacy beats consensus (Schweiger et al., 1986; Schwenk, 1990), yet a role-played devil’s advocate mainly provokes bolstering, unlike authentic dissent (Nemeth et al., 2001). We build the structure from copies of one model and find the production half reproduced and the uptake half all but absent, since the seats never converge and a collective fed its own dissent does not move, so truth, where the collective was wrong, entered from outside the graph, and our objectivity is a rule for reading dissent, the objection no role explains carrying the information.

Cooperative AI asks what a third node can do, and correlated equilibrium sets the limit. Cooperative AI asks how agents with divergent interests reach joint welfare (Dafoe et al., 2020), and

in repeated dilemmas language models play with model-specific dispositions, several leaning cooperative (Akata et al., 2025; Fontana et al., 2025; Willis et al., 2025; Zheng et al., 2025), fail a commons without foresight (Piatti et al., 2024), and shift with an assigned persona (Leon et al., 2026; Jia et al., 2025). Mediation, contracting and one-word talk sustain cooperation (Tewolde et al., 2026; Wyse et al., 2026; Madmoun and Lahlou, 2026), mediator design moves consensus (Liu et al., 2026), and negotiation among role-assigned stakeholders is scorable (Abdelnabi et al., 2024). The formal basis is Aumann’s mediator, whose correlated recommendation reaches outcomes outside the Nash set (Aumann, 1974), though a direct mechanism with equal commitment achieves the same (Myerson, 1986). Our advocates lock at the opening statement wherever interests diverge and yield at the vote on every payoff table, leaving the flag nothing to read. Our games found what that theory predicts and no more, since without enforcement every node and the solo recommend the equilibrium in a one-shot prisoner’s dilemma, whose correlated-equilibrium set is its Nash set, and under a binding plan the model needs no help, so no seating made cooperation self-enforcing, while with two equilibria the collective coordinates where the solo declines, the mediator adding nothing over a naive reader, and where no plan is unique it adds, given the mediator, an unreplicated restraint.

Theory-of-mind benchmarks measure the representation of other minds that narration is asked to express. [TODO: ToM strand. ToMBench (Chen et al., 2024b) and the sources named in the origins review (lines 193 to 202), each fetched at a primary page before citing]

6 Limitations

The pillar-one outcome is judge-scored. Agreement among three judges on the validation construct is moderate, each judge agrees with human gold labels at a kappa of 0.73 to 0.86 on the 29 gold items, and a per-judge threshold absorbs most of the disagreement about levels while a judge-by-arm interaction remains. The instrument has no stance manipulation, so it measures the register of one response and not a change of verdict under pressure, and the mechanically scored instrument that does manipulate stance shows no removal of the asker’s pull. The scaffold’s mechanism at the level of compliance was retracted after a missing-

data analysis and is not claimed. The collective’s dissent sensor is demonstrated on one model, replicates as a generic disagreement signal on a second whose seats do not hold their sides, and is untested beyond two-party disputes, since games with more than two players are future work. Every cost figure is a list price at the time of the run.

References

Sahar Abdelnabi, Amr Goma, Sarath Sivaprasad, Lea Schönherr, and Mario Fritz. 2024. [Cooperation, competition, and maliciousness: LLM-stakeholders interactive negotiation](#). In *Advances in Neural Information Processing Systems 37 (NeurIPS 2024), Datasets and Benchmarks Track*.

Pranjal Aggarwal, Aman Madaan, Ankit Anand, Srividya Pranavi Potharaju, Swaroop Mishra, Pei Zhou, Aditya Gupta, Dheeraj Rajagopal, Karthik Kappaganthu, Yiming Yang, Shyam Upadhyay, Manaal Faruqui, and Mausam. 2024. [AutoMix: Automatically mixing language models](#). In *Advances in Neural Information Processing Systems 37 (NeurIPS 2024)*.

Elif Akata, Lion Schulz, Julian Coda-Forno, Seong Joon Oh, Matthias Bethge, and Eric Schulz. 2025. [Playing repeated games with large language models](#). *Nature Human Behaviour*, 9:1380–1390.

Anthropic. 2025. [Claude Haiku 4.5](#).

Anthropic. 2026. [Claude Sonnet 4.6](#).

Samuel Arnesen, David Rein, and Julian Michael. 2024. [Training language models to win debates with self-play improves judge accuracy](#). *Preprint*, arXiv:2409.16636. ArXiv preprint arXiv:2409.16636, submitted 25 September 2024.

Robert J. Aumann. 1974. [Subjectivity and correlation in randomized strategies](#). *Journal of Mathematical Economics*, 1(1):67–96.

Marco Battaglini. 2002. [Multiple referrals and multi-dimensional cheap talk](#). *Econometrica*, 70(4):1379–1401.

Jan Batzner, Volker Stocker, Stefan Schmid, and Gjergji Kasneci. 2025. [Sycophancy claims about language models: The missing human-in-the-loop](#). *Preprint*, arXiv:2512.00656. NeurIPS 2025 Workshop on LLM Evaluation and ICLR 2025 Workshop on Bi-Directional Human-AI Alignment.

Blaž Bertalaníč and Carolina Fortuna. 2026a. [The cost of consensus: Isolated self-correction prevails over unguided homogeneous multi-agent debate](#). *Preprint*, arXiv:2605.00914.

Blaž Bertalaníč and Carolina Fortuna. 2026b. [The ringelmann effect in multi-agent LLM systems: A scaling law for effective team size](#). *Preprint*, arXiv:2606.02646.

Philip J. Bolland. 1989. [Majority systems and the Condorcet jury theorem](#). *Journal of the Royal Statistical Society. Series D (The Statistician)*, 38(3):181–189.

Jonah Brown-Cohen, Geoffrey Irving, and Georgios Piliouras. 2024. [Scalable AI safety via doubly-efficient debate](#). In *Proceedings of the 41st International Conference on Machine Learning*, volume 235 of *PMLR*, pages 4585–4602.

Jonah Brown-Cohen, Geoffrey Irving, Georgios Piliouras, Lijie Chen, Jiawei Li, and Zhiyang Xun. 2025. [Avoiding obfuscation with prover-estimator debate](#). *Preprint*, arXiv:2506.13609. ArXiv preprint arXiv:2506.13609 (v1 16 June 2025, v2 21 July 2026).

Chi-Min Chan, Weize Chen, Yusheng Su, Jianxuan Yu, Wei Xue, Shanghang Zhang, Jie Fu, and Zhiyuan Liu. 2024. [ChatEval: Towards better LLM-based evaluators through multi-agent debate](#). In *The Twelfth International Conference on Learning Representations (ICLR 2024)*.

Lingjiao Chen, Matei Zaharia, and James Zou. 2023. [FrugalGPT: How to use large language models while reducing cost and improving performance](#). *Preprint*, arXiv:2305.05176. ArXiv preprint; a TMLR 2024 publication is reported by secondary indexes but its OpenReview page (id cSimKw5p6R) sits behind a human-verification check and was not read, so the arXiv version is the one cited.

Runjin Chen, Andy Arditi, Henry Sleight, Owain Evans, and Jack Lindsey. 2025. [Persona vectors: Monitoring and controlling character traits in language models](#). *Preprint*, arXiv:2507.21509.

Wei Chen, Zhen Huang, Liang Xie, Binbin Lin, Houqiang Li, Le Lu, Xinmei Tian, Deng Cai, Yonggang Zhang, Wenxiao Wang, Xu Shen, and Jieping Ye. 2024a. [From yes-men to truth-tellers: Addressing sycophancy in large language models with pinpoint tuning](#). In *Proceedings of the 41st International Conference on Machine Learning*, volume 235 of *Proceedings of Machine Learning Research*, pages 6950–6972.

Yongqiang Chen, Gang Niu, James Cheng, Bo Han, and Masashi Sugiyama. 2026a. [When and why does multi-agent debate fail and does it really underperform?](#) *Preprint*, arXiv:2510.20963. V2, July 2026; v1 (October 2025) was titled “Towards Scalable Oversight with Collaborative Multi-Agent Debate in Error Detection”.

Zhuang Chen, Jincenzi Wu, Jinfeng Zhou, Bosi Wen, Guanqun Bi, Gongyao Jiang, Yaru Cao, Mengting Hu, Yunghwei Lai, Zexuan Xiong, and Minlie Huang. 2024b. [ToMBench: Benchmarking theory of mind in large language models](#). In *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*.

Zixuan Chen, Hao Lin, Zizhe Chen, Yizhou Tian, Garry Yang, Depeng Wang, Ya Guo, Huijia Zhu, and James

1539	Cheng. 2026b. Knowing but not correcting: Routine task requests suppress factual correction in LLMs. <i>Preprint</i> , arXiv:2605.05957.	1593
1540		1594
1541		1595
1542	Myra Cheng, Esin Durmus, Miranda Zhang, Tomasz Korbak, Ethan Perez, Samuel R. Bowman, and 1 others. 2025a. ELEPHANT: Measuring and understanding social sycophancy in LLMs. <i>arXiv preprint arXiv:2505.13995</i> .	1596
1543		1597
1544		1598
1545		
1546		
1547	Myra Cheng, Cinoo Lee, Pranav Khadpe, Sunny Yu, Dyllan Han, and Dan Jurafsky. 2025b. Sycophantic AI decreases prosocial intentions and promotes dependence. <i>Preprint</i> , arXiv:2510.01395. ArXiv preprint arXiv:2510.01395 (v1, 1 October 2025). A journal version in Science (DOI 10.1126/science.aec8352) is reported by Crossref, but its landing page could not be read, so the arXiv version is the one cited.	1599
1548		1600
1549		1601
1550		1602
1551		1603
1552		
1553		
1554		
1555		
1556	Hyeong Kyu Choi, Xiaojin Zhu, and Sharon Li. 2025. Debate or vote: Which yields better decisions in multi-agent large language models? In <i>Advances in Neural Information Processing Systems 38 (NeurIPS 2025)</i> .	1604
1557		1605
1558		1606
1559		1607
1560		
1561	Roi Cohen, May Hamri, Mor Geva, and Amir Globerson. 2023. LM vs LM: Detecting factual errors via cross examination. In <i>Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing</i> , pages 12621–12640.	1608
1562		1609
1563		1610
1564		1611
1565		1612
1566	Vincent P. Crawford and Joel Sobel. 1982. Strategic information transmission. <i>Econometrica</i> , 50(6):1431–1451.	1613
1567		1614
1568		
1569	Allan Dafoe, Edward Hughes, Yoram Bachrach, Tantum Collins, Kevin R. McKee, Joel Z. Leibo, Kate Larson, and Thore Graepel. 2020. Open problems in cooperative AI. <i>Preprint</i> , arXiv:2012.08630.	1615
1570		1616
1571		1617
1572		1618
1573	Saloni Dash, Amélie Reymond, Emma Spiro, and Aylin Caliskan. 2026. Persona-assigned large language models exhibit human-like motivated reasoning. In <i>Findings of the Association for Computational Linguistics: ACL 2026</i> , pages 12043–12069.	1619
1574		1620
1575		1621
1576		1622
1577		
1578	Morris H. DeGroot. 1974. Reaching a consensus. <i>Journal of the American Statistical Association</i> , 69(345):118–121.	1623
1579		1624
1580		1625
1581	Ameet Deshpande, Vishvak Murahari, Tanmay Rajpurohit, Ashwin Kalyan, and Karthik Narasimhan. 2023. Toxicity in ChatGPT: Analyzing persona-assigned language models. In <i>Findings of the Association for Computational Linguistics: EMNLP 2023</i> , pages 1236–1270.	1626
1582		
1583		
1584		
1585		
1586		
1587	Mathias Dewatripont and Jean Tirole. 1999. Advocates. <i>Journal of Political Economy</i> , 107(1):1–39.	1627
1588		1628
1589	Franz Dietrich and Kai Spiekermann. 2013. Independent opinions? On the causal foundations of belief formation and jury theorems. <i>Mind</i> , 122(487):655–685.	1629
1590		1630
1591		1631
1592		1632
	Juergen Dietrich. 2026. When roles fail: Epistemic constraints on advocate role fidelity in LLM-based political statement analysis. <i>Preprint</i> , arXiv:2604.27228.	1633
		1634
		1635
	Igor Douven. 2026. Wisdom of LLM crowds: Aggregation and contamination in language model ensembles. <i>Preprint</i> , arXiv:2607.18269.	1636
		1637
	Yilun Du, Shuang Li, Antonio Torralba, Joshua B. Tenenbaum, and Igor Mordatch. 2023. Improving Factuality and Reasoning in Language Models through Multiagent Debate. <i>arXiv preprint arXiv:2305.14325</i> .	1638
		1639
		1640
		1641
	Ethan Elasky, Frank Nakasako, and Naman Goyal. 2026. Debate helps weak judges reward stronger models. <i>Preprint</i> , arXiv:2605.27483. ArXiv preprint arXiv:2605.27483, submitted 26 May 2026.	1642
		1643
		1644
	Joshua Engels, David D. Baek, Subhash Kantamneni, and Max Tegmark. 2025. Scaling laws for scalable oversight. In <i>Advances in Neural Information Processing Systems 38 (NeurIPS 2025)</i> . Spotlight poster, confirmed at https://neurips.cc/virtual/2025/poster/115536 ; arXiv:2504.18530.	1645
		1646
		1647
	Andrew Estornell and Yang Liu. 2024. Multi-LLM debate: Framework, principals, and interventions. In <i>Advances in Neural Information Processing Systems 37 (NeurIPS 2024)</i> .	
	Ahmed Fanous and 1 others. 2025. SycEval: Evaluating sycophancy in large language models. <i>arXiv preprint</i> . Graded false-premise probes; reports 29–58% sycophancy on frontier models.	
	Sebastian Farquhar, Jannik Kossen, Lorenz Kuhn, and Yarin Gal. 2024. Detecting hallucinations in large language models using semantic entropy. <i>Nature</i> , 630(8017):625–630.	
	Federal Rules of Evidence. 2024. Federal rules of evidence, rule 804(b)(3): Statement against interest. Legal Information Institute, Cornell Law School. Rule 804, Hearsay Exceptions; Declarant Unavailable, subsection (b)(3), with the Advisory Committee Notes; as amended through December 1, 2024.	
	Nicolò Fontana, Francesco Pierri, and Luca Maria Aiello. 2025. Nicer than humans: How do large language models behave in the prisoner’s dilemma? <i>Proceedings of the International AAAI Conference on Web and Social Media</i> , 19(1):522–535.	
	Luke M. Froeb and Bruce H. Kobayashi. 1996. Naive, biased, yet Bayesian: Can juries interpret selectively produced evidence? <i>The Journal of Law, Economics, and Organization</i> , 12(1):257–276.	
	Matthew Gentzkow and Emir Kamenica. 2017. Competition in persuasion. <i>The Review of Economic Studies</i> , 84(1):300–322.	
	Jacob Glazer and Ariel Rubinstein. 2001. Debates and decisions: On a rationale of argumentation rules. <i>Games and Economic Behavior</i> , 36(2):158–173.	

- Jacob Glazer and Ariel Rubinstein. 2004. [On optimal rules of persuasion](#). *Econometrica*, 72(6):1715–1736.
- Shashwat Goel, Joschka Strüder, Ilze Amanda Auzina, Karuna K Chandra, Ponnurangam Kumaraguru, Douwe Kiela, Ameya Prabhu, Matthias Bethge, and Jonas Geiping. 2025. [Great models think alike and this undermines AI oversight](#). In *Proceedings of the 42nd International Conference on Machine Learning*, volume 267 of *PMLR*, pages 19621–19678.
- Benjamin Golub and Matthew O. Jackson. 2010. [Naïve learning in social networks and the wisdom of crowds](#). *American Economic Journal: Microeconomics*, 2(1):112–149.
- Neha Gupta, Harikrishna Narasimhan, Wittawat Jitkrittum, Ankit Singh Rawat, Aditya Krishna Menon, and Sanjiv Kumar. 2024a. [Language model cascades: Token-level uncertainty and beyond](#). In *The Twelfth International Conference on Learning Representations (ICLR 2024)*.
- Shashank Gupta, Vaishnavi Shrivastava, Ameet Deshpande, Ashwin Kalyan, Peter Clark, Ashish Sabharwal, and Tushar Khot. 2024b. [Bias runs deep: Implicit reasoning biases in persona-assigned LLMs](#). In *The Twelfth International Conference on Learning Representations (ICLR 2024)*.
- Lewis Hammond and Sam Adam-Day. 2025. [Neural interactive proofs](#). In *International Conference on Learning Representations (ICLR 2025)*. ICLR 2025 poster, confirmed at <https://iclr.cc/virtual/2025/poster/29664>; arXiv:2412.08897.
- Chen Han, Jin Tan, Bohan Yu, Wenzhen Zheng, and Xijin Tang. 2026. [Conformity dynamics in LLM multi-agent systems: The roles of topology and self-social weighting](#). *Preprint*, arXiv:2601.05606.
- Chuan He, Zebin Chen, Zhengyi Yang, Shaobo Qiao, Mingchen Ju, Jiate Liu, Dong Wen, and Guanfeng Liu. 2026. [Minority sentinel: When to overturn majority voting in multi-agent LLM debates](#). *Preprint*, arXiv:2606.29270. AgentSearch Workshop at SIGIR 2026.
- Jiseung Hong, Grace Byun, Seungone Kim, and Kai Shu. 2025. [Measuring sycophancy of language models in multi-turn dialogues](#). In *Findings of the Association for Computational Linguistics: EMNLP 2025*, pages 2239–2259.
- Geoffrey Irving, Paul Christiano, and Dario Amodei. 2018. [AI Safety via Debate](#). *arXiv preprint arXiv:1805.00899*.
- Bernard S. Jefferson. 1944. [Declarations against interest: An exception to the hearsay rule](#). *Harvard Law Review*, 58(1):1–69.
- Jingru Jia, Zehua Yuan, Junhao Pan, Paul E. McNamee, and Deming Chen. 2025. [LLM strategic reasoning: Agentic study through behavioral game theory](#). In *Advances in Neural Information Processing Systems 38 (NeurIPS 2025)*. NeurIPS 2025 poster; arXiv:2502.20432.
- Yuyang Jiang, Chacha Chen, Teng Wu, Liwen Sun, Han Liu, Shi Feng, and Chenhao Tan. 2026. [Collaborative disagreement resolution for scalable oversight](#). *Preprint*, arXiv:2607.01251. ArXiv preprint arXiv:2607.01251; arXiv comments field states “Accepted to ICML 2026” (not verified on a PMLR page).
- Wittawat Jitkrittum, Neha Gupta, Aditya K. Menon, Harikrishna Narasimhan, Ankit Rawat, and Sanjiv Kumar. 2023. [When does confidence-based cascade deferral suffice?](#) In *Advances in Neural Information Processing Systems 36 (NeurIPS 2023)*.
- Jaehun Jung, Faeze Brahman, and Yejin Choi. 2025. [Trust or escalate: LLM judges with provable guarantees for human agreement](#). In *The Thirteenth International Conference on Learning Representations (ICLR 2025)*.
- Ishaan Kelkar, Nebras Alam, Vikram Kakaria, Madhur Panwar, Vasu Sharma, and Maheep Chaudhary. 2026. [Playing devil’s advocate: Off-the-shelf persona vectors rival targeted steering for sycophancy](#). *Preprint*, arXiv:2605.21006. ICML 2026 Pluralistic Alignment Workshop.
- Zachary Kenton, Noah Y. Siegel, János Kramár, Jonah Brown-Cohen, Samuel Albanie, Jannis Bulian, Rishabh Agarwal, David Lindner, Yunhao Tang, Noah D. Goodman, and Rohin Shah. 2024. [On scalable oversight with weak LLMs judging strong LLMs](#). In *Advances in Neural Information Processing Systems 37 (NeurIPS 2024)*.
- Akbir Khan, John Hughes, Dan Valentine, Laura Ruis, Kshitij Sachan, Ansh Radhakrishnan, Edward Grefenstette, Samuel R. Bowman, Tim Rocktäschel, and Ethan Perez. 2024. [Debating with more persuasive LLMs leads to more truthful answers](#). In *Proceedings of the 41st International Conference on Machine Learning*, volume 235 of *PMLR*, pages 23662–23733.
- Donghwan Kim. 2026. [Are diversity metrics measuring diversity? a capability-controlled audit of majority-vote gain in LLM ensembles](#). *Preprint*, arXiv:2607.20768.
- Elliot Myunghoon Kim, Avi Garg, Kenny Peng, and Nikhil Garg. 2025a. [Correlated errors in large language models](#). In *Proceedings of the 42nd International Conference on Machine Learning*, volume 267 of *PMLR*, pages 30038–30066.
- Junseok Kim, Nakyeong Yang, and Kyomin Jung. 2025b. [Persona is a double-edged sword: Rethinking the impact of role-play prompts in zero-shot reasoning tasks](#). In *Proceedings of the 14th International*

1758			1814
1759			1815
1760			1816
1761			
1762	Jan Hendrik Kirchner, Yining Chen, Harri Edwards, Jan		1817
1763	Leike, Nat McAleese, and Yuri Burda. 2024. Prover-		1818
1764	verifier games improve legibility of LLM outputs.		1819
1765	<i>Preprint</i> , arXiv:2407.13692.		1820
1766	Aobo Kong, Shiwan Zhao, Hao Chen, Qicheng Li, Yong		1821
1767	Qin, Ruiqi Sun, Xin Zhou, Enzhi Wang, and Xiaohang		1822
1768	Dong. 2024. Better zero-shot reasoning with		1823
1769	role-play prompting. In <i>Proceedings of the 2024</i>		1824
1770	<i>Conference of the North American Chapter of the</i>		1825
1771	<i>Association for Computational Linguistics: Human</i>		
1772	<i>Language Technologies (Volume 1: Long Papers)</i> ,		
1773	pages 4099–4113.		
1774	Vijay Krishna and John Morgan. 2001. A model of		
1775	expertise. <i>The Quarterly Journal of Economics</i> ,		
1776	116(2):747–775.		
1777	Krishna K. Ladha. 1992. The Condorcet jury theorem,		
1778	free speech, and correlated votes. <i>American Journal</i>		
1779	<i>of Political Science</i> , 36(3):617–634.		
1780	Hao Lang, Fei Huang, and Yongbin Li. 2025. Debate		
1781	helps weak-to-strong generalization. <i>Proceedings</i>		
1782	<i>of the AAAI Conference on Artificial Intelligence</i> ,		
1783	39(26):27410–27418.		
1784	Carlos Leon, Alexandre Rodrigues, Pedro Gamito,		
1785	and Thomas D. Parsons. 2026. How personas can		
1786	influence agents to play split or steal. <i>Preprint</i> ,		
1787	arXiv:2607.05398.		
1788	Guohao Li, Hasan Hammoud, Hani Itani, Dmitrii		
1789	Khizbullin, and Bernard Ghanem. 2023. CAMEL:		
1790	Communicative agents for “mind” exploration of		
1791	large language model society. In <i>Advances in Neural</i>		
1792	<i>Information Processing Systems 36 (NeurIPS 2023)</i> .		
1793	Wenzhe Li, Yong Lin, Mengzhou Xia, and Chi Jin.		
1794	2025. Rethinking mixture-of-agents: Is mixing dif-		
1795	ferent large language models beneficial? <i>Preprint</i> ,		
1796	arXiv:2502.00674.		
1797	Yubo Li, Ramayya Krishnan, and Rema Padman. 2026.		
1798	The chain holds, the answer folds: Trace-answer		
1799	dissociation in reasoning models under adversarial		
1800	pressure. <i>Preprint</i> , arXiv:2605.29087.		
1801	Yunxuan Li, Yibing Du, Jiageng Zhang, Le Hou, Peter		
1802	Grabowski, Yeqing Li, and Eugene Ie. 2024. Improv-		
1803	ing multi-agent debate with sparse communication		
1804	topology. In <i>Findings of the Association for Com-</i>		
1805	<i>putational Linguistics: EMNLP 2024</i> , pages 7281–		
1806	7294.		
1807	Tian Liang, Zhiwei He, Wenxiang Jiao, Xing Wang,		
1808	Yan Wang, Rui Wang, Yujiu Yang, Shuming Shi, and		
1809	Zhaopeng Tu. 2024. Encouraging divergent thinking		
1810	in large language models through multi-agent debate.		
1811	In <i>Proceedings of the 2024 Conference on Empiri-</i>		
1812	<i>cal Methods in Natural Language Processing</i> , pages		
1813	17889–17904.		
	Barton L. Lipman and Duane J. Seppi. 1995. Robust		
	inference in communication games with partial prov-		
	ability. <i>Journal of Economic Theory</i> , 66(2):370–405.		
	Andy Liu, Mona Diab, and Daniel Fried. 2024a. Evalu-		
	ating large language model biases in persona-steered		
	generation. In <i>Findings of the Association for Com-</i>		
	<i>putational Linguistics: ACL 2024</i> , pages 9832–9850.		
	Zijun Liu, Yanzhe Zhang, Peng Li, Yang Liu, and Diyi		
	Yang. 2024b. A dynamic LLM-powered agent net-		
	work for task-oriented agent collaboration. <i>Preprint</i> ,		
	arXiv:2310.02170. Conference on Language Model-		
	ing (COLM) 2024.		
	Ziyi Liu, Bahareh Sarrafzadeh, Pei Zhou, Longqi Yang,		
	Jieyu Zhao, and Ashish Sharma. 2026. ProMediate:		
	A simulation testbed for evaluating proactive media-		
	tion in multi-party negotiation. In <i>Findings of the As-</i>		
	<i>sociation for Computational Linguistics: ACL 2026</i> ,		
	pages 29570–29598. Association for Computational		
	Linguistics.		
	Helen E. Longino. 1990. <i>Science as Social Knowledge:</i>		
	<i>Values and Objectivity in Scientific Inquiry.</i> Prince-		
	ton University Press.		
	Helen E. Longino. 2002. <i>The Fate of Knowledge.</i>		
	Princeton University Press.		
	Hachem Madmoun and Salem Lahlou. 2026. Com-		
	munication enables cooperation in LLM agents: A		
	comparison with curriculum-based approaches. In		
	<i>Proceedings of the 19th Conference of the European</i>		
	<i>Chapter of the Association for Computational Lin-</i>		
	<i>guistics (Volume 2: Short Papers)</i> , pages 307–321,		
	Rabat, Morocco. Association for Computational Lin-		
	guistics.		
	Lars Malmqvist. 2024. Sycophancy in Large Language		
	Models: Causes and Mitigations. <i>arXiv preprint</i>		
	arXiv:2411.15287.		
	Potsawee Manakul, Adian Liusie, and Mark Gales. 2023.		
	SelfCheckGPT: Zero-resource black-box hallucina-		
	tion detection for generative large language models.		
	In <i>Proceedings of the 2023 Conference on Empiri-</i>		
	<i>cal Methods in Natural Language Processing</i> , pages		
	9004–9017, Singapore. Association for Computa-		
	tional Linguistics.		
	Hugo Mercier. 2016. The argumentative theory: Pre-		
	dictions and empirical evidence. <i>Trends in Cognitive</i>		
	<i>Sciences</i> , 20(9):689–700.		
	Hugo Mercier and Dan Sperber. 2011. Why do humans		
	reason? Arguments for an argumentative theory. <i>Be-</i>		
	<i>havioral and Brain Sciences</i> , 34(2):57–74.		
	Hugo Mercier and Dan Sperber. 2017. <i>The Enigma of</i>		
	<i>Reason.</i> Harvard University Press.		
	Julian Michael, Salsabila Mahdi, David Rein, Jack-		
	son Petty, Julien Dirani, Vishakh Padmakumar, and		
	Samuel R. Bowman. 2023. Debate helps supervise		
	unreliable experts. <i>Preprint</i> , arXiv:2311.08702.		

- Erik Miehl, Michael Desmond, Karthikeyan Nathe-
san Ramamurthy, Elizabeth M. Daly, Kush R. Varsh-
ney, Eitan Farchi, Pierre Dognin, Jesus Rios, Djallel
Bouneffouf, Miao Liu, and Prasanna Sattigeri. 2025.
[Evaluating the prompt steerability of large language
models](#). In *Proceedings of the 2025 Conference of
the Nations of the Americas Chapter of the Associ-
ation for Computational Linguistics: Human Lan-
guage Technologies (Volume 1: Long Papers)*, pages
7874–7900.
- Paul Milgrom and John Roberts. 1986. [Relying on the
information of interested parties](#). *The RAND Journal
of Economics*, 17(1):18–32.
- Roger B. Myerson. 1986. [Multistage games with com-
munication](#). *Econometrica*, 54(2):323–358.
- Charlan Nemeth, Keith Brown, and John Rogers. 2001.
[Devil’s advocate versus authentic dissent: stimulating
quantity and quality](#). *European Journal of Social
Psychology*, 31(6):707–720.
- Nafiseh Nikeghbal, Amir Hossein Kargaran,
Shaghayegh Kolli, and Jana Diesner. 2026.
[Who flips? self- and cross-model counterarguments
reveal answer instability in LLMs](#). *Preprint*,
arXiv:2606.16011. Findings of EMNLP 2026, to
appear.
- Shmuel Nitzan and Jacob Paroush. 1982. [Optimal deci-
sion rules in uncertain dichotomous choice situations](#).
International Economic Review, 23(2):289–297.
- OpenAI. 2025. [Sycophancy in GPT-4o: What Hap-
pened and What We’re Doing about It](#).
- OpenAI. 2026. [GPT-5.4 Nano](#).
- Ethan Perez, Sam Ringer, Kamile Lukosiute, Karina
Nguyen, Edwin Chen, Scott Heiner, Craig Pettit,
Catherine Olsson, Sandipan Kundu, Saurav Kada-
vath, and 1 others. 2023. [Discovering language
model behaviors with model-written evaluations](#). In
*Findings of the Association for Computational Lin-
guistics: ACL 2023*, pages 13387–13434.
- Ivo Petrov, Jasper Dekoninck, and Martin Vechev.
2025. BrokenMath: A benchmark for sycophancy
in theorem proving with LLMs. *arXiv preprint
arXiv:2510.04721*. 451 expert-verified adversarial
math problems; four-category sycophancy judge.
- Giorgio Piatti, Zhijing Jin, Max Kleiman-Weiner, Bern-
hard Schölkopf, Mrinmaya Sachan, and Rada Mi-
halcea. 2024. [Cooperate or collapse: Emergence of
sustainable cooperation in a society of LLM agents](#).
In *Advances in Neural Information Processing Sys-
tems 37 (NeurIPS 2024)*.
- Karl Popper. 2011. *The Open Society and Its Enemies*.
Routledge Classics. Routledge. First published 1945.
- Pradyumna Shyama Prasad and Minh Nhat Nguyen.
2025. [When two LLMs debate, both think they’ll
win](#). *Preprint*, arXiv:2505.19184. ArXiv preprint
arXiv:2505.19184 (v3, 9 June 2025).
- Dražen Prelec, H. Sebastian Seung, and John McCoy.
2017. [A solution to the single-question crowd wis-
dom problem](#). *Nature*, 541(7638):532–535.
- Chen Qian. 2026. [A layered analysis of disagree-
ment and answer quality in multi-agent LLM de-
bate](#). *Preprint*, arXiv:2609.08016. ArXiv preprint
arXiv:2609.08016, submitted 7 September 2026.
- Chen Qian, Zihao Xie, YiFei Wang, Wei Liu, Kunlun
Zhu, Hanchen Xia, Yufan Dang, Zhuoyun Du, Weize
Chen, Cheng Yang, Zhiyuan Liu, and Maosong
Sun. 2025. [Scaling large language model-based
multi-agent collaboration](#). In *The Thirteenth Inter-
national Conference on Learning Representations
(ICLR 2025)*.
- Salman Rahman, Sherif Issaka, Ashima Suvarna,
Genglin Liu, James Shiffer, Jaeyoung Lee,
Md Rizwan Parvez, Hamid Palangi, Shi Feng,
Nanyun Peng, Yejin Choi, Julian Michael, Liwei
Jiang, and Saadia Gabriel. 2025. [AI debate
aids assessment of controversial claims](#). In
*Advances in Neural Information Processing
Systems 38 (NeurIPS 2025)*. Poster, confirmed
at <https://neurips.cc/virtual/2025/poster/117257>;
arXiv:2506.02175.
- Nina Rimskey, Nick Gabrieli, Julian Schulz, Meg Tong,
Evan Hubinger, and Alexander Turner. 2024. [Steer-
ing llama 2 via contrastive activation addition](#). In
*Proceedings of the 62nd Annual Meeting of the As-
sociation for Computational Linguistics (Volume 1:
Long Papers)*, pages 15504–15522.
- Leonard Salewski, Stephan Alaniz, Isabel Rio-Torto,
Eric Schulz, and Zeynep Akata. 2023. [In-context im-
personation reveals large language models’ strengths
and biases](#). In *Advances in Neural Information Pro-
cessing Systems 36 (NeurIPS 2023)*.
- Chris William Sanchirico. 2001. [Relying on the in-
formation of interested—and potentially dishonest—
parties](#). *American Law and Economics Review*,
3(2):320–357.
- Shibani Santurkar, Esin Durmus, Faisal Ladhak, Cino
Lee, Percy Liang, and Tatsunori Hashimoto. 2023.
[Whose opinions do language models reflect?](#) In
*Proceedings of the 40th International Conference on
Machine Learning*, volume 202 of *Proceedings of
Machine Learning Research*, pages 29971–30004.
- David M. Schweiger, William R. Sandberg, and
James W. Ragan. 1986. [Group approaches for
improving strategic decision making: A compara-
tive analysis of dialectical inquiry, devil’s advocacy,
and consensus](#). *Academy of Management Journal*,
29(1):51–71.
- Charles R. Schwenk. 1990. [Effects of devil’s advo-
cacy and dialectical inquiry on decision making: A
meta-analysis](#). *Organizational Behavior and Human
Decision Processes*, 47(1):161–176.

1978	Arya Shah, Deepali Mishra, and Chaklam Silpasuwanchai. 2026. Too nice to tell the truth: Quantifying agreeableness-driven sycophancy in role-playing language models . In <i>Proceedings of the 64th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)</i> , pages 30788–30801.	2032
1979		2033
1980		2034
1981		2035
1982		2036
1983		
1984		
1985	Murray Shanahan, Kyle McDonell, and Laria Reynolds. 2023. Role play with large language models . <i>Nature</i> , 623:493–498.	2037
1986		2038
1987		2039
1988		2040
1989	Lloyd Shapley and Bernard Grofman. 1984. Optimizing group judgmental accuracy in the presence of interdependencies . <i>Public Choice</i> , 43(3):329–343.	2041
1990		2042
1991		2043
1992		2044
1993		2045
1994		
1995	Mrinank Sharma, Meg Tong, Tomasz Korbak, David Duvenaud, Amanda Askell, Samuel R. Bowman, Newton Cheng, Esin Durmus, Zac Hatfield-Dodds, Scott R. Johnston, Shauna Kravec, Timothy Maxwell, Sam McCandlish, Kamal Ndousse, Oliver Rausch, Nicholas Schiefer, Da Yan, Miranda Zhang, and Ethan Perez. 2024. Towards Understanding Sycophancy in Language Models. In <i>International Conference on Learning Representations (ICLR)</i> . ArXiv:2310.13548.	2046
1996		2047
1997		2048
1998		2049
1999		2050
2000		2051
2001		
2002	Xu Shen, Yixin Liu, Yiwei Dai, Yili Wang, Rui Miao, Yue Tan, Shirui Pan, and Xin Wang. 2025. Understanding the information propagation effects of communication topologies in LLM-based multi-agent systems . In <i>Proceedings of the 2025 Conference on Empirical Methods in Natural Language Processing</i> , pages 12347–12361.	2052
2003		2053
2004		2054
2005		2055
2006		2056
2007		2057
2008	Hyun Song Shin. 1998. Adversarial and inquisitorial procedures in arbitration . <i>The RAND Journal of Economics</i> , 29(2):378–405.	2058
2009		2059
2010		2060
2011		2061
2012	Andries Petrus Smit, Nathan Grinsztajn, Paul Duckworth, Thomas D. Barrett, and Arnu Pretorius. 2024. Should we be going MAD? a look at multi-agent debate strategies for LLMs . In <i>Proceedings of the 41st International Conference on Machine Learning</i> , volume 235 of <i>PMLR</i> , pages 45883–45905.	2062
2013		2063
2014		2064
2015		
2016		
2017	Joel Sobel. 1985. A theory of credibility . <i>The Review of Economic Studies</i> , 52(4):557–573.	2065
2018		2066
2019		2067
2020	Amir Taubenfeld, Yaniv Dover, Roi Reichart, and Ariel Goldstein. 2024. Systematic biases in LLM simulations of debates . In <i>Proceedings of the 2024 Conference on Empirical Methods in Natural Language Processing</i> , pages 251–267.	2068
2021		
2022		
2023		
2024	Emanuel Tewolde, Xiao Zhang, David Guzman Piedrahita, Vincent Conitzer, and Zhijing Jin. 2026. CoopEval: Benchmarking cooperation-sustaining mechanisms and LLM agents in social dilemmas . <i>Preprint</i> , arXiv:2604.15267. ICML 2026.	2069
2025		2070
2026		2071
2027		2072
2028		2073
2029	Emmanuel Trouche, Petter Johansson, Lars Hall, and Hugo Mercier. 2016. The selective laziness of reasoning . <i>Cognitive Science</i> , 40(8):2122–2136.	2074
2030		2075
2031		2076
	Emmanuel Trouche, Emmanuel Sander, and Hugo Mercier. 2014. Arguments, more than confidence, explain the good performance of reasoning groups . <i>Journal of Experimental Psychology: General</i> , 143(5):1958–1971.	2077
	Daniel Vennemeyer, Phan Anh Duong, Tiffany Zhan, and Tianyu Jiang. 2025. Sycophancy is not one thing: Causal separation of sycophantic behaviors in LLMs . <i>Preprint</i> , arXiv:2509.21305.	2078
		2079
		2080
		2081
		2082
	Junlin Wang, Jue Wang, Ben Athiwaratkun, Ce Zhang, and James Y. Zou. 2025. Mixture-of-agents enhances large language model capabilities . In <i>The Thirteenth International Conference on Learning Representations (ICLR 2025)</i> .	2083
		2084
		2085
		2086
	Mengdie Flora Wang, Haochen Xie, Guanghui Wang, Devin Zhang, and Jae Oh Woo. 2026. Budgeted act-or-defer multi-agent LLM deliberation with local reliability bounds . <i>Preprint</i> , arXiv:2606.29654. ArXiv preprint arXiv:2606.29654, submitted 28 June 2026.	
	Qineng Wang, Zihao Wang, Ying Su, Hanghang Tong, and Yangqiu Song. 2024. Rethinking the bounds of LLM reasoning: Are multi-agent discussions the key? In <i>Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)</i> , pages 6106–6131.	
	Xuezhi Wang, Jason Wei, Dale Schuurmans, Quoc V. Le, Ed H. Chi, Sharan Narang, Aakanksha Chowdhery, and Denny Zhou. 2023. Self-consistency improves chain of thought reasoning in language models . In <i>The Eleventh International Conference on Learning Representations (ICLR 2023)</i> . ArXiv:2203.11171.	
	Jerry Wei, Da Huang, Yifeng Lu, Denny Zhou, and Quoc V. Le. 2023. Simple synthetic data reduces sycophancy in large language models . <i>Preprint</i> , arXiv:2308.03958.	
	Zhiyuan Weng, Guikun Chen, and Wenguan Wang. 2025. Do as we do, not as you think: the conformity of large language models . In <i>The Thirteenth International Conference on Learning Representations (ICLR 2025)</i> .	
	Richard Willis, Yali Du, Joel Z. Leibo, and Michael Luck. 2025. Will systems of LLM agents cooperate: An investigation into a social dilemma . <i>Preprint</i> , arXiv:2501.16173.	
	Andrea Wynn, Harsh Satija, and Gillian Hadfield. 2025. Talk isn’t always cheap: Understanding failure modes in multi-agent debate . <i>Preprint</i> , arXiv:2509.05396. ICML 2025 Workshop on Multi-Agent Systems (per the arXiv comments field); arXiv:2509.05396.	
	Tim Wyse, Kaitlin Bustos, Yulia Volkova, and Max Kleiman-Weiner. 2026. Commitment to cooperation with self-negotiated contracts . <i>Preprint</i> , arXiv:2607.22750.	

- xAI. 2025. [Grok 4.1 Fast Reasoning](#).
- Miao Xiong, Zhiyuan Hu, Xinyang Lu, Yifei Li, Jie Fu, Junxian He, and Bryan Hooi. 2024. [Can LLMs express their uncertainty? an empirical evaluation of confidence elicitation in LLMs](#). In *The Twelfth International Conference on Learning Representations (ICLR 2024)*.
- Murong Yue, Jie Zhao, Min Zhang, Liang Du, and Ziyu Yao. 2024. [Large language model cascades with mixture of thought representations for cost-efficient reasoning](#). In *The Twelfth International Conference on Learning Representations (ICLR 2024)*.
- Hangfan Zhang, Zhiyao Cui, Jianhao Chen, Xinrun Wang, Qiaosheng Zhang, Zhen Wang, Dinghao Wu, and Shuyue Hu. 2025. [Stop overvaluing multi-agent debate – we must rethink evaluation and embrace model heterogeneity](#). *Preprint*, arXiv:2502.08788. V3, June 2025; v1 (February 2025) was titled “If Multi-Agent Debate is the Answer, What is the Question?”.
- Jintian Zhang, Xin Xu, Ningyu Zhang, Ruibo Liu, Bryan Hooi, and Shumin Deng. 2024. [Exploring collaboration mechanisms for LLM agents: A social psychology view](#). In *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 14544–14607.
- Kehan Zheng, Jinfeng Zhou, and Hongning Wang. 2025. [Beyond Nash equilibrium: Bounded rationality of LLMs and humans in strategic decision-making](#). *Preprint*, arXiv:2506.09390.
- Mingqian Zheng, Jiaxin Pei, Lajanugen Logeswaran, Moontae Lee, and David Jurgens. 2024. [When “a helpful assistant” is not really helpful: Personas in system prompts do not improve performances of large language models](#). In *Findings of the Association for Computational Linguistics: EMNLP 2024*, pages 15126–15154.
- Xiaochen Zhu, Caiqi Zhang, Yizhou Chi, Tom Stafford, Nigel Collier, and Andreas Vlachos. 2026. [Demystifying multi-agent debate: The role of confidence and diversity](#). In *Findings of the Association for Computational Linguistics: ACL 2026*, pages 33909–33930.
- Mingchen Zhuge, Wenyi Wang, Louis Kirsch, Francesco Faccio, Dmitrii Khizbullin, and Jürgen Schmidhuber. 2024. [GPTSwarm: Language agents as optimizable graphs](#). In *Proceedings of the 41st International Conference on Machine Learning*, volume 235 of *PMLR*, pages 62743–62767.